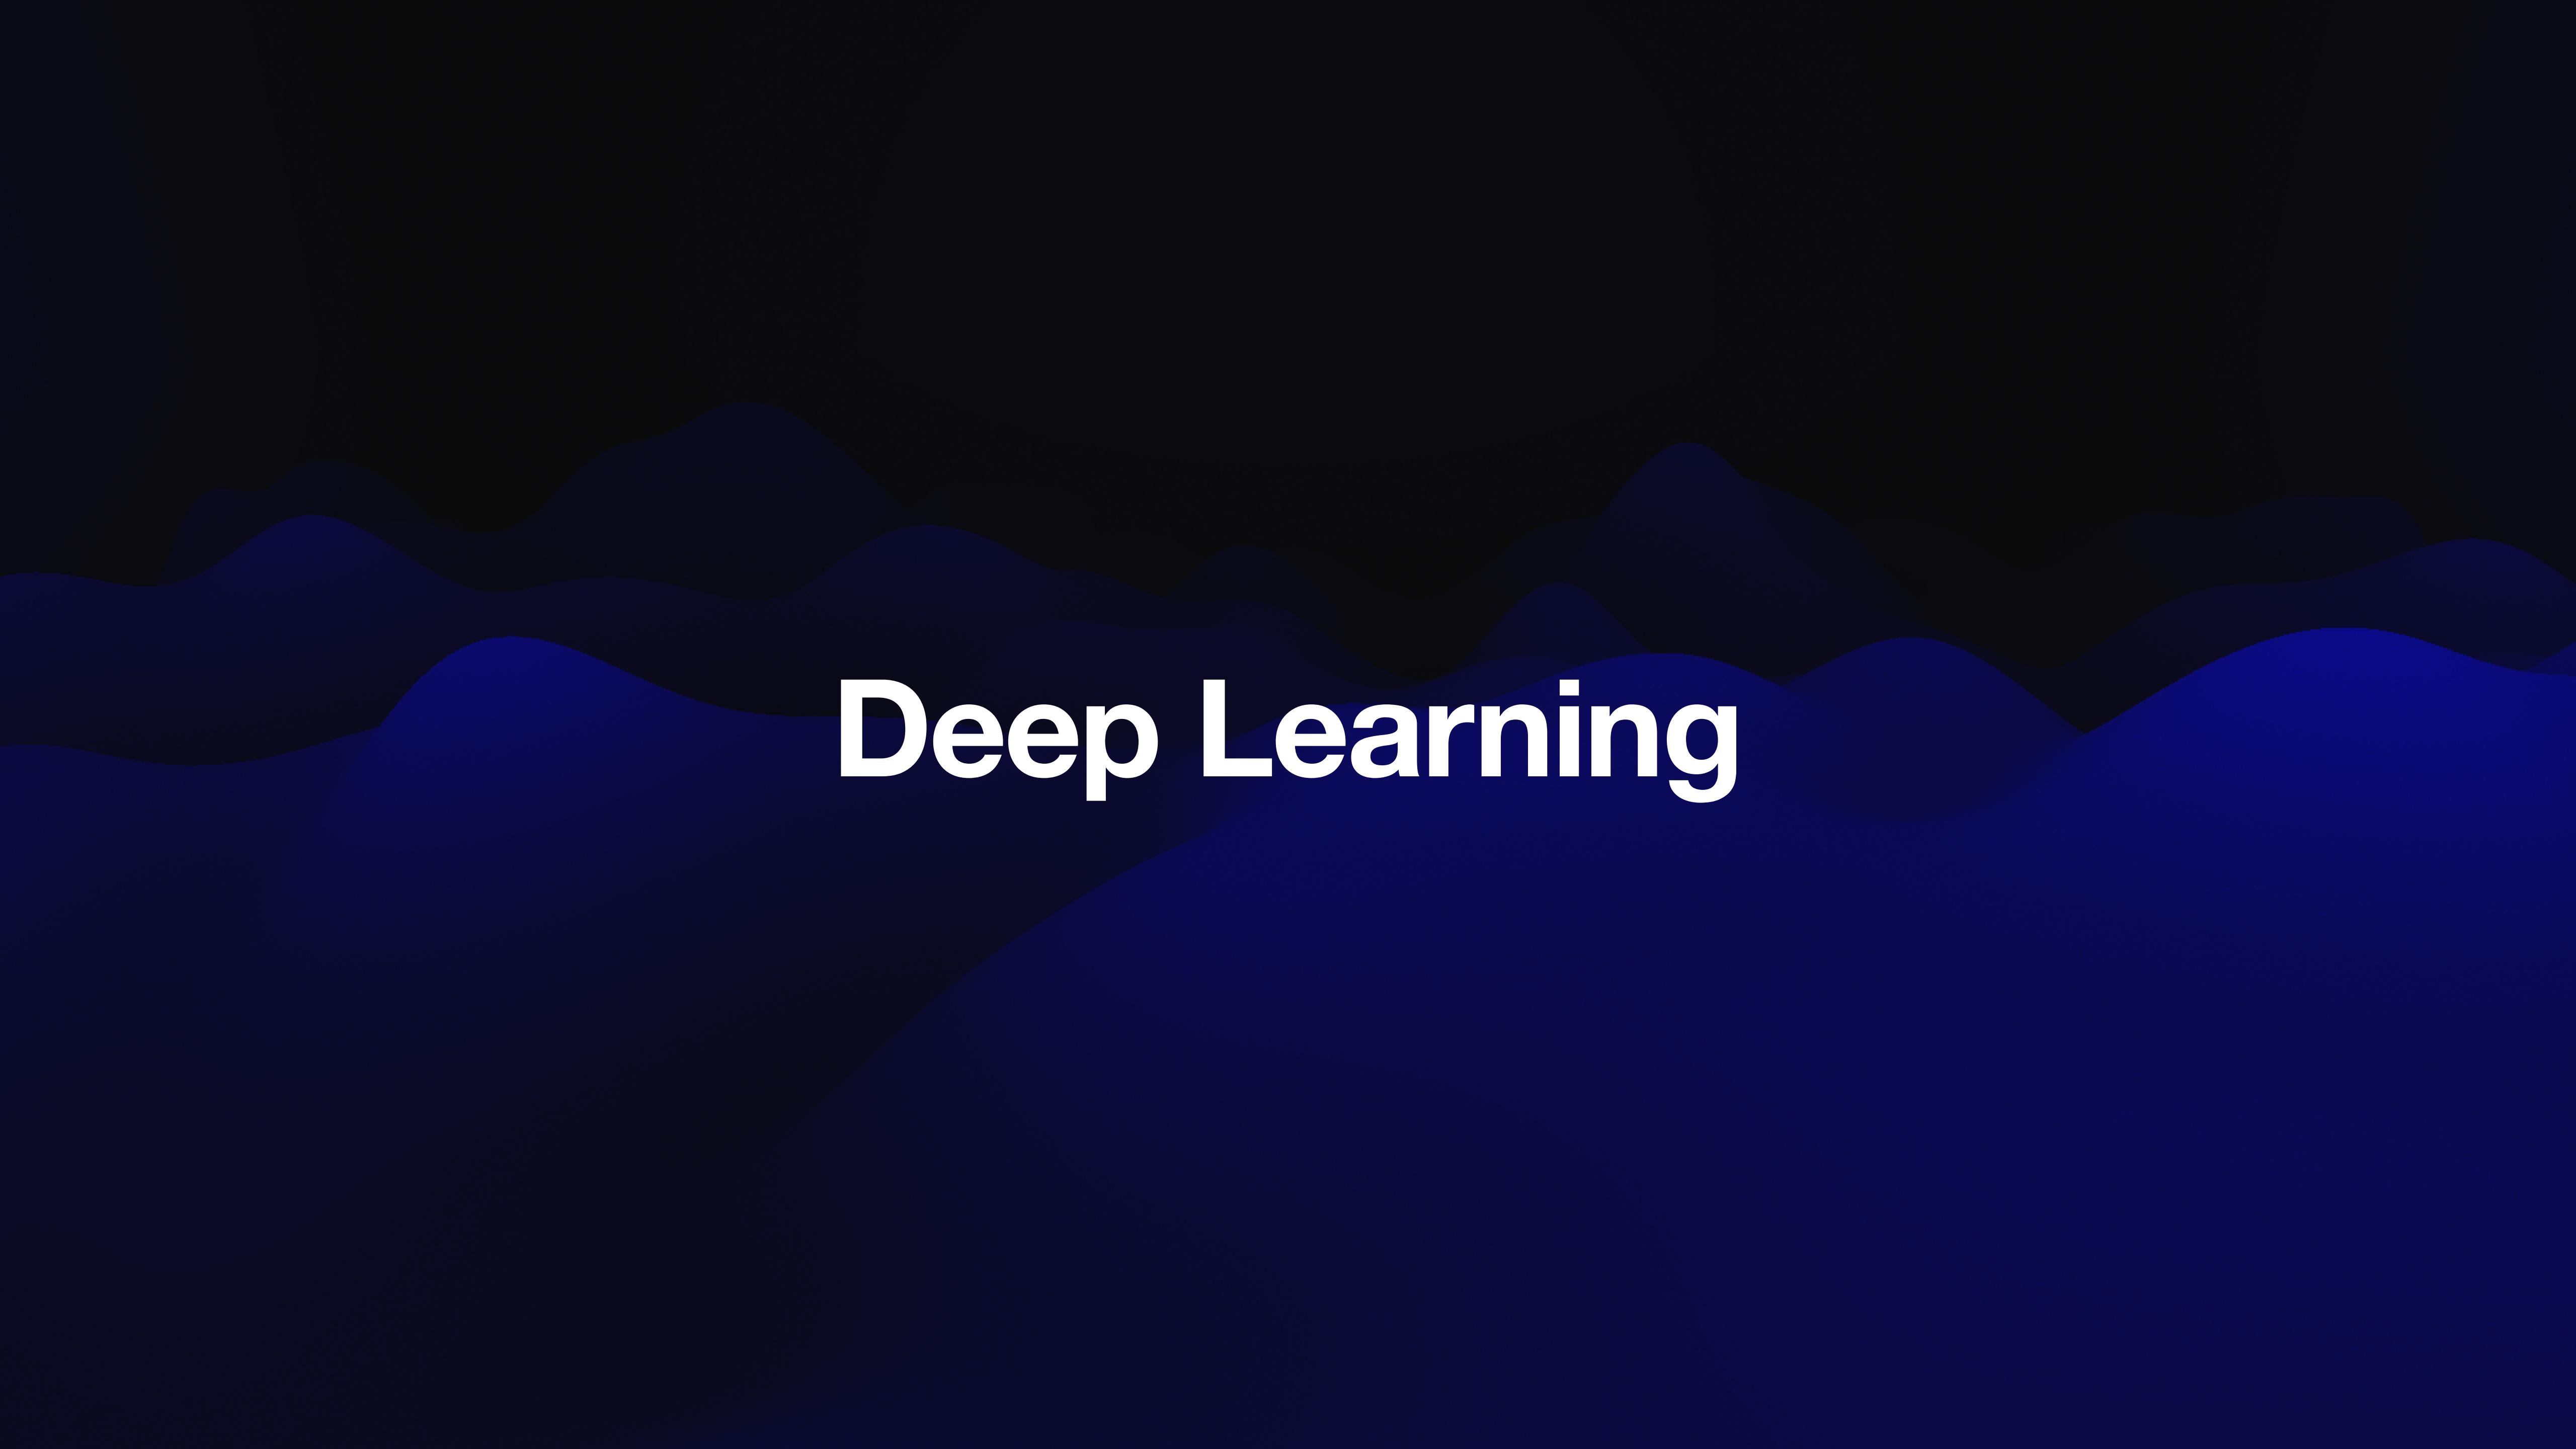


The Neural Network Zoo

A Survey of Neural Network Architectures
PICSciE/PRC Workshop Series

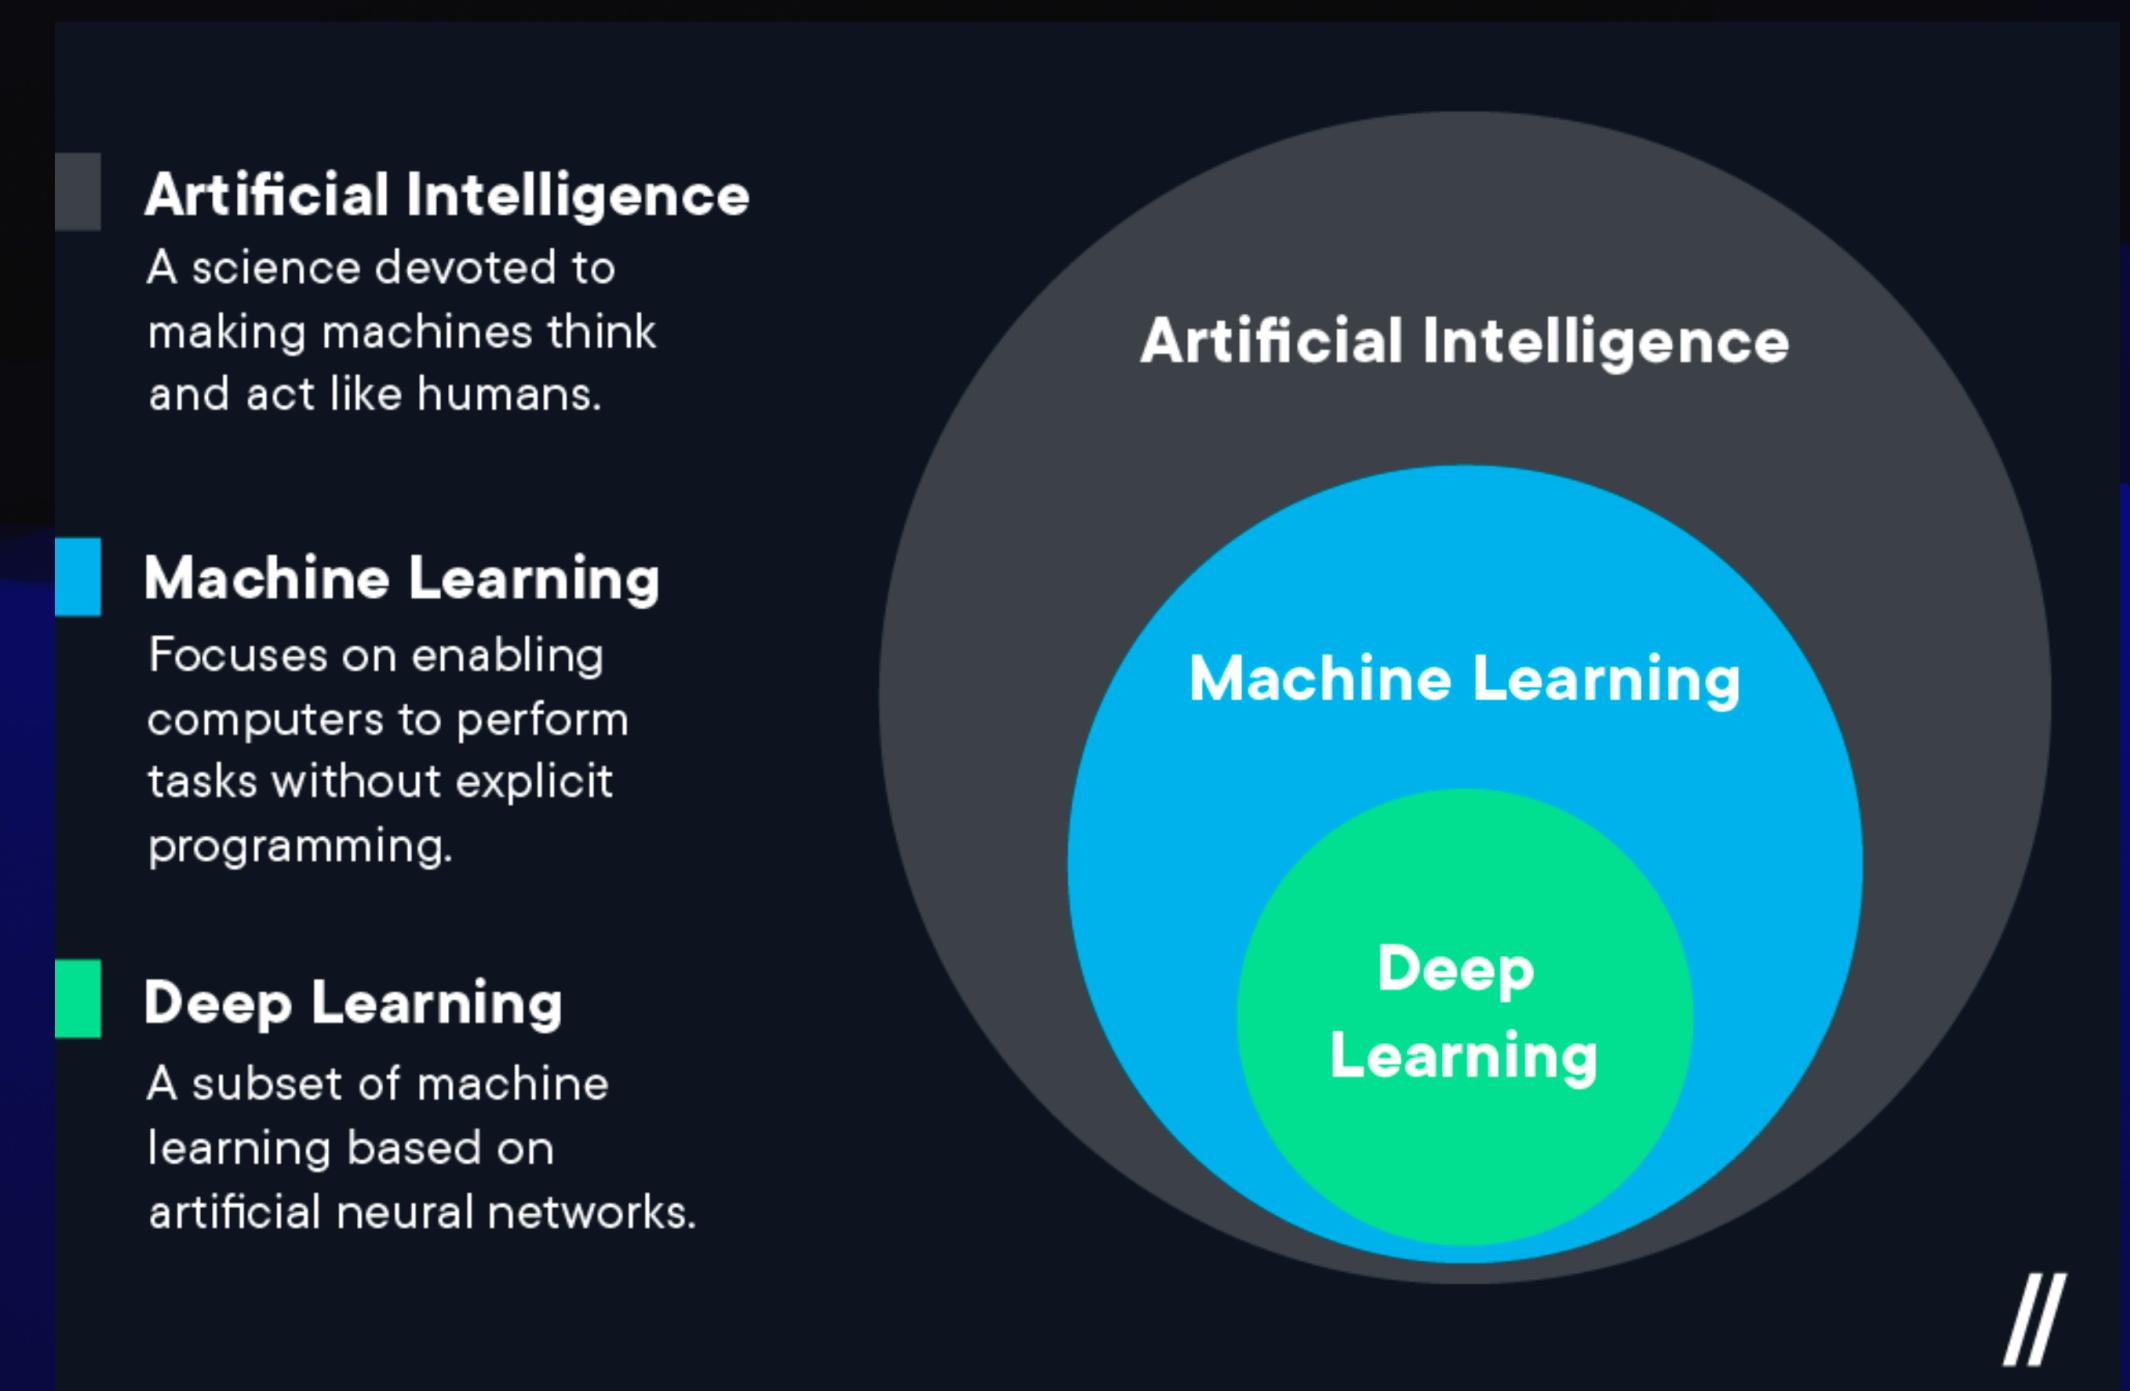
Gage DeZoort | 3/25/24

1. Intro to DL Concepts
2. The Neural Network Zoo
 - A. Deep Neural Networks (DNNs)
 - B. Convolutional Neural Networks (CNNs)
 - C. Recurrent Neural Networks (RNNs)
 - D. Graph Neural Networks (GNNs)
 - E. Generative Adversarial Network (GANs)
 - F. Autoencoders (AEs)
 - G. Variational Autoencoders (VAEs)

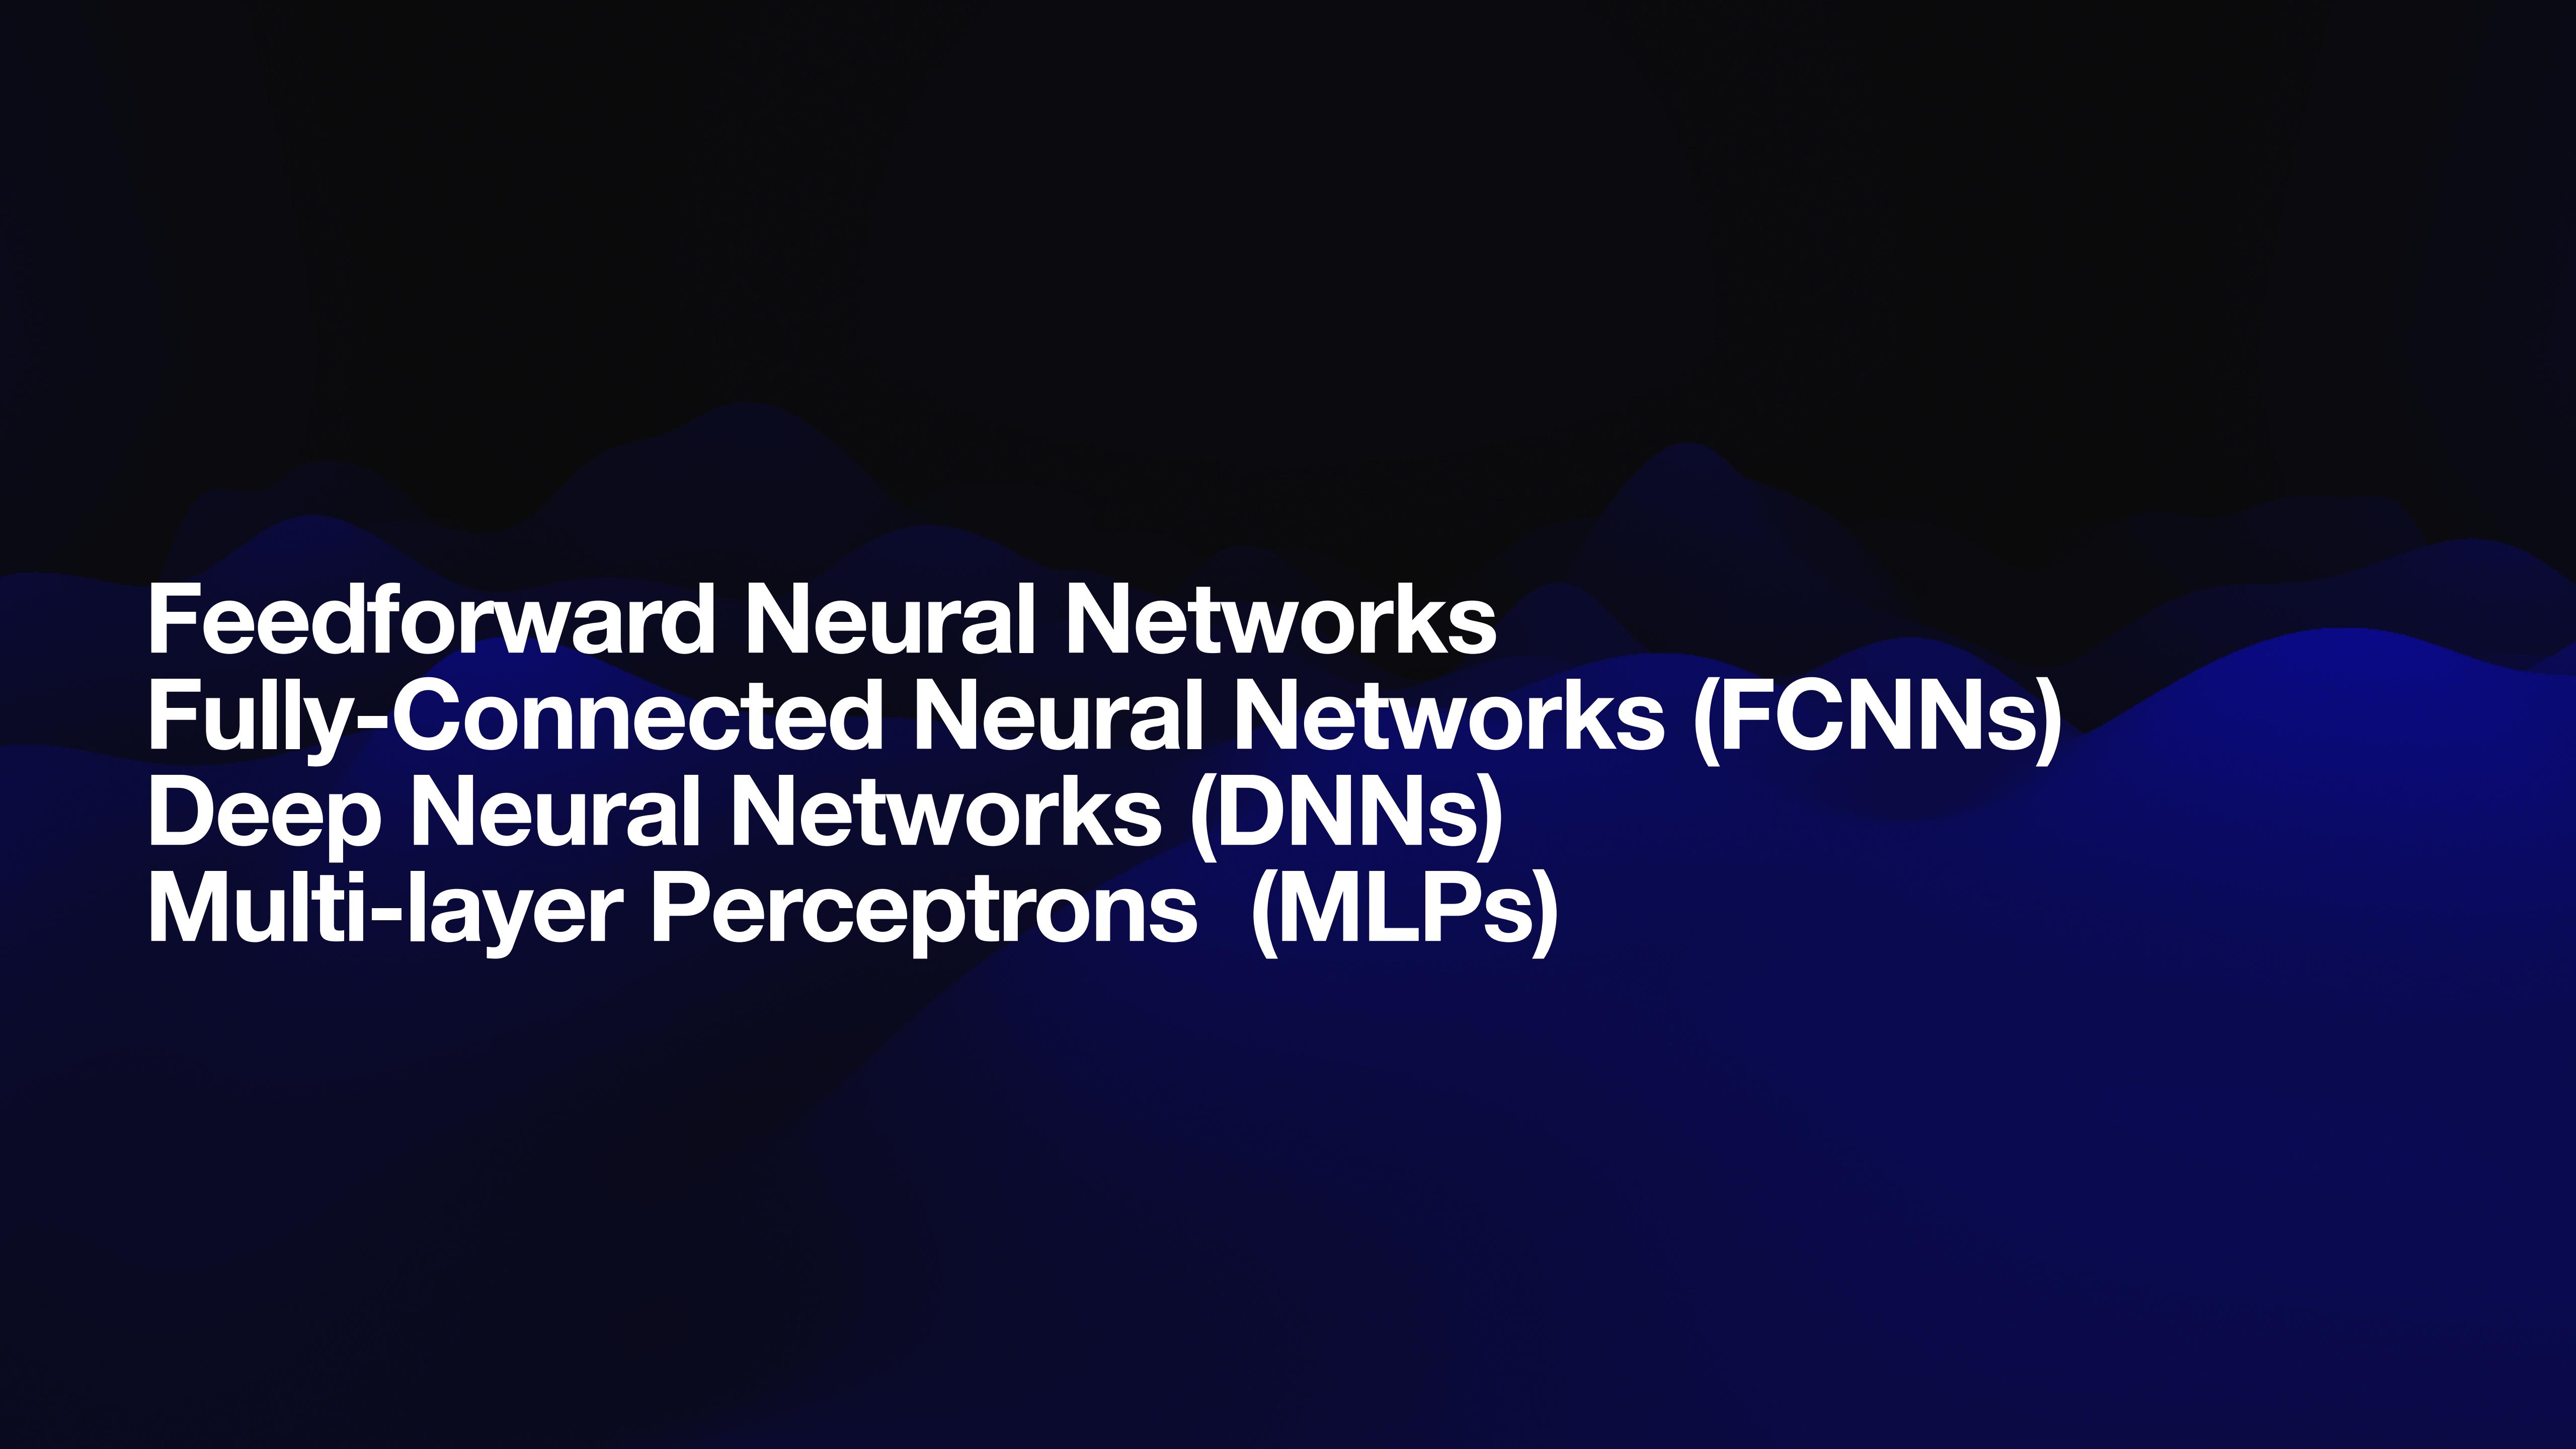
The background features a dark navy blue gradient with three distinct wavy layers. The top layer is a solid dark blue. Below it is a layer with a medium-dark blue gradient, and the bottom layer has a medium-light blue gradient. The waves are smooth and organic in shape.

Deep Learning

- Deep Learning (DL) algorithms leverage *depth* to sequentially construct complicated concepts from simpler concepts
- Importance of depth:
 - Mathematical view → composition of many simple functions leads to rich expressivity
 - Computational view → sequential computer program with communicative intermediate states



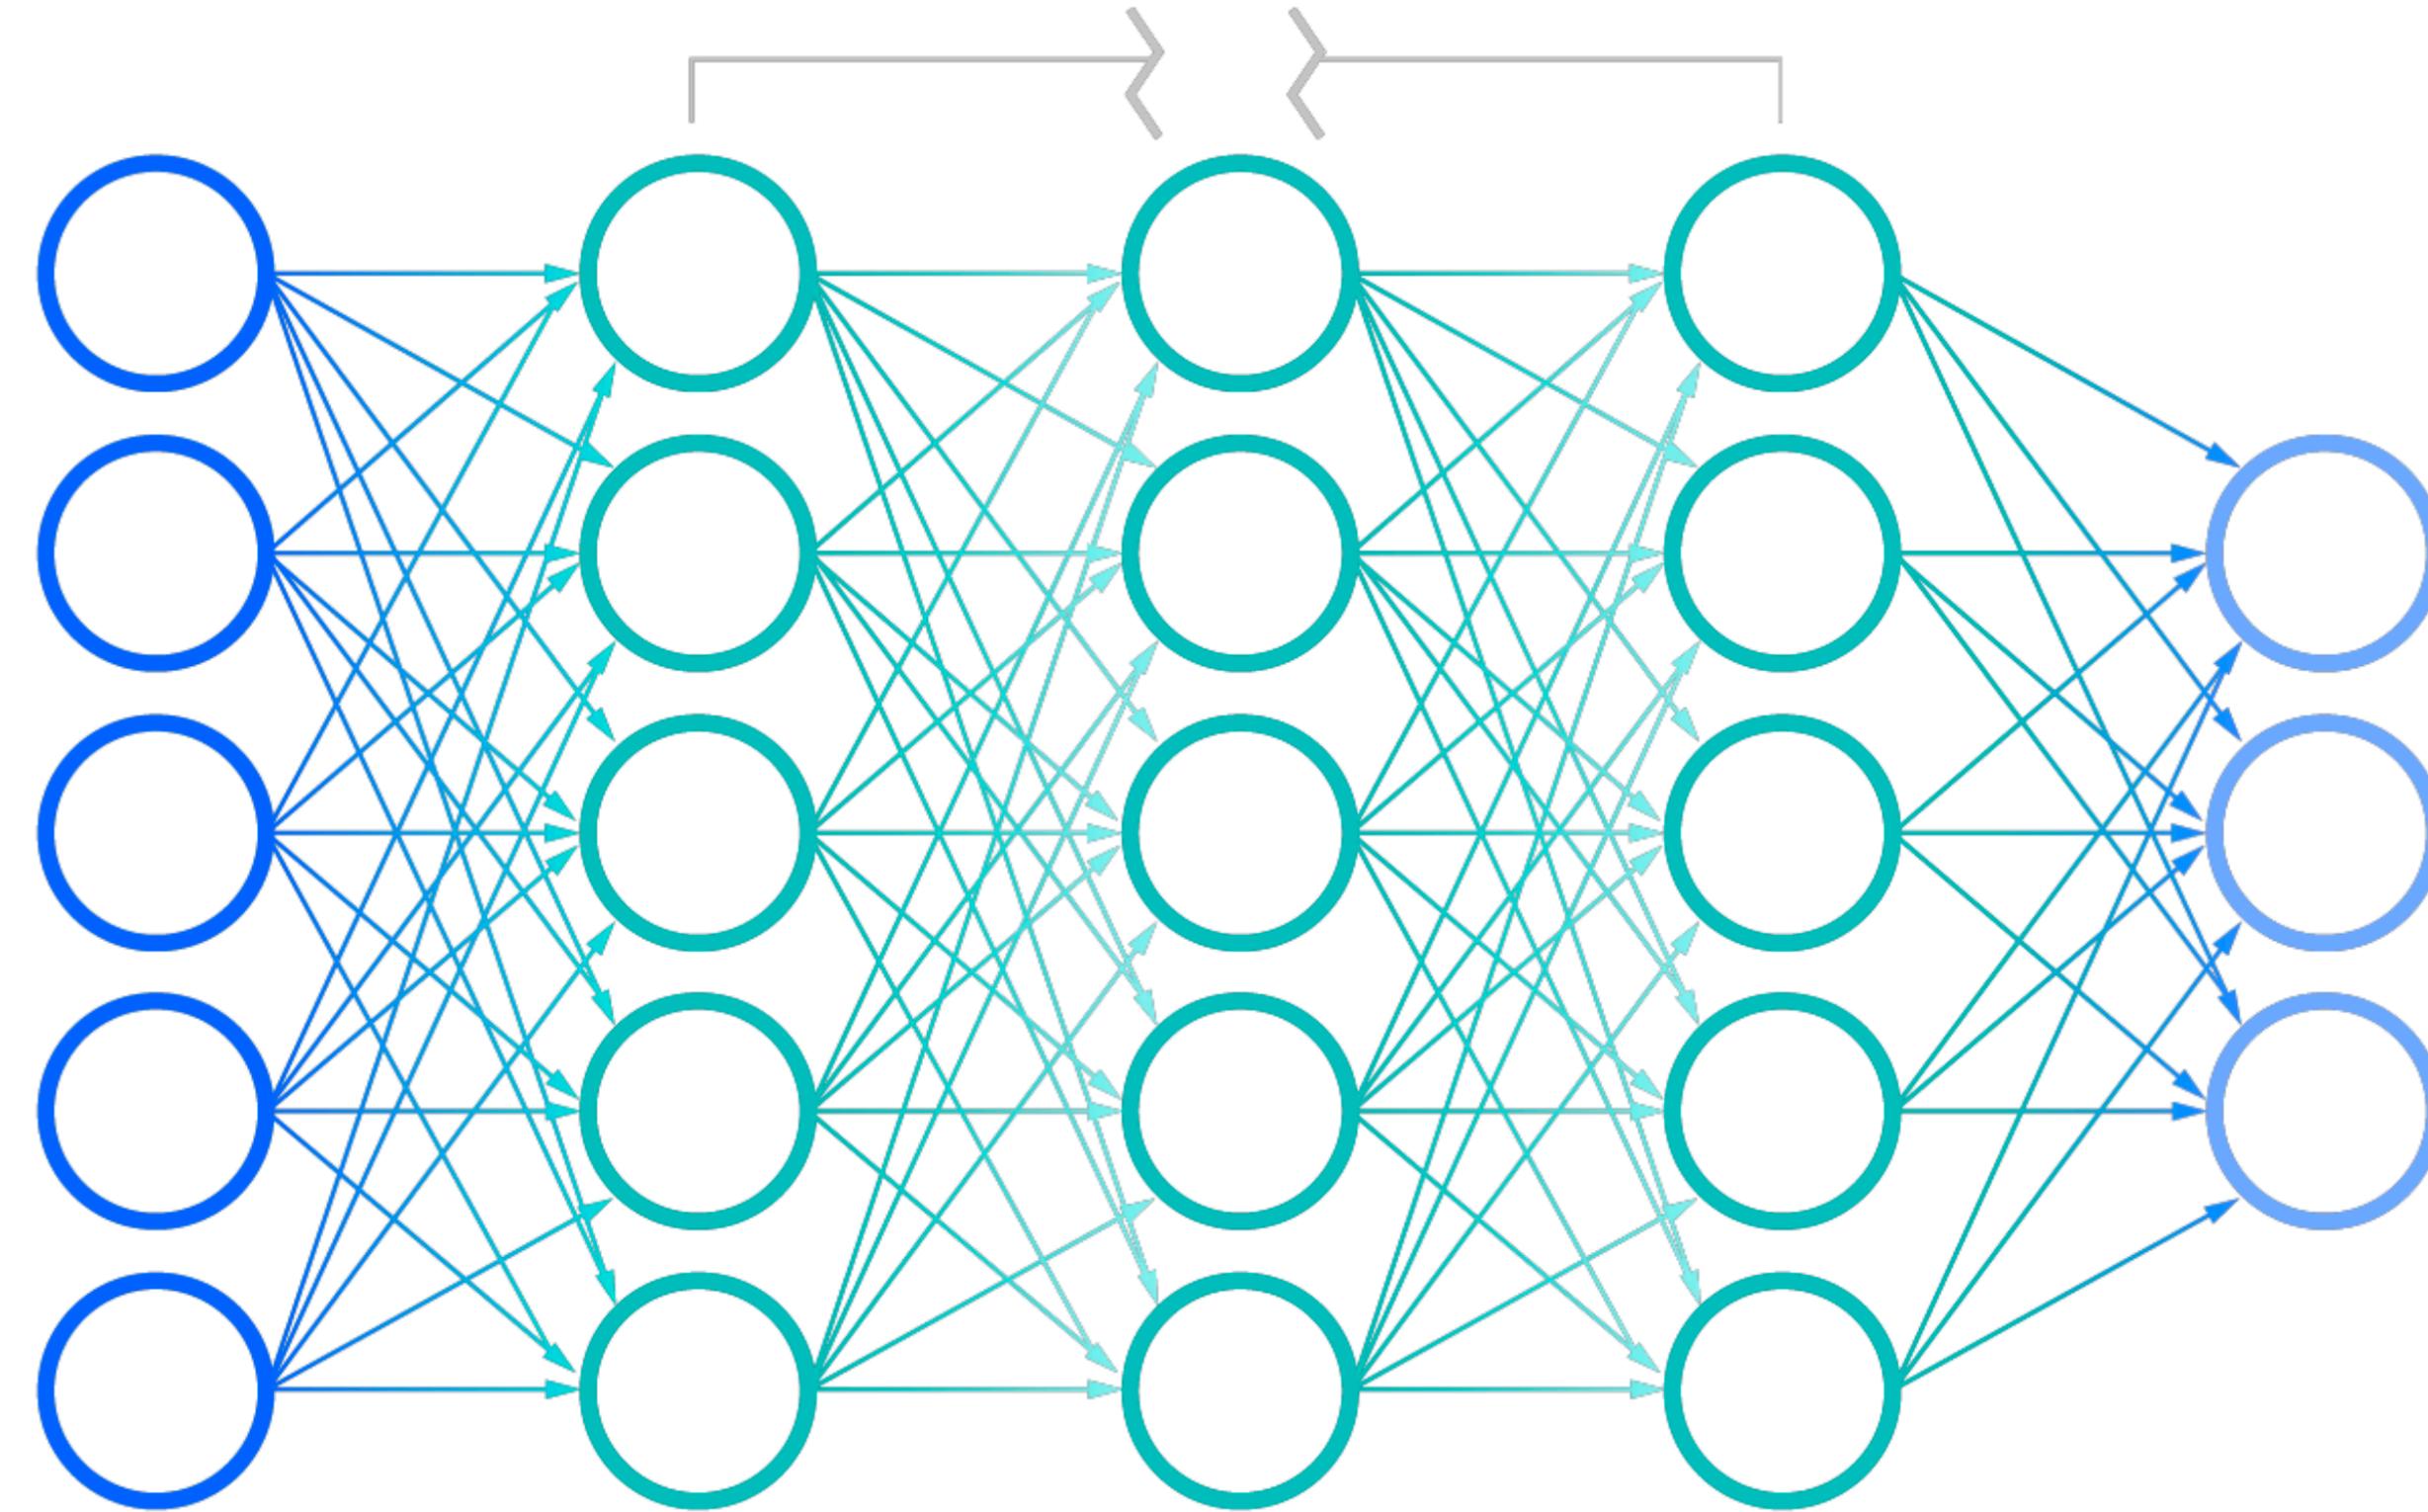
//

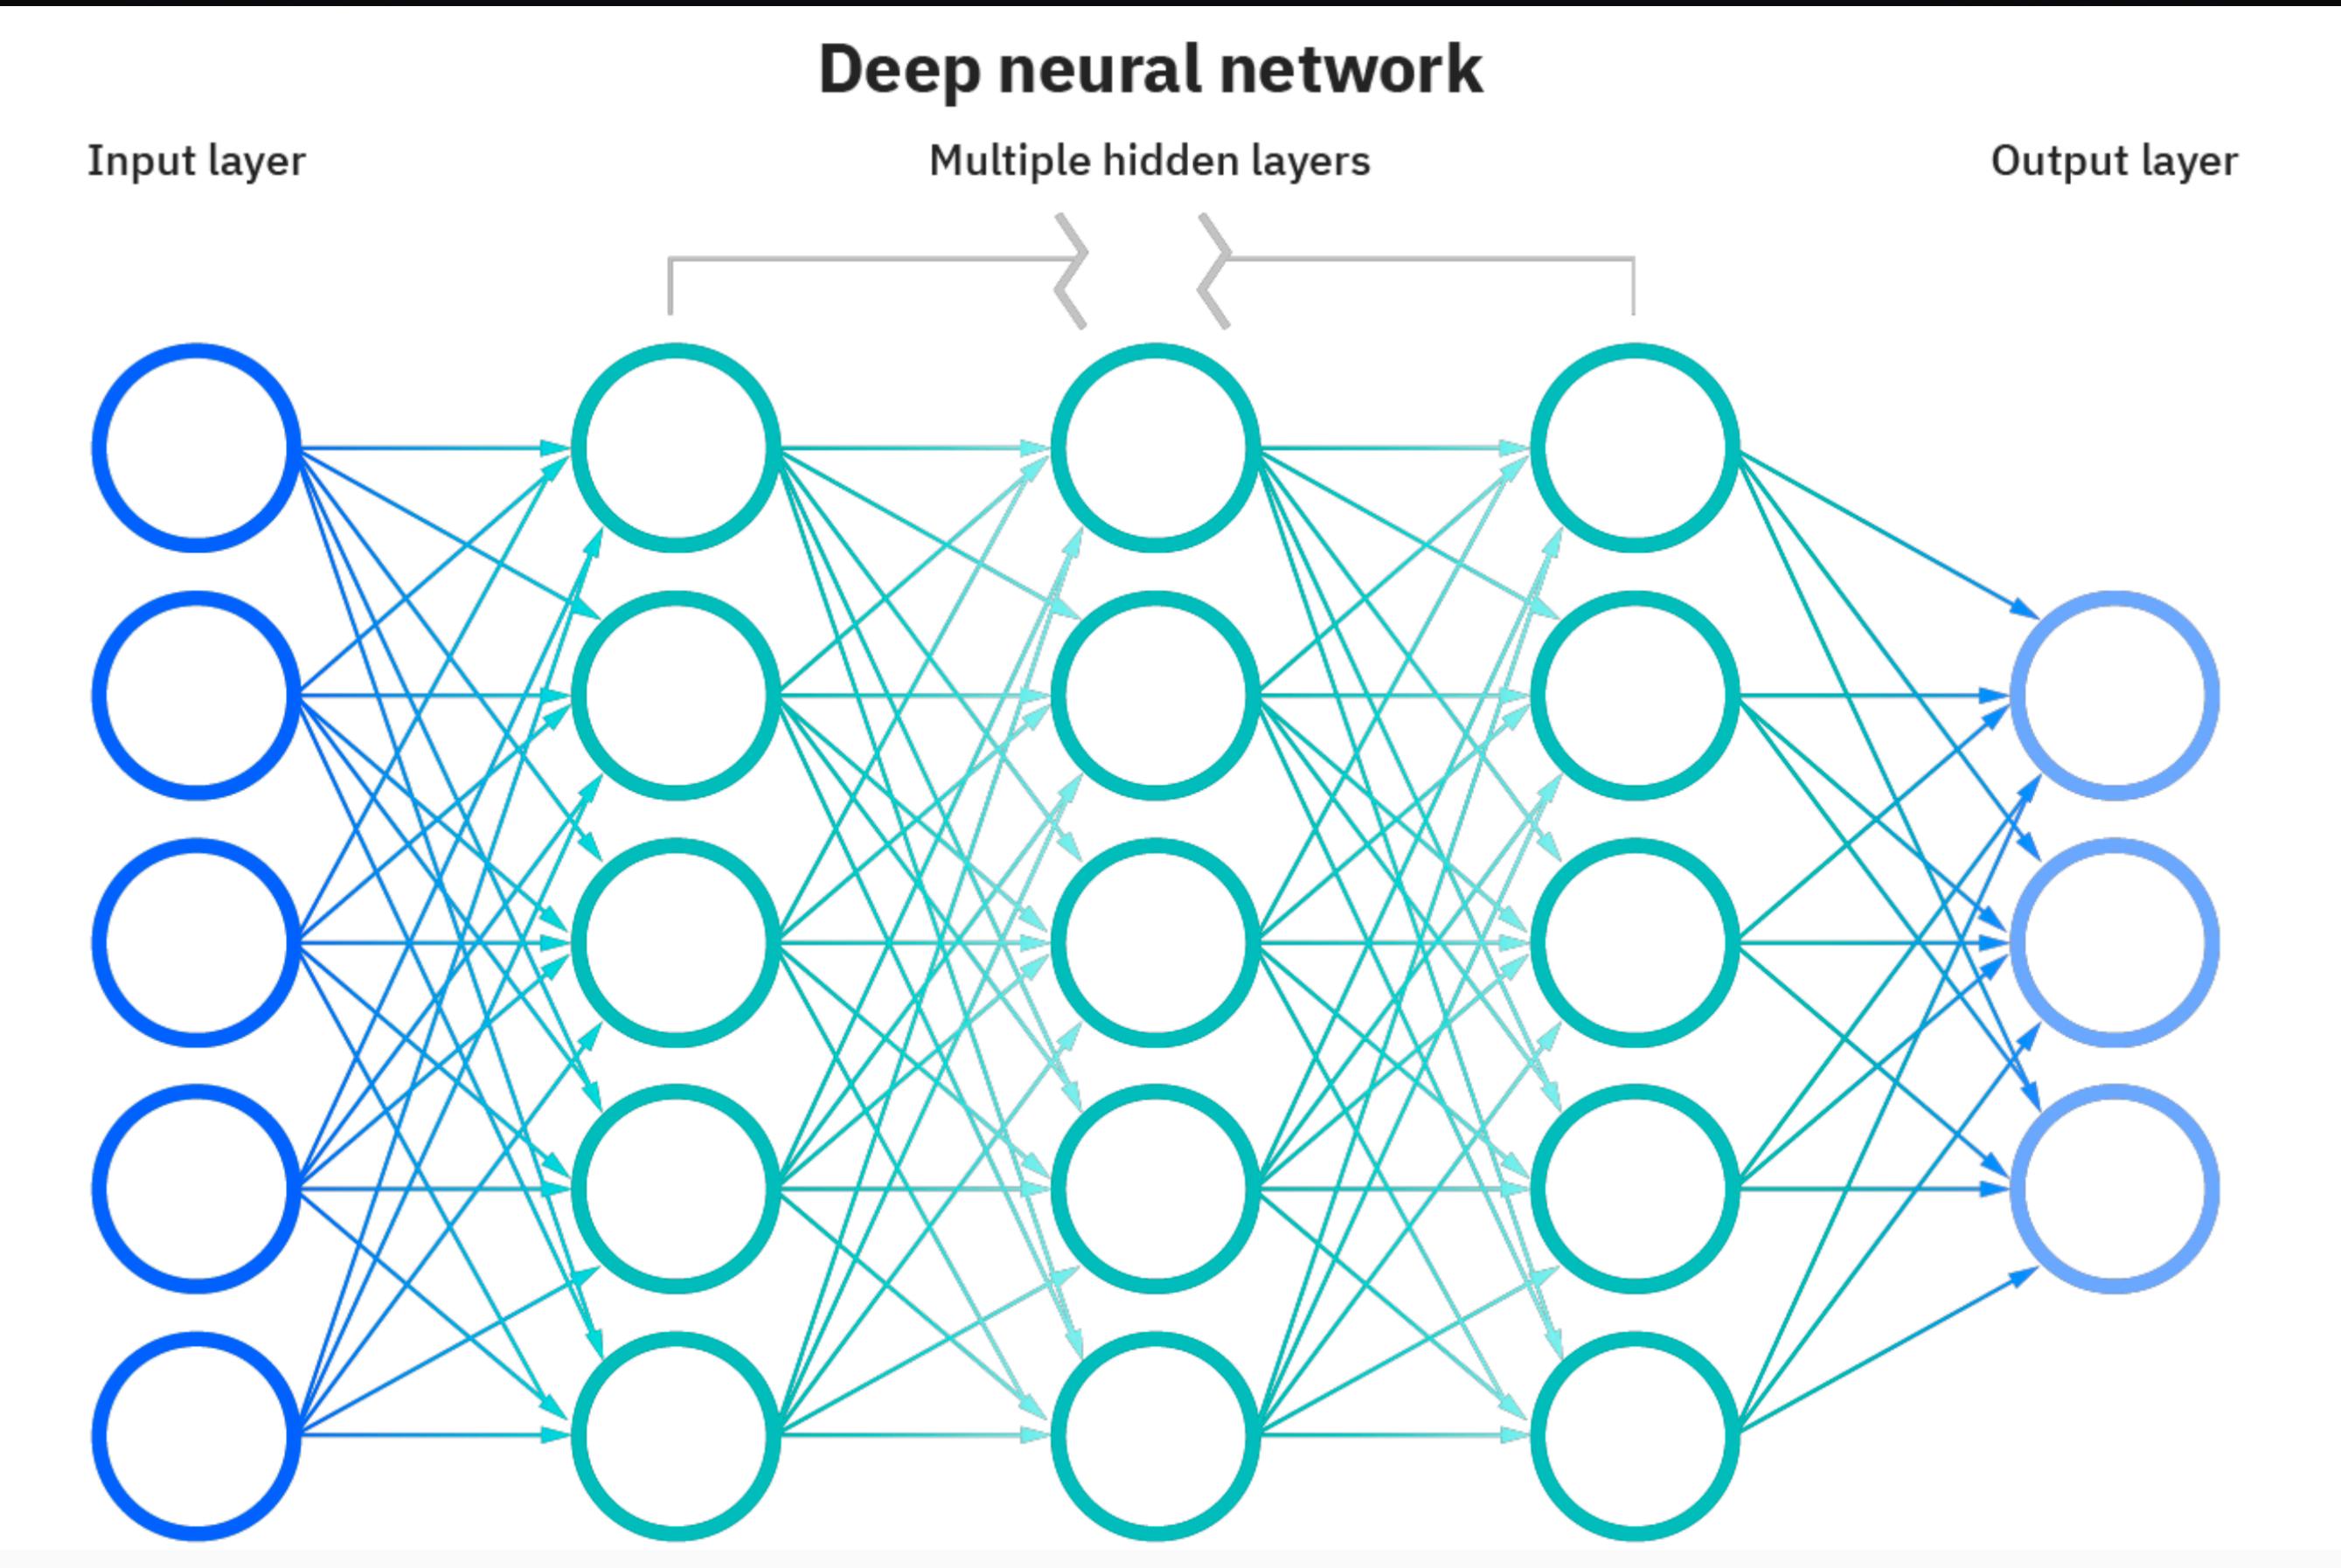


Feedforward Neural Networks
Fully-Connected Neural Networks (FCNNs)
Deep Neural Networks (DNNs)
Multi-layer Perceptrons (MLPs)

Deep neural network

Input layer Multiple hidden layers Output layer

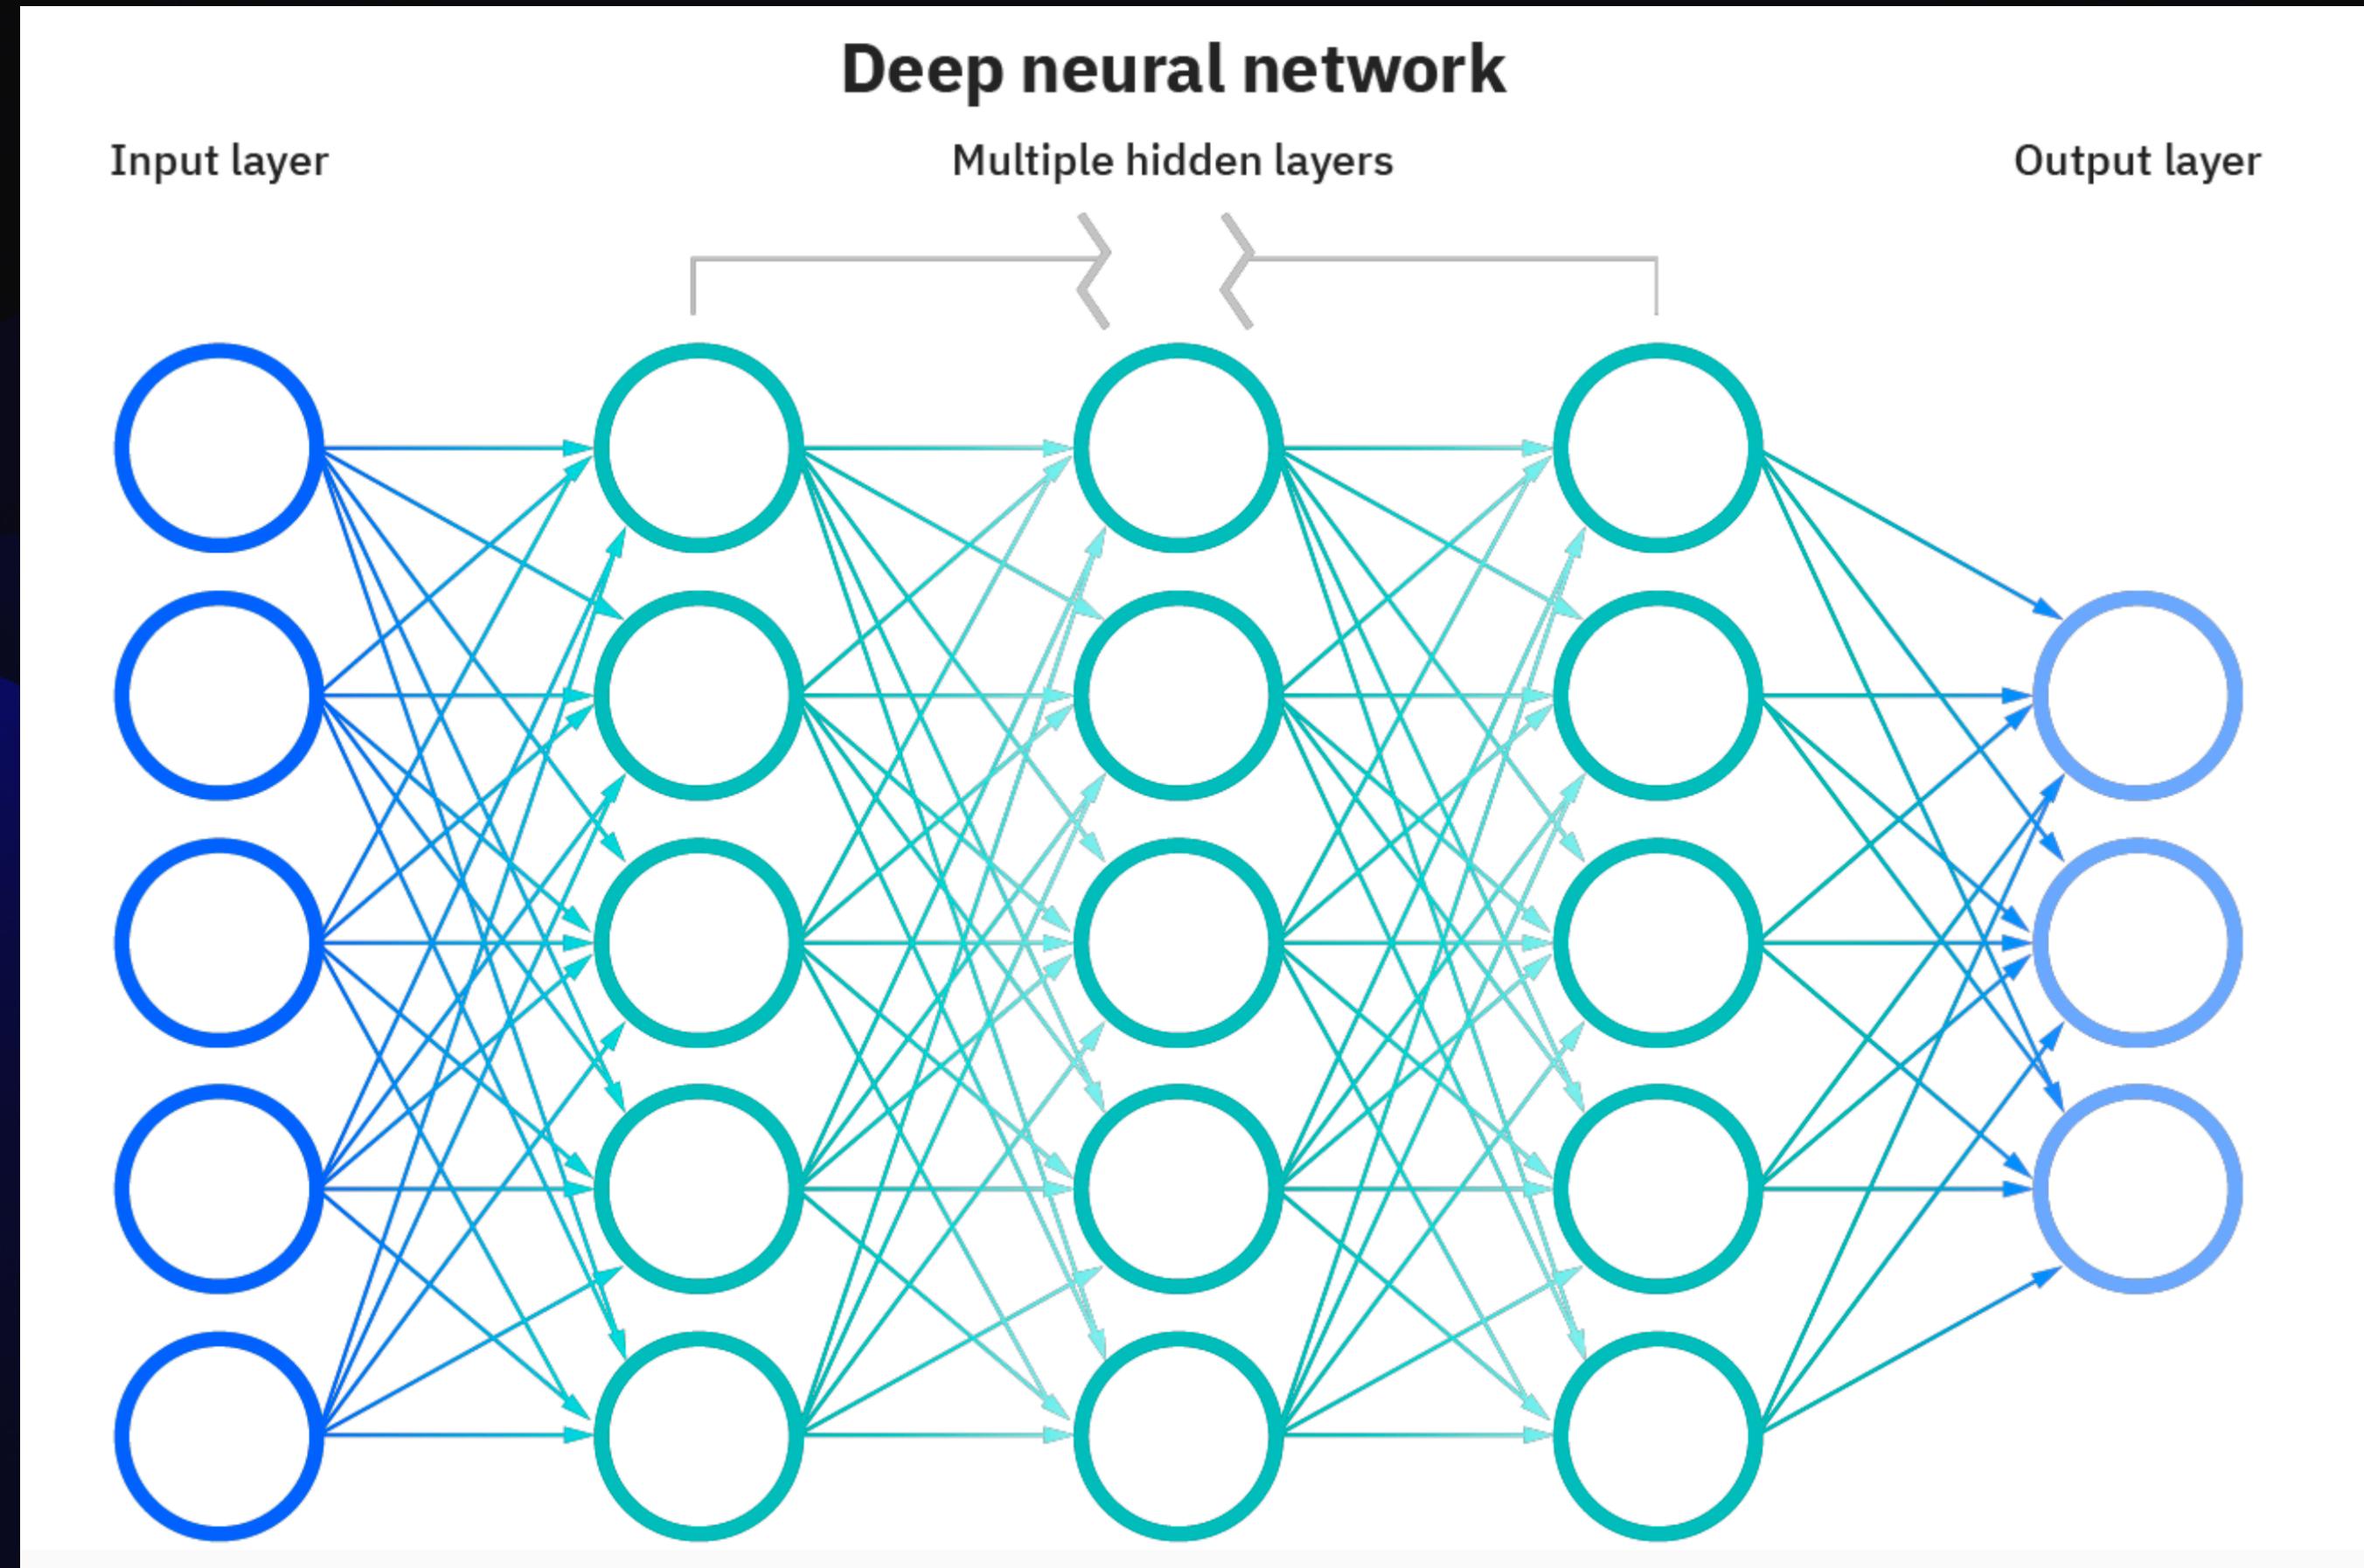




Inputs: images, vector features,
words, sentences, etc

E.g. a flattened greyscale image

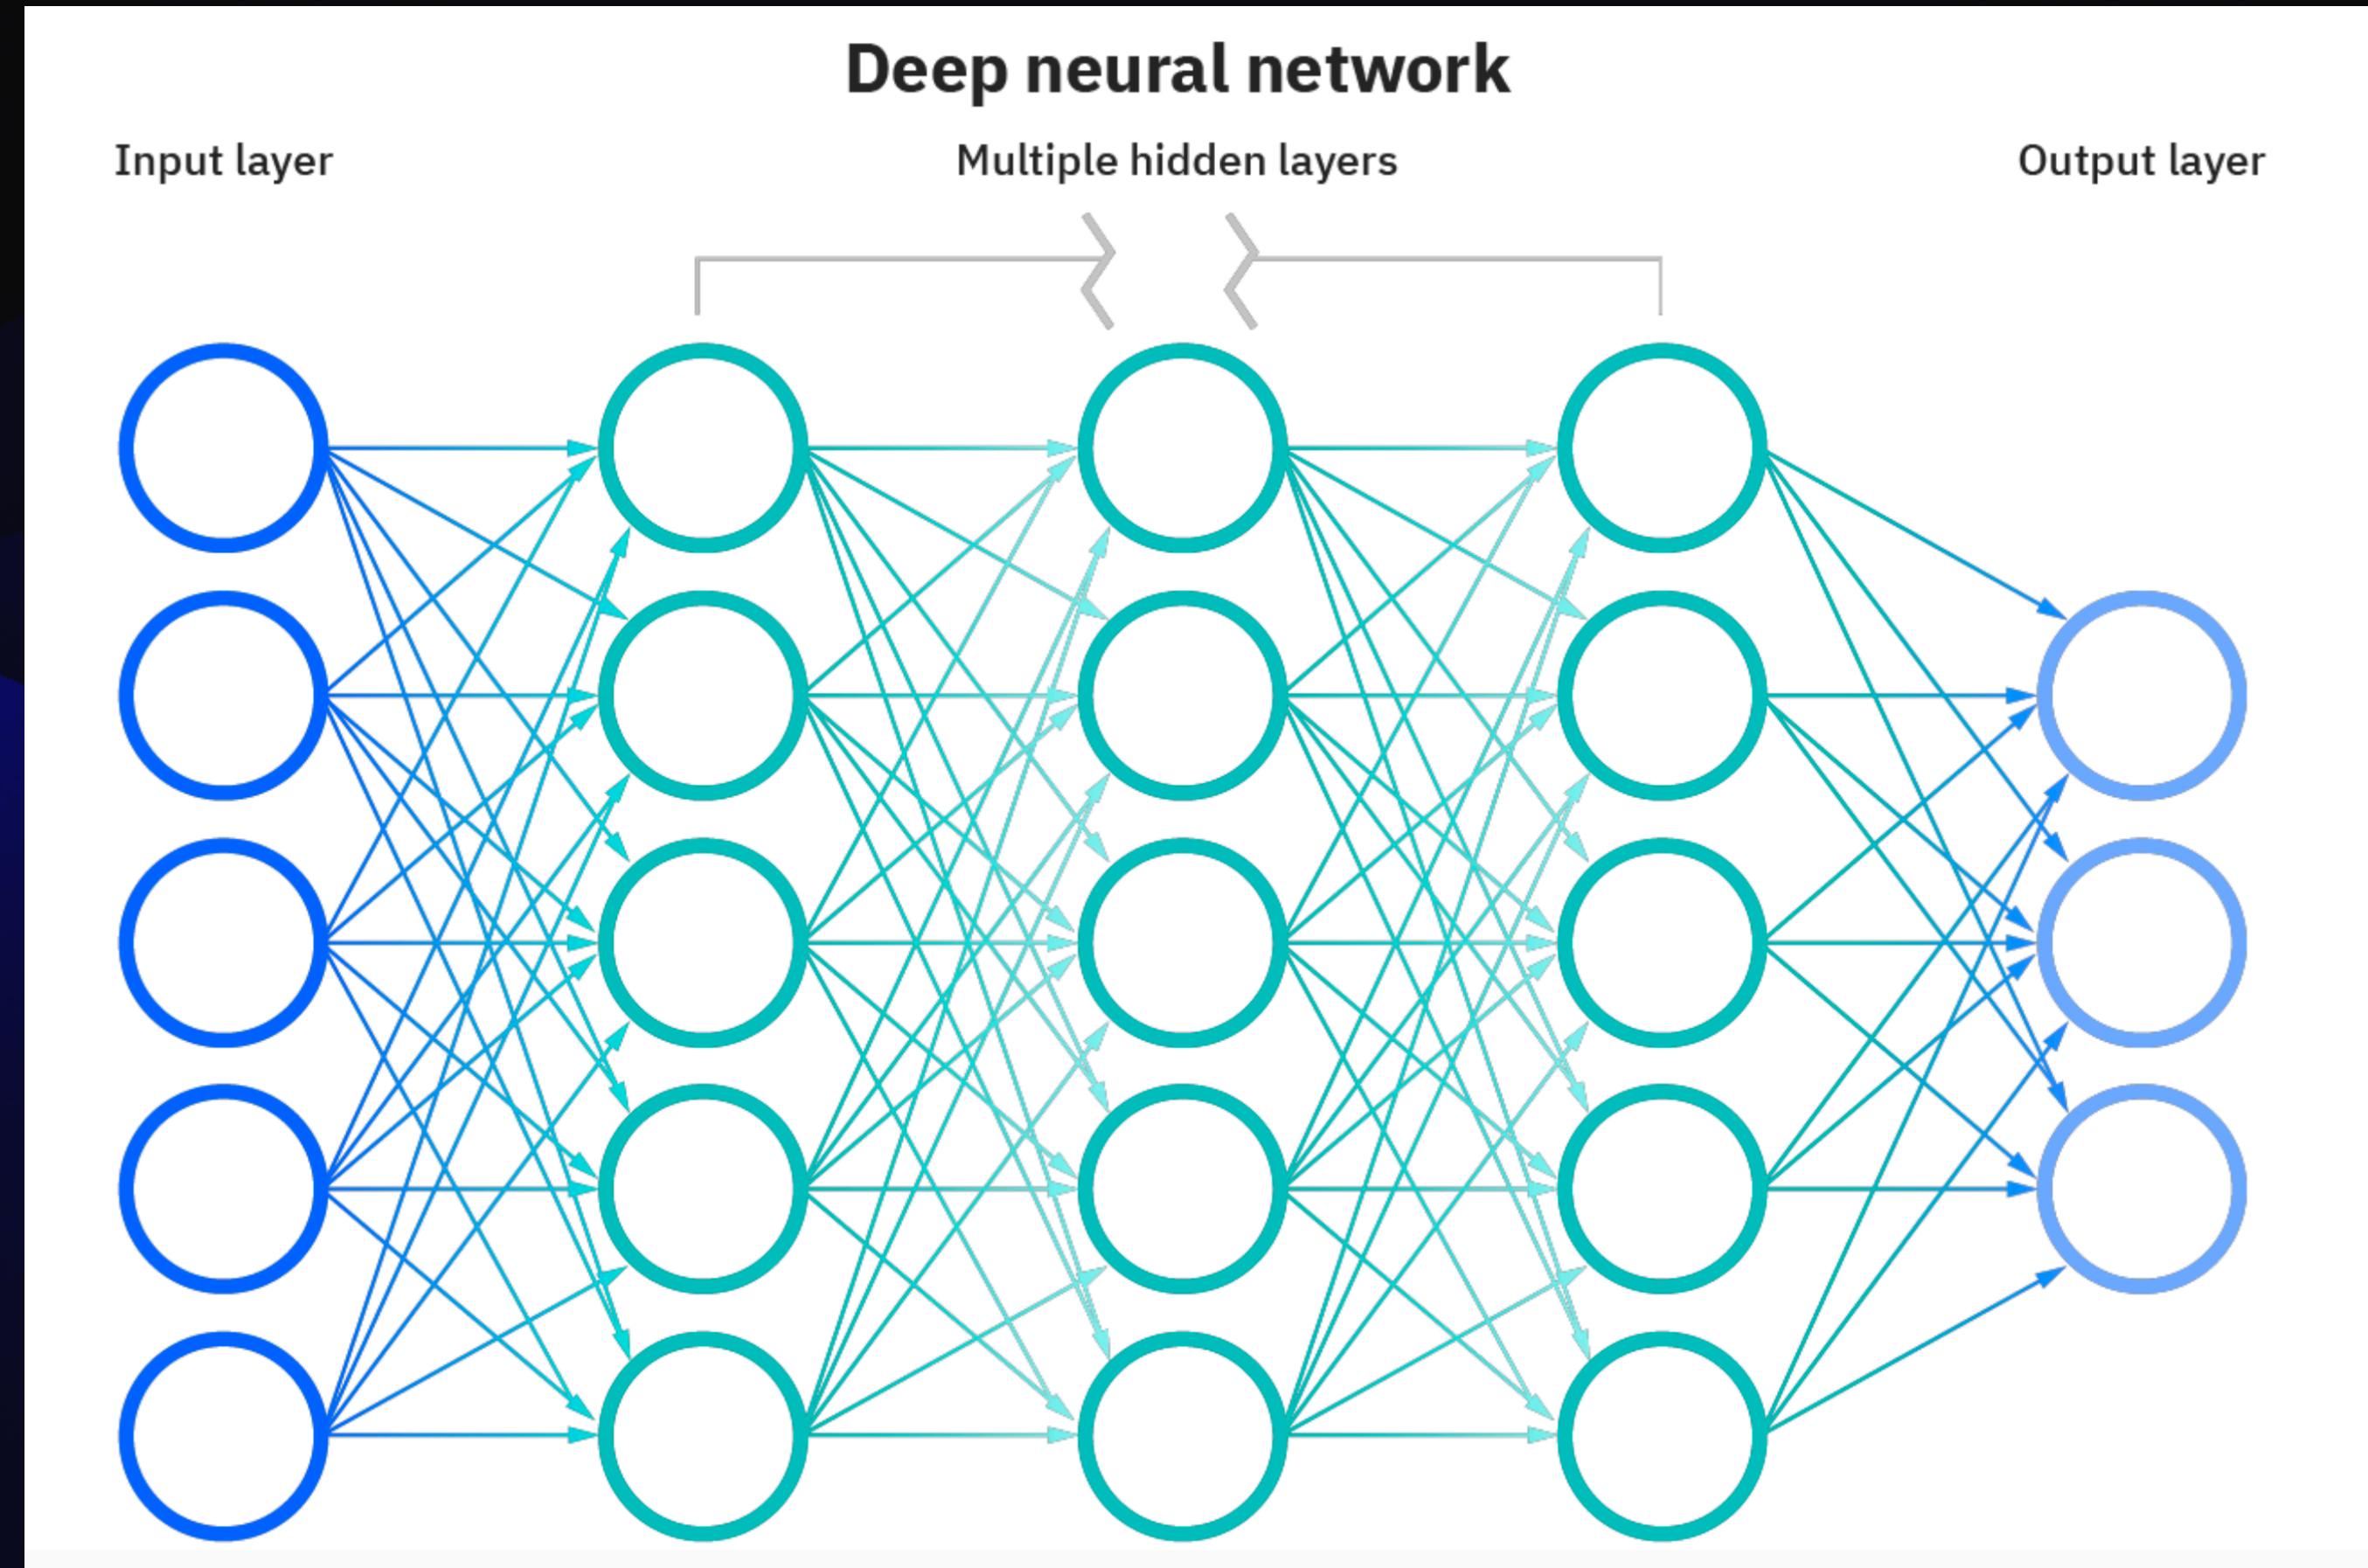
Pixel 0
Pixel 1
Pixel 2
⋮
Pixel N



Inputs: images, vector features,
words, sentences, etc

E.g. a flattened greyscale image

Pixel 0
Pixel 1
Pixel 2
⋮
Pixel N



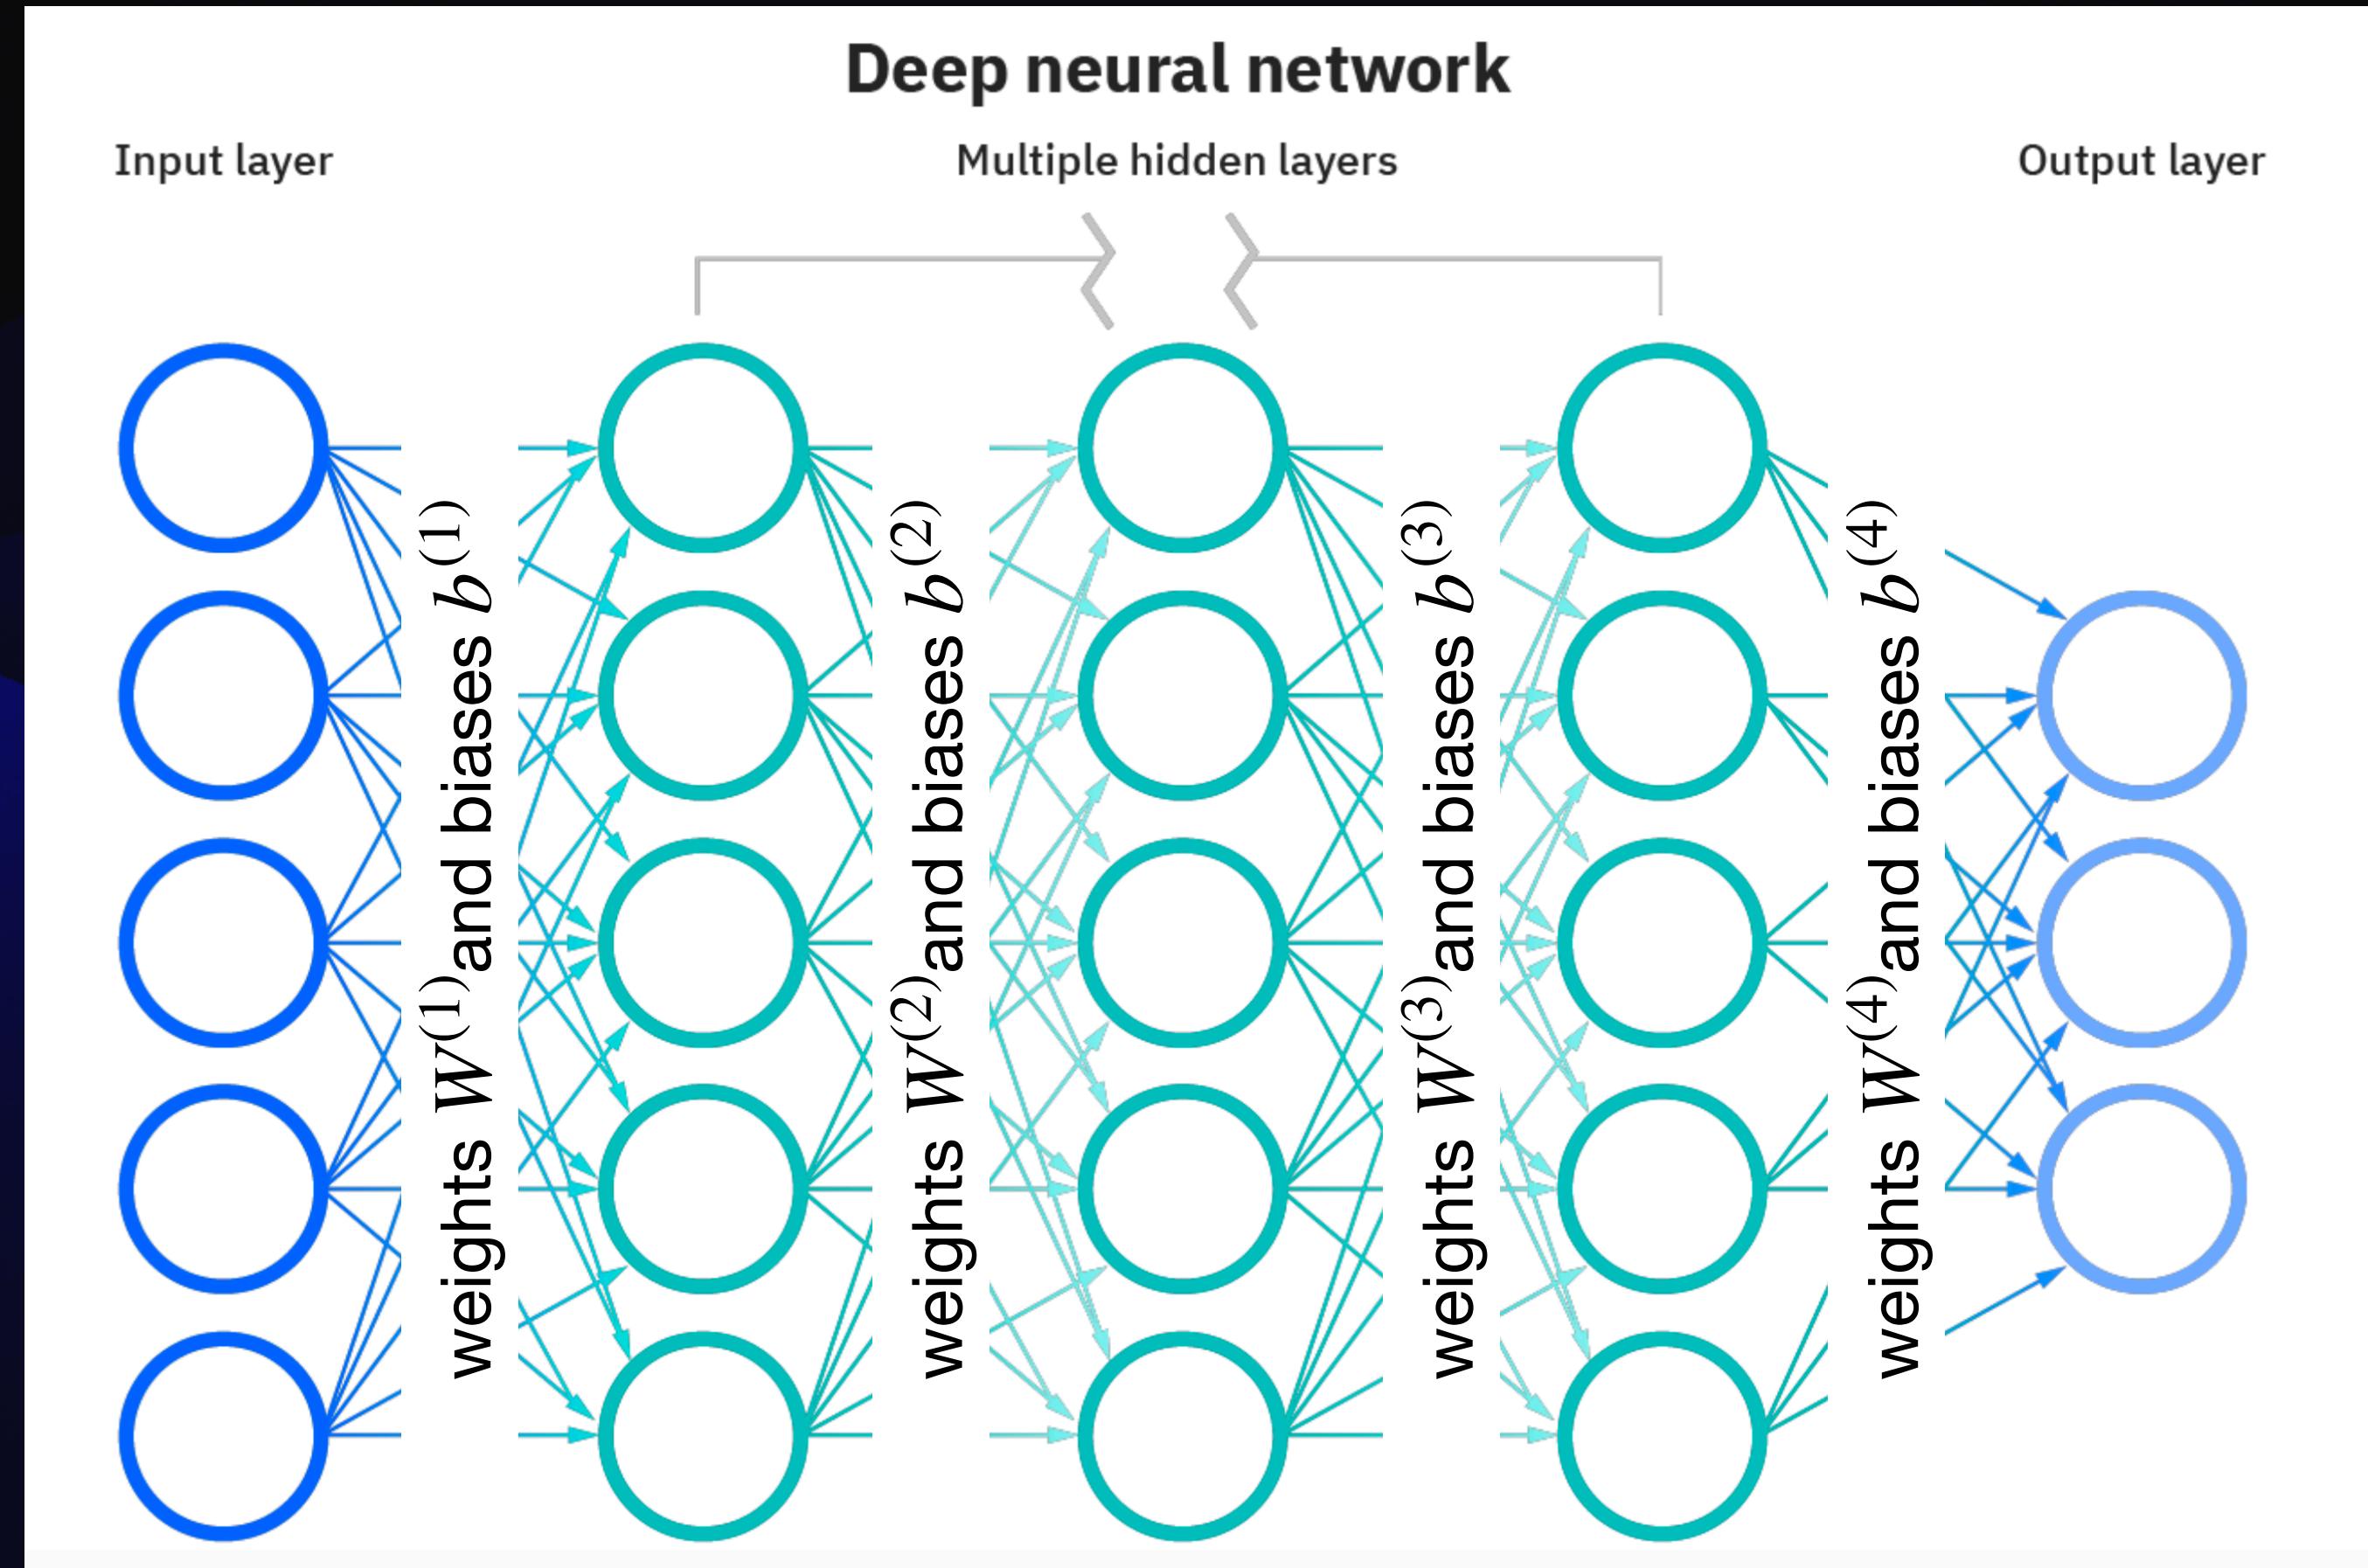
Inputs: images, vector features, words, sentences, etc

Hidden States: learnable parameters are used to transform the input sequentially

E.g. a flattened greyscale image

3 “hidden layers”
Updates of the form
 $x' = Wx + b$

Pixel 0
Pixel 1
Pixel 2
⋮
Pixel N

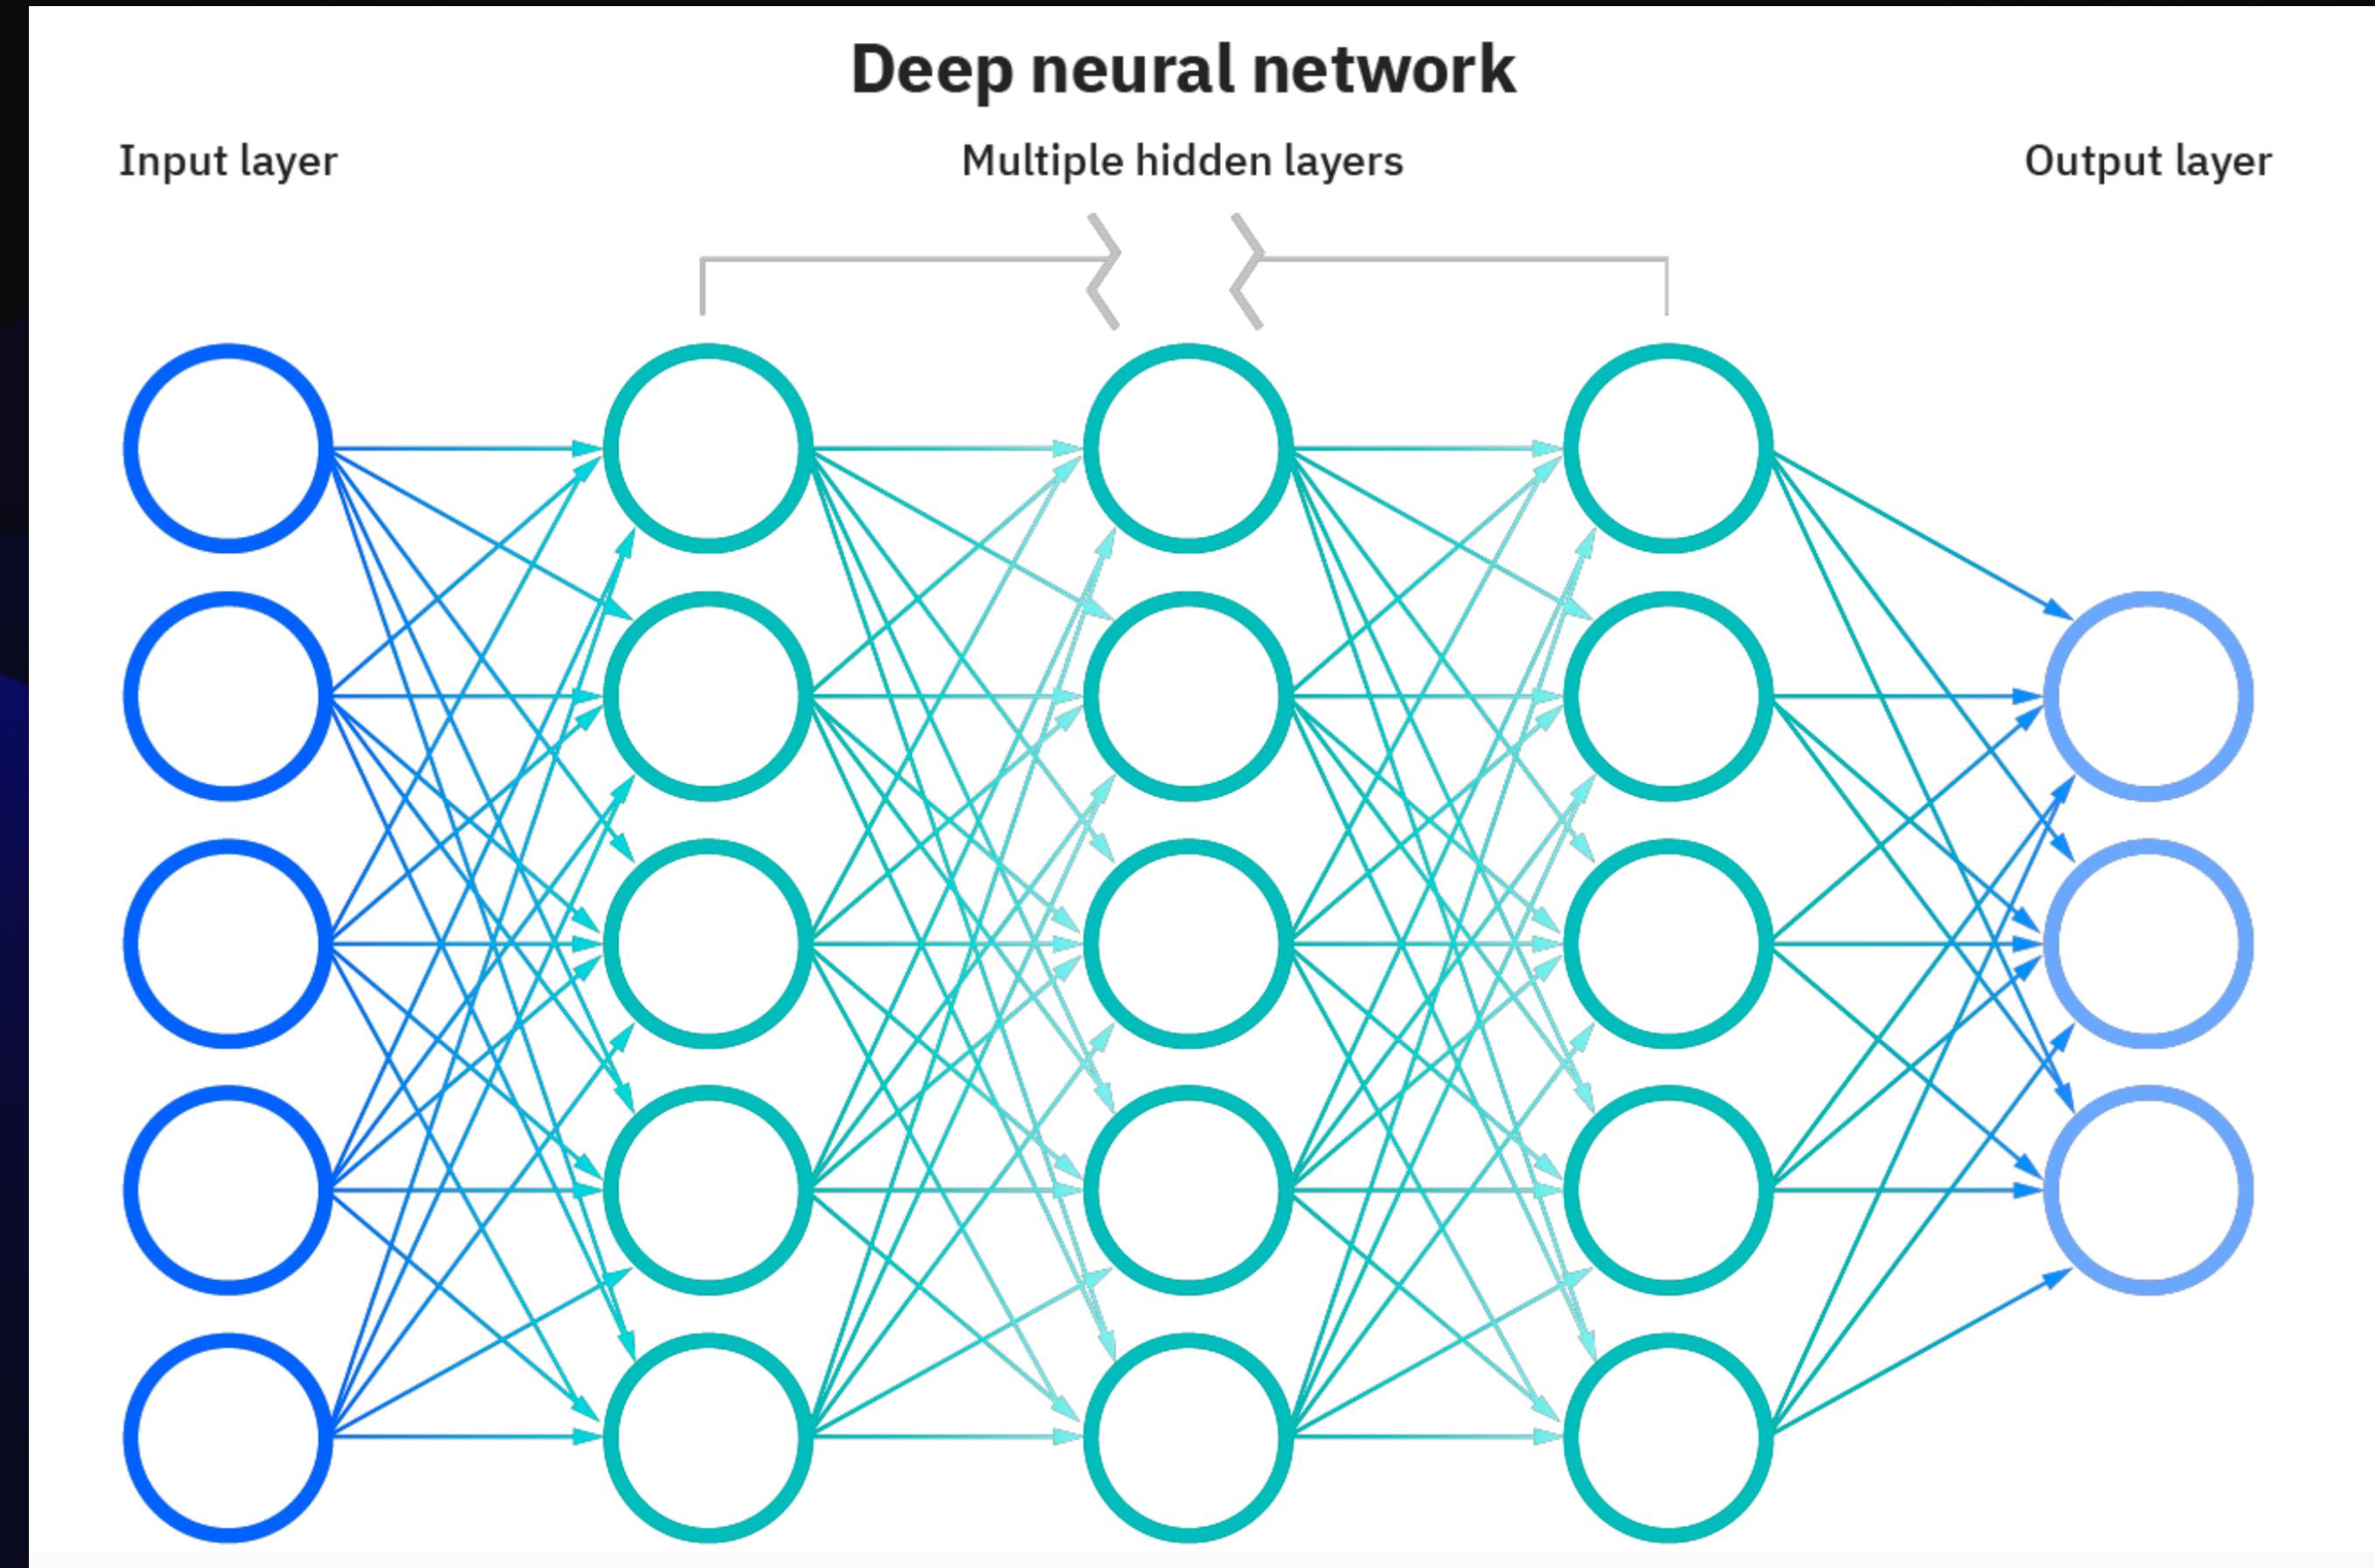


Inputs: images, vector features, words, sentences, etc

Hidden States: learnable parameters are used to transform the input sequentially

E.g. a flattened greyscale image

Pixel 0
Pixel 1
Pixel 2
⋮
Pixel N

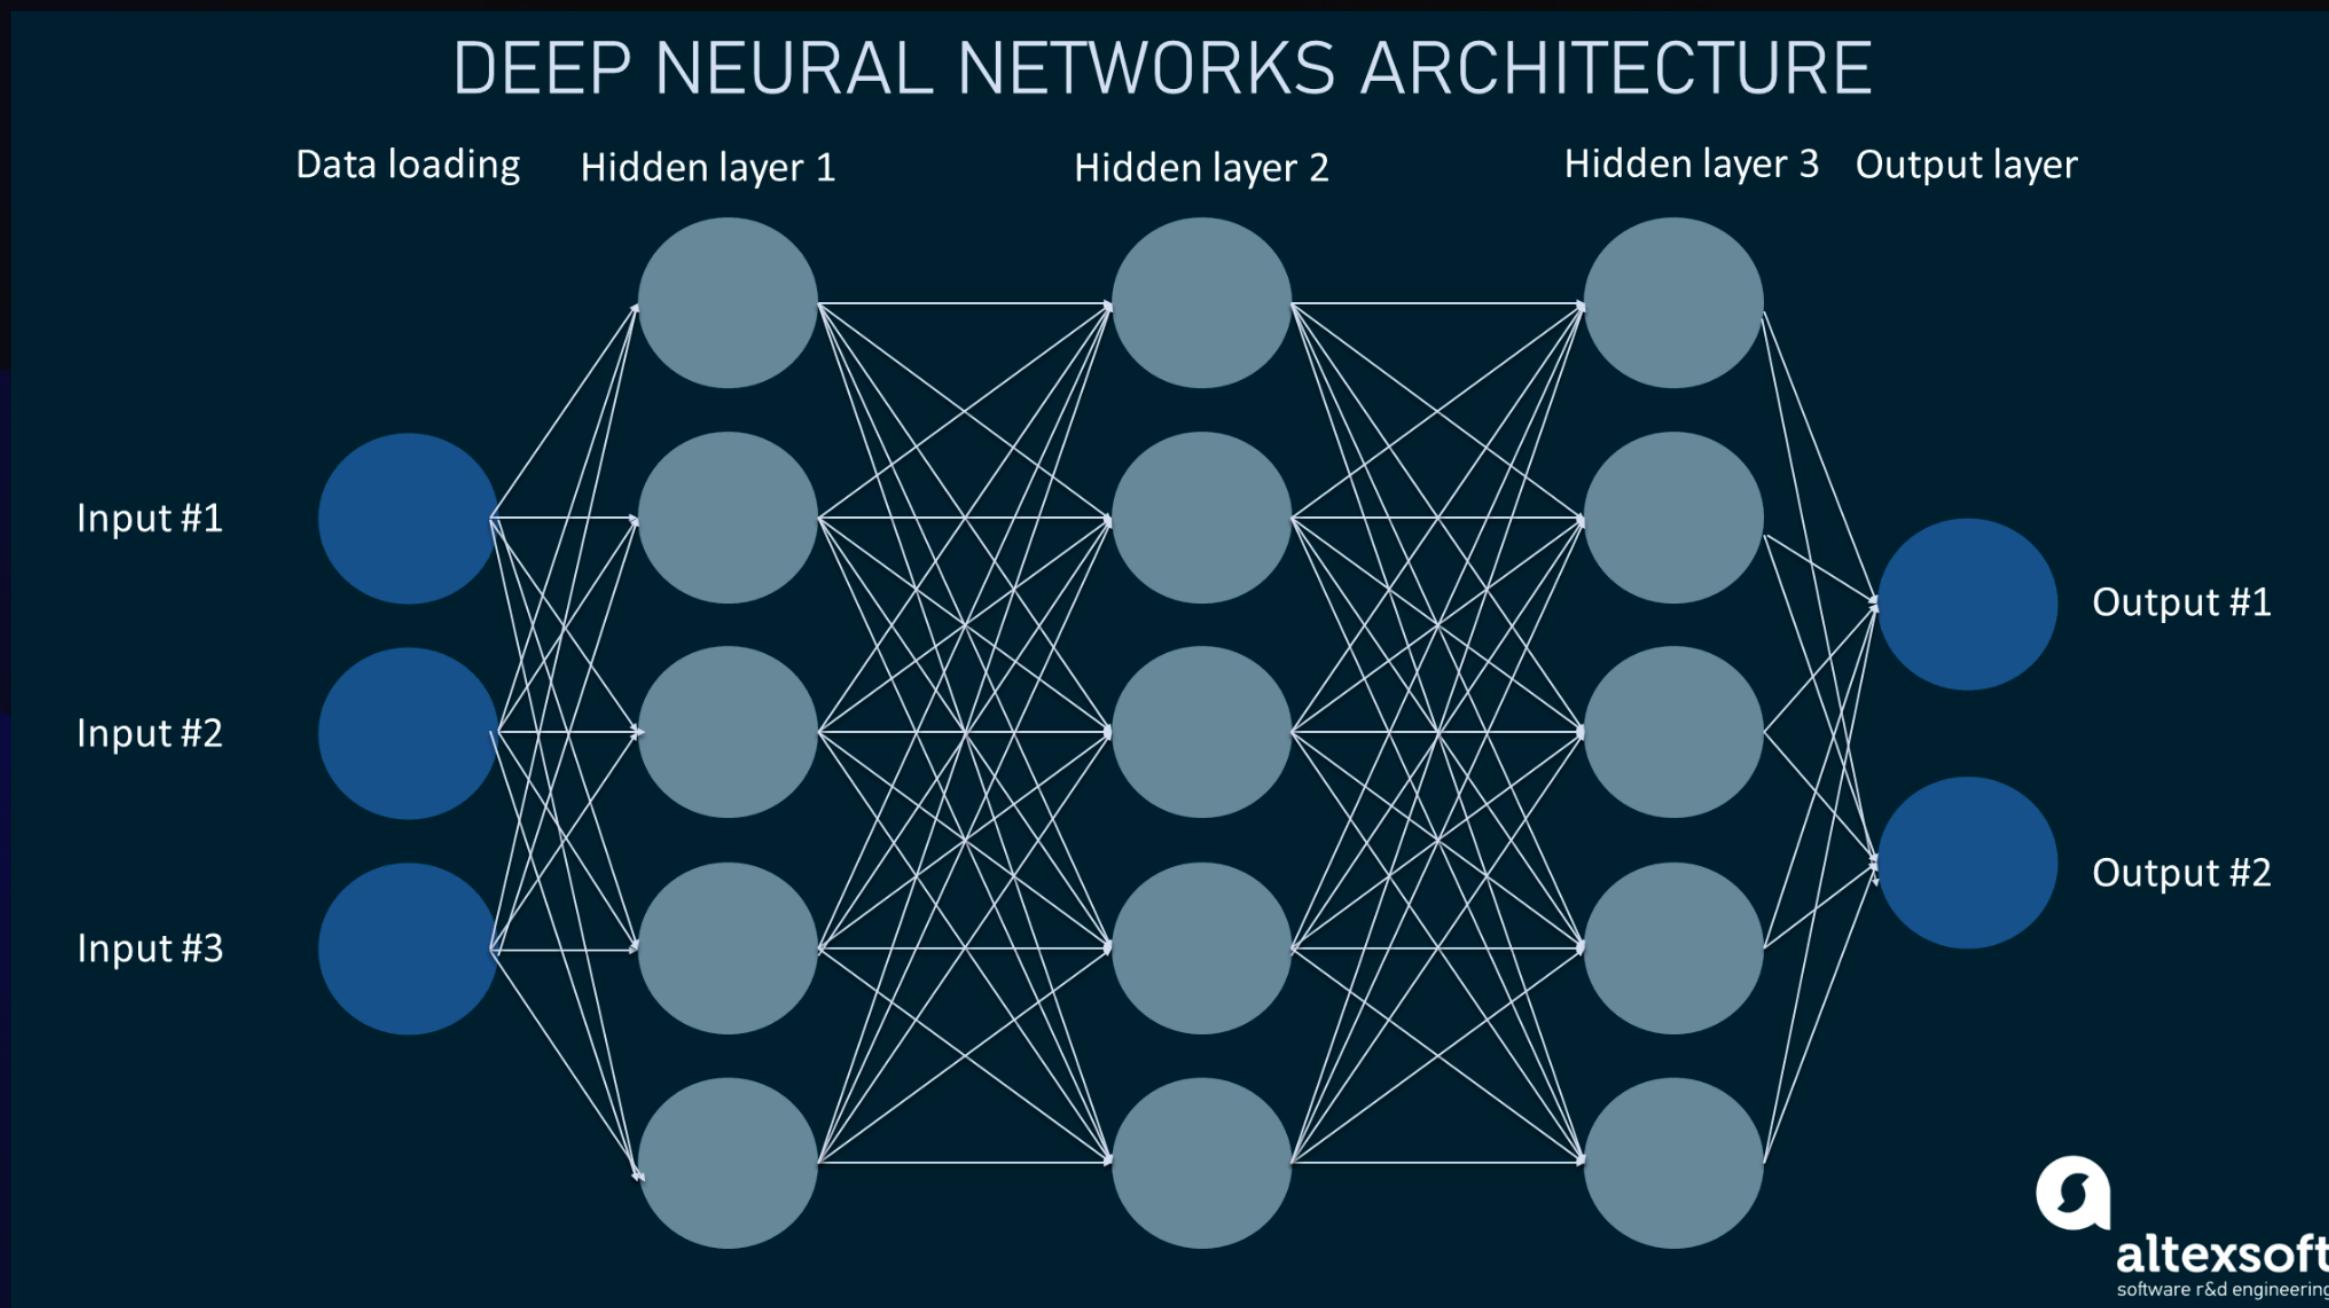


Inputs: images, vector features, words, sentences, etc

Hidden States: learnable parameters are used to transform the input sequentially

Output: model's predictions

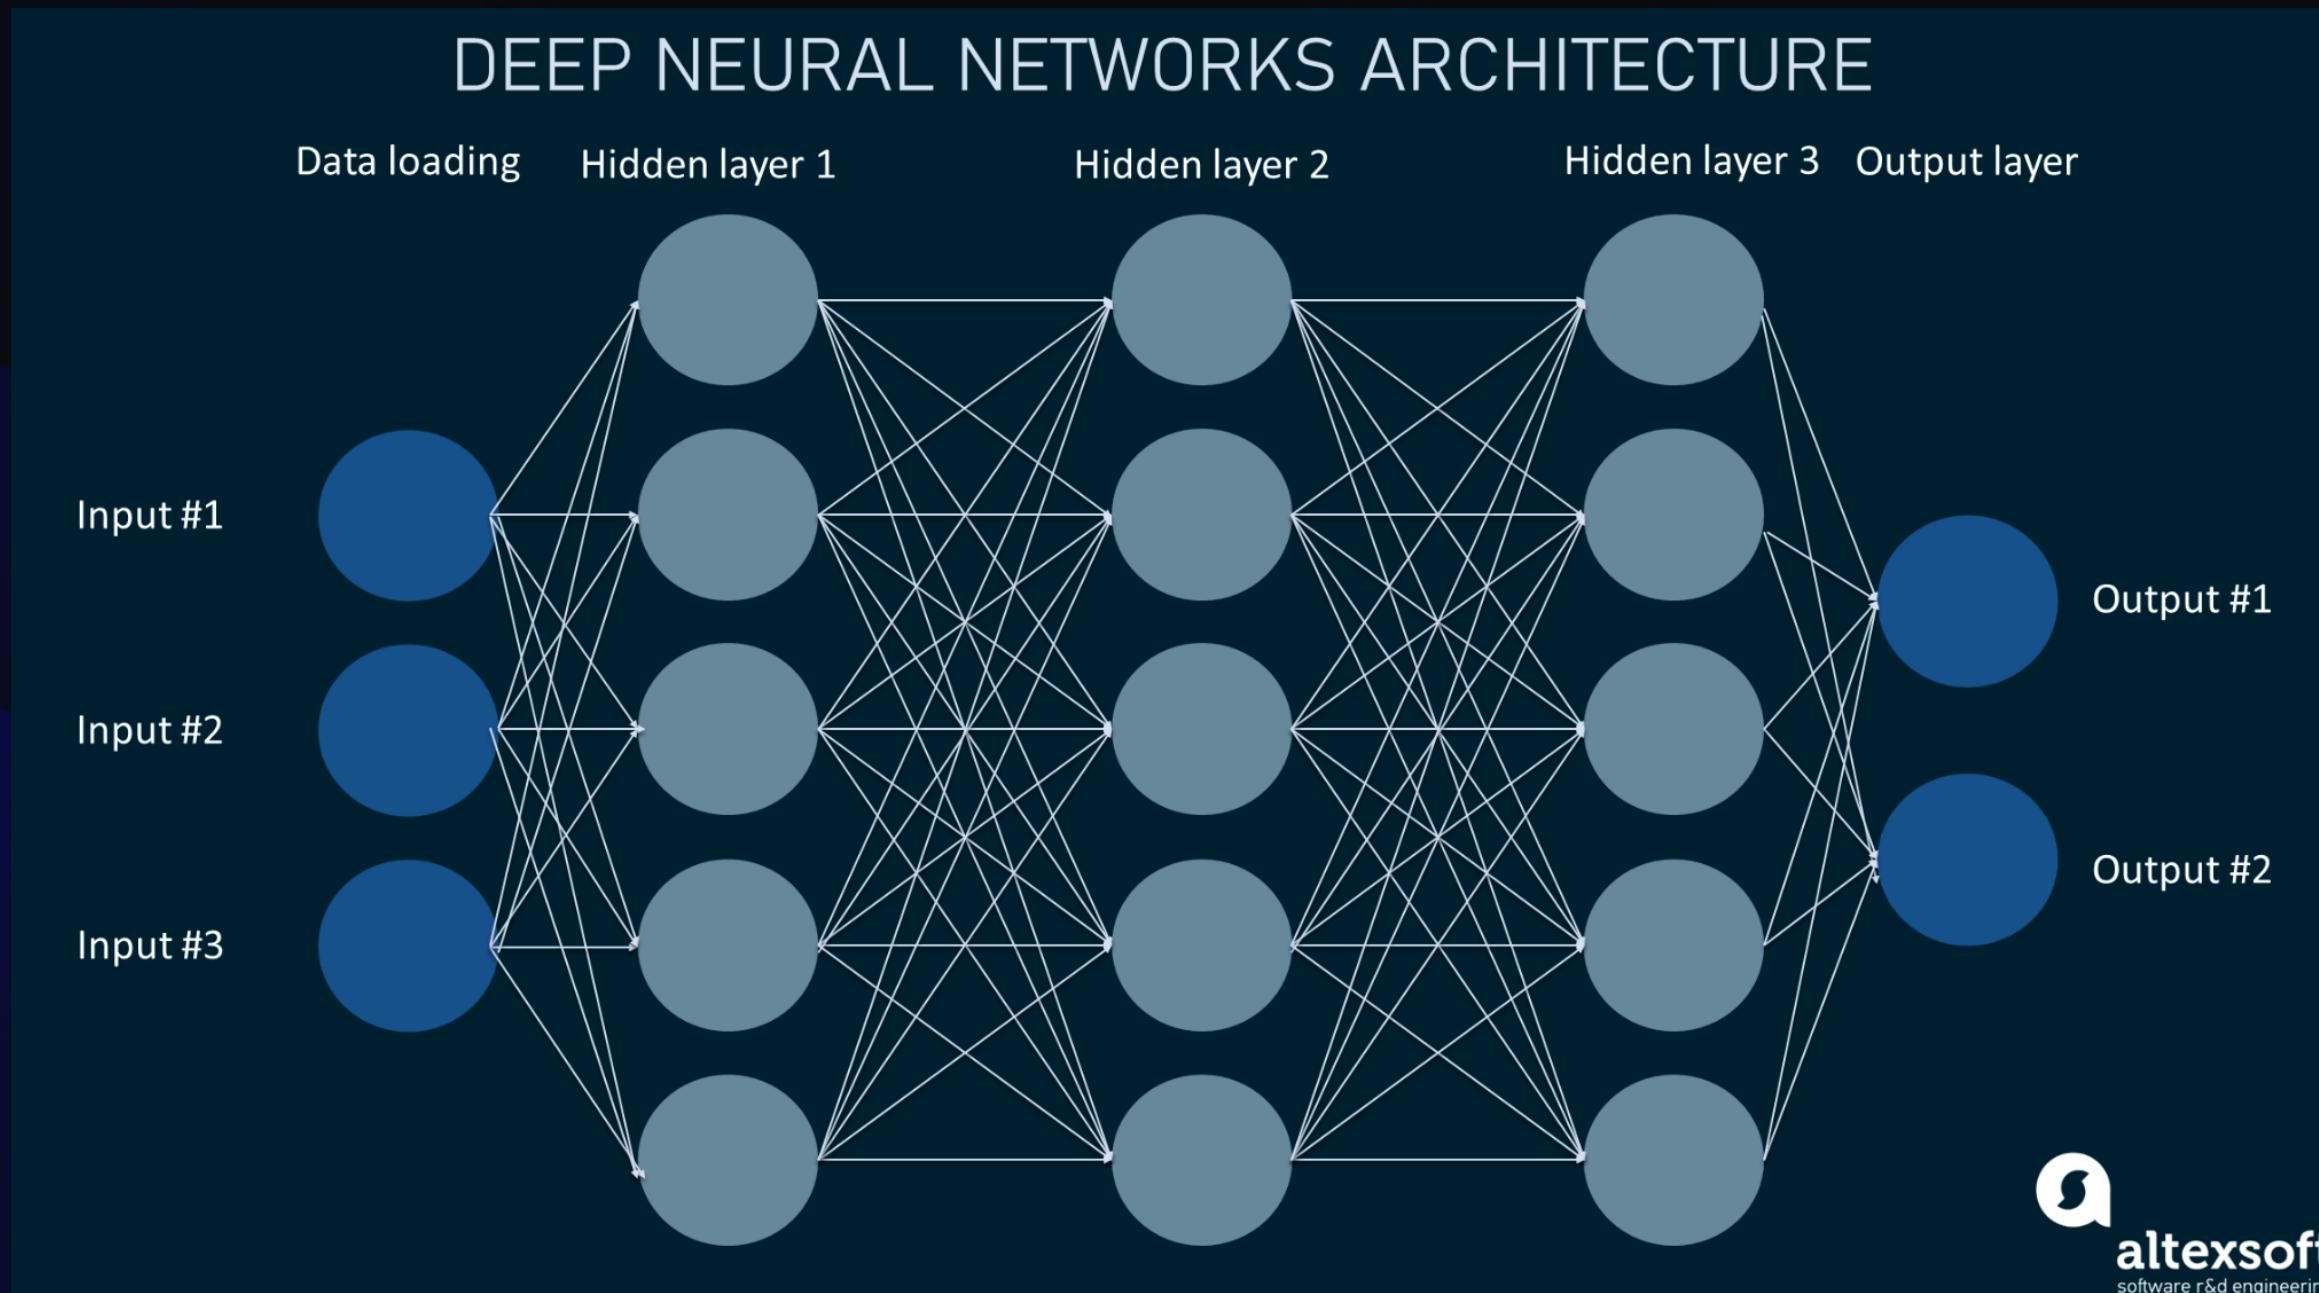
Make a prediction



Update the model

Objective / Loss Function describes how well the prediction matches a target

Make a prediction



Update the model

- Forward pass (inference):

$$z_i^{(\ell+1)} = \sum_{j=1}^{n_\ell} W_{ij}^{(\ell+1)} \sigma(z_j^{(\ell)}) + b_i^{(\ell+1)}$$

- Backward pass (training):

$$W_{ij}^{(\ell+1)} = W_{ij}^{(\ell)} - \gamma \frac{\partial L}{\partial W_{ij}} \Big|_{W_{ij}^{(\ell)}}$$

$$b_i^{(\ell+1)} = b_i^{(\ell)} - \gamma \frac{\partial L}{\partial b_i} \Big|_{b_i^{(\ell)}} \\ (\text{gradient descent})$$

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

```

$$z_i^{(\ell+1)} = \sum_{j=1}^{n_\ell} W_{ij}^{(\ell+1)} \sigma(z_j^{(\ell)}) + b_i^{(\ell+1)}$$

```

def train_one_epoch(epoch_index, tb_writer):
    running_loss = 0.
    last_loss = 0.

    # Here, we use enumerate(training_loader) instead of
    # iter(training_loader) so that we can track the batch
    # index and do some intra-epoch reporting
    for i, data in enumerate(training_loader):
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()

        # Gather data and report
        running_loss += loss.item()
        if i % 1000 == 999:
            last_loss = running_loss / 1000 # loss per batch
            print(' batch {} loss: {}'.format(i + 1, last_loss))
            tb_x = epoch_index * len(training_loader) + i + 1
            tb_writer.add_scalar('Loss/train', last_loss, tb_x)
            running_loss = 0.

    return last_loss

```

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

```

$$W_{ij}^{(\ell+1)} = W_{ij}^{(\ell)} - \gamma \frac{\partial L}{\partial W_{ij}} \Big|_{W_{ij}^{(\ell)}} \quad b_i^{(\ell+1)} = b_i^{(\ell)} - \gamma \frac{\partial L}{\partial b_i} \Big|_{b_i^{(\ell)}}$$

```

def train_one_epoch(epoch_index, tb_writer):
    running_loss = 0.
    last_loss = 0.

    # Here, we use enumerate(training_loader) instead of
    # iter(training_loader) so that we can track the batch
    # index and do some intra-epoch reporting
    for i, data in enumerate(training_loader):
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()

        # Gather data and report
        running_loss += loss.item()
        if i % 1000 == 999:
            last_loss = running_loss / 1000 # loss per batch
            print(' batch {} loss: {}'.format(i + 1, last_loss))
            tb_x = epoch_index * len(training_loader) + i + 1
            tb_writer.add_scalar('Loss/train', last_loss, tb_x)
            running_loss = 0.

    return last_loss

```

DNN Notebook

neural-network-zoo_DNNs.ipynb



<https://tinyurl.com/2fx8smzk>

<https://github.com/GageDeZoort/neural-network-zoo/tree/main>

Convolutional Neural Networks

Computer Vision

- **Classification:** predict one image label
- **Object detection:** predict bounding boxes for objects in image
- **Instance segmentation:** predict pixel boundaries for objects in images
- **Semantic segmentation:** predict class labels per pixel of objects in image

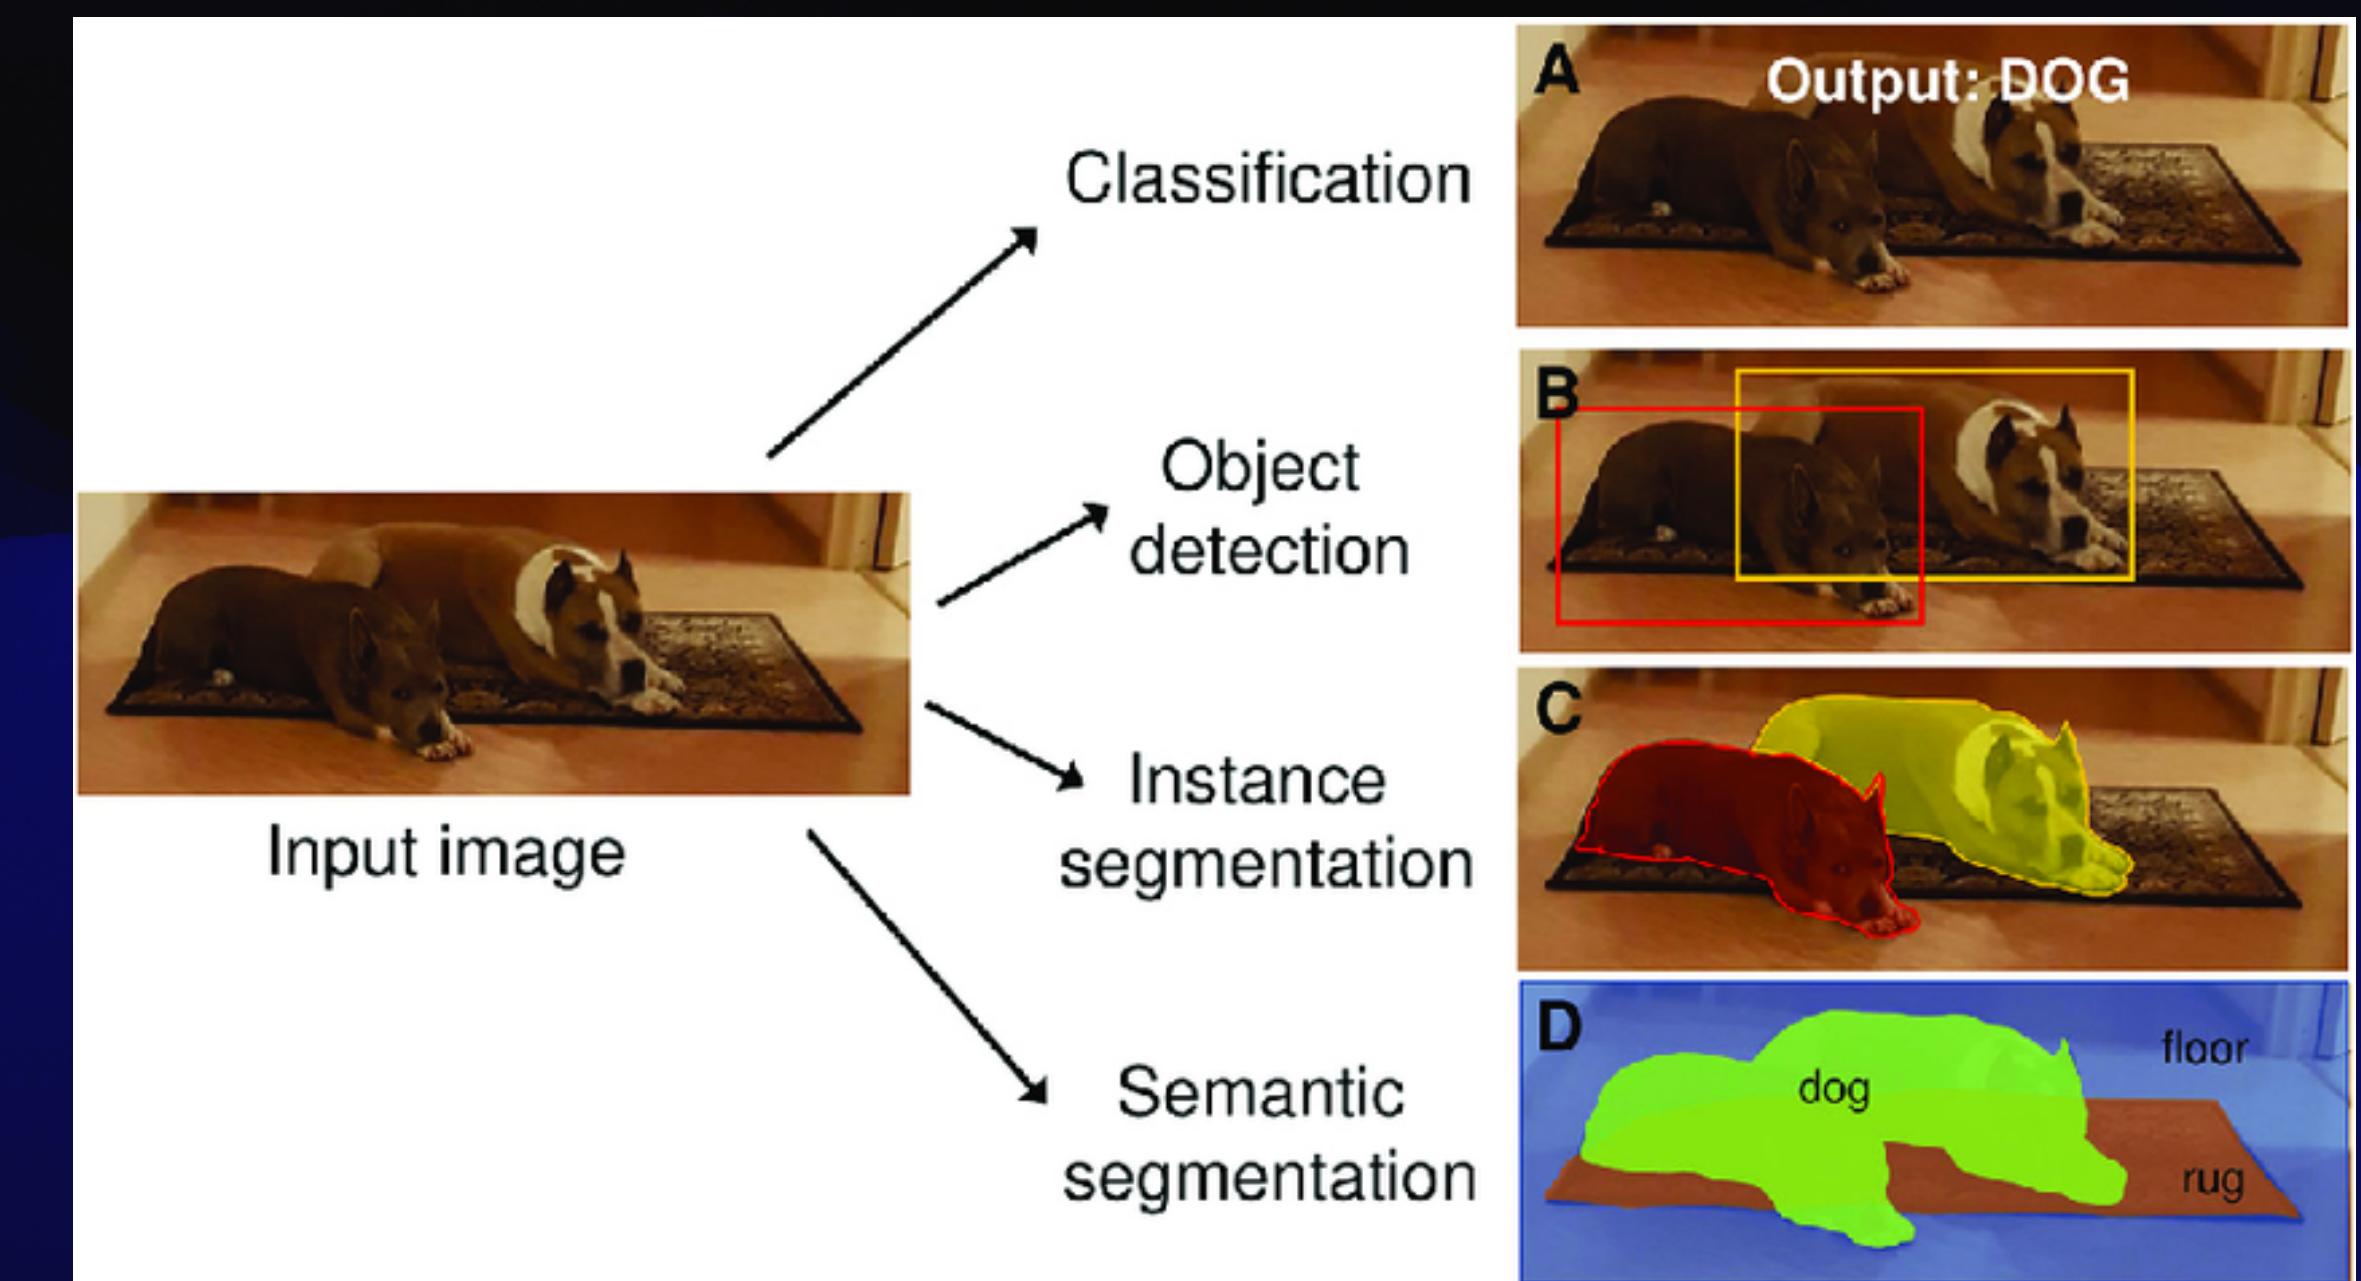
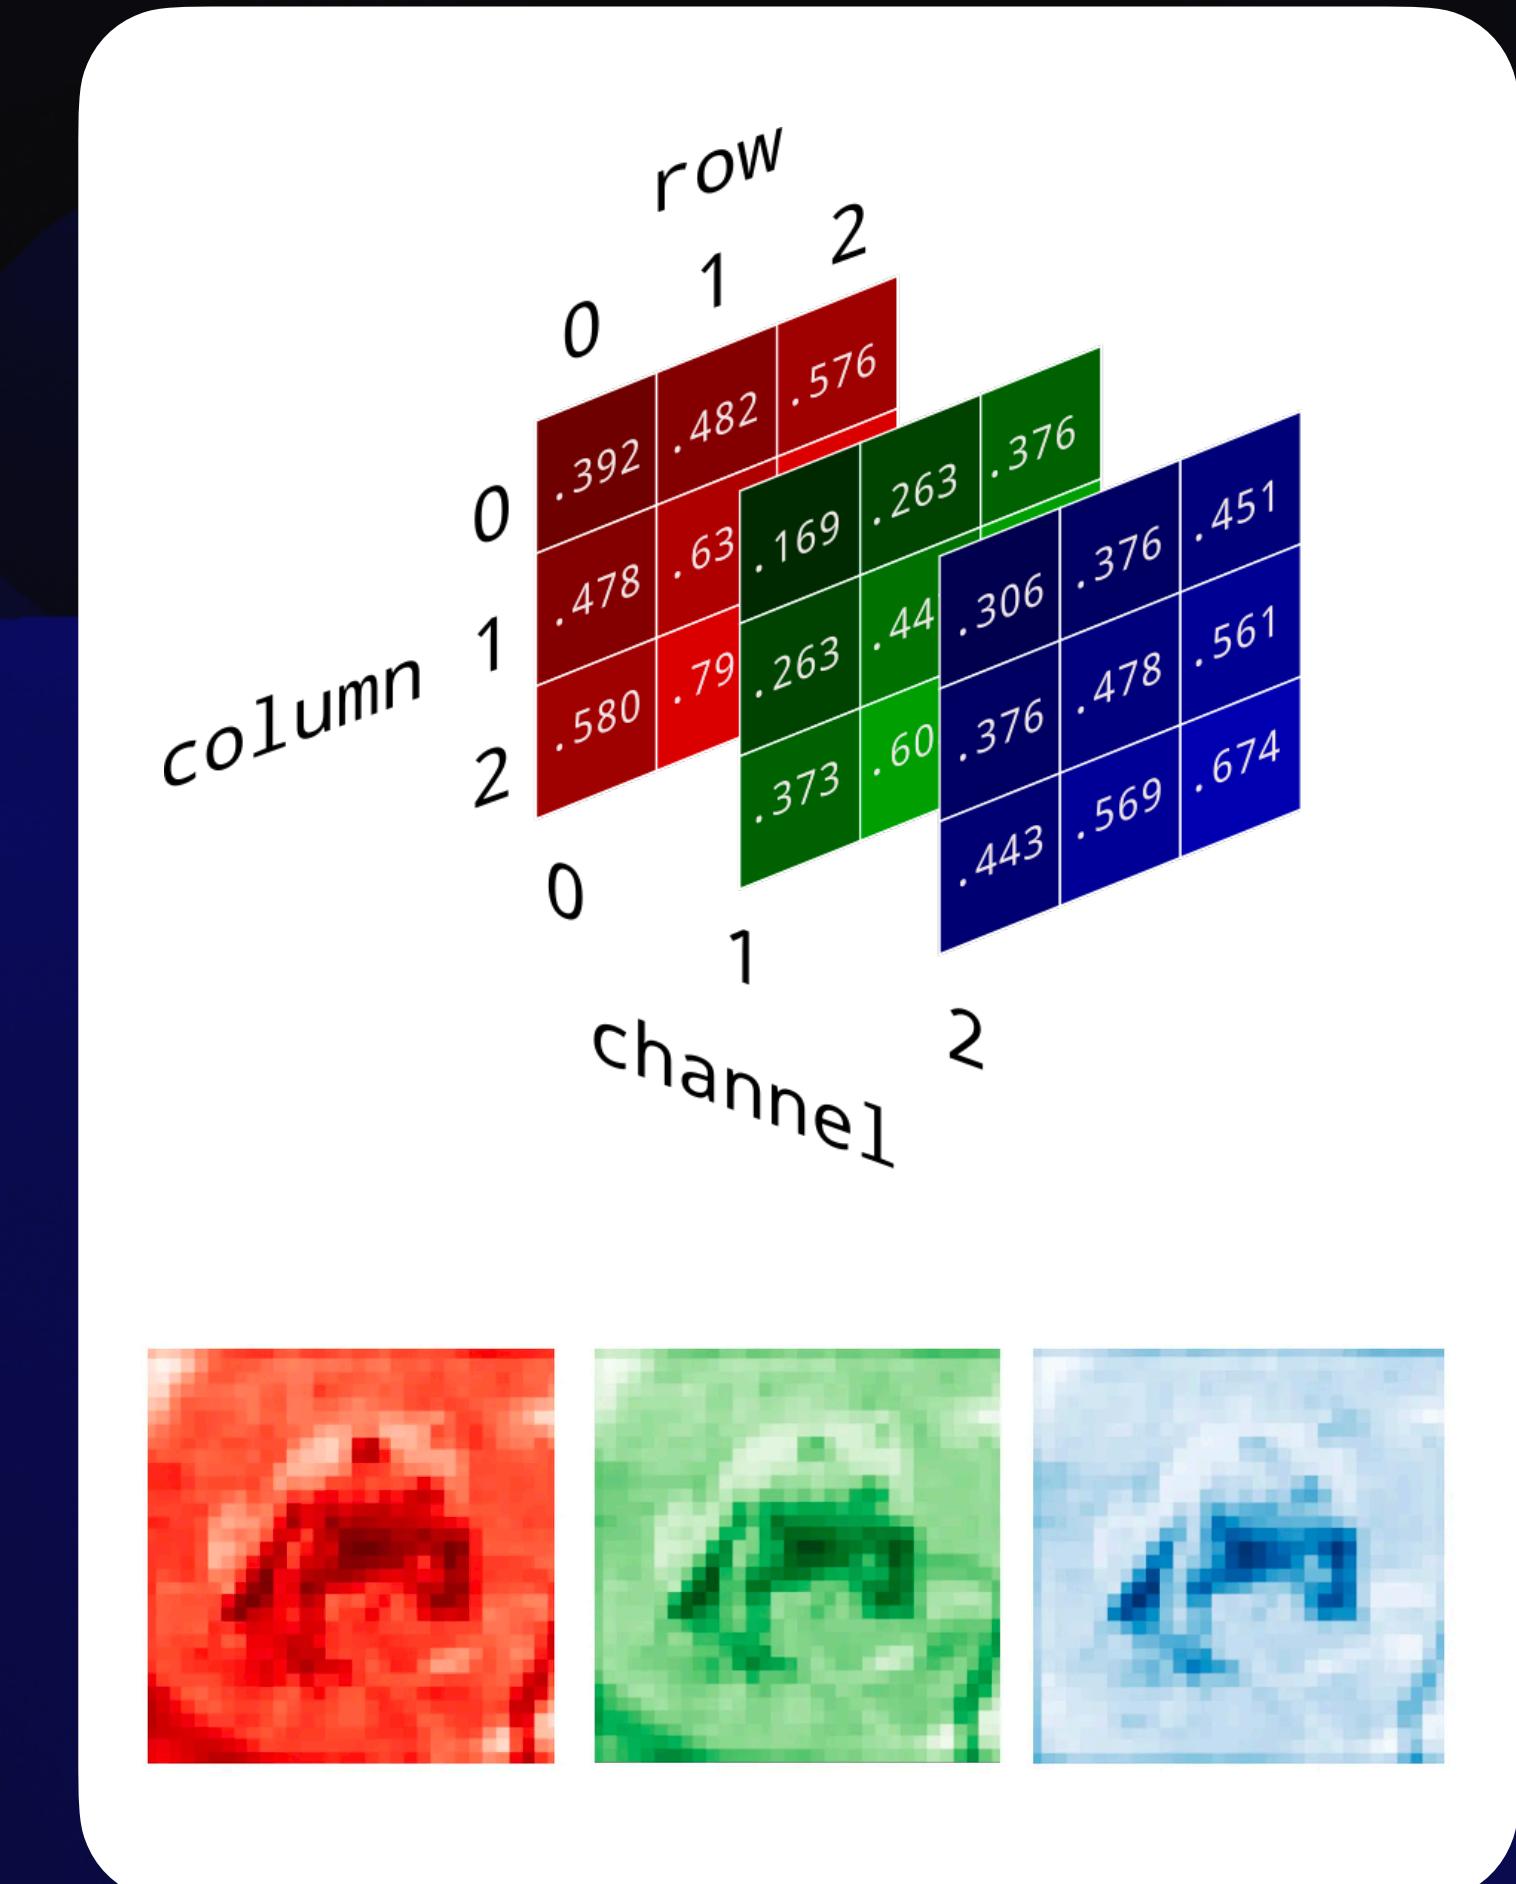


Image Classification

Problem Description

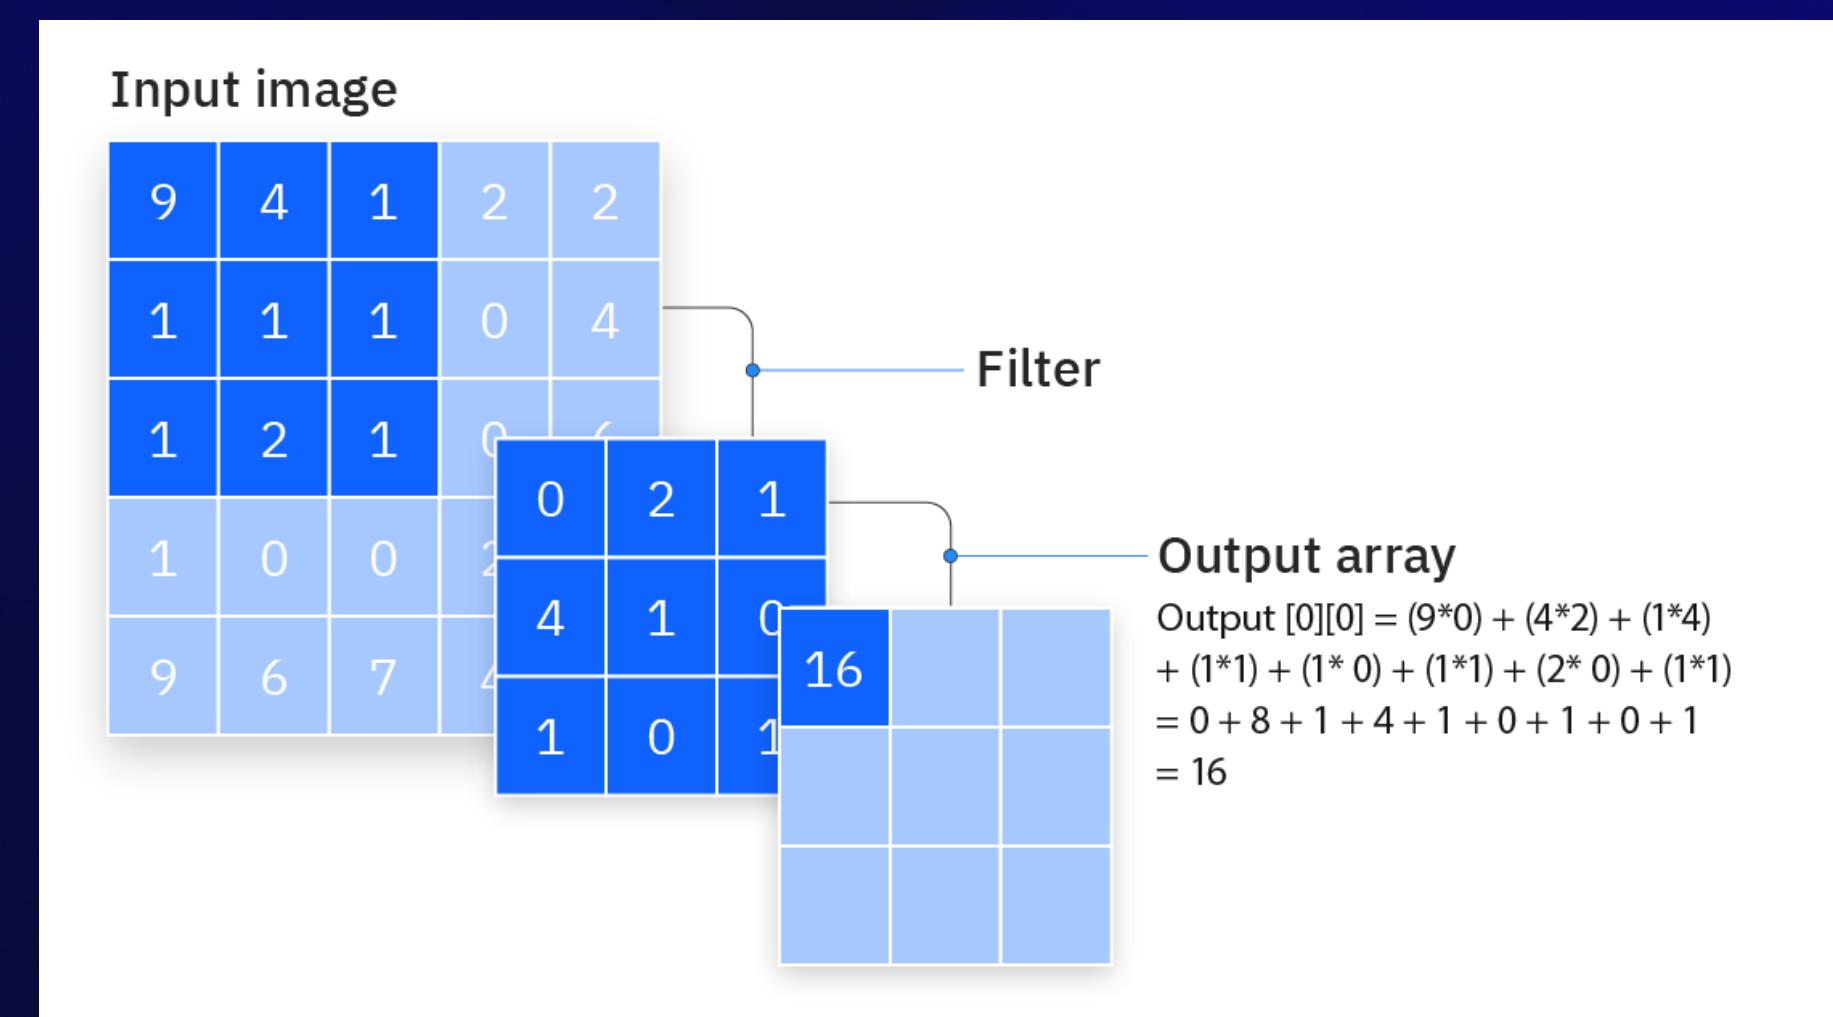
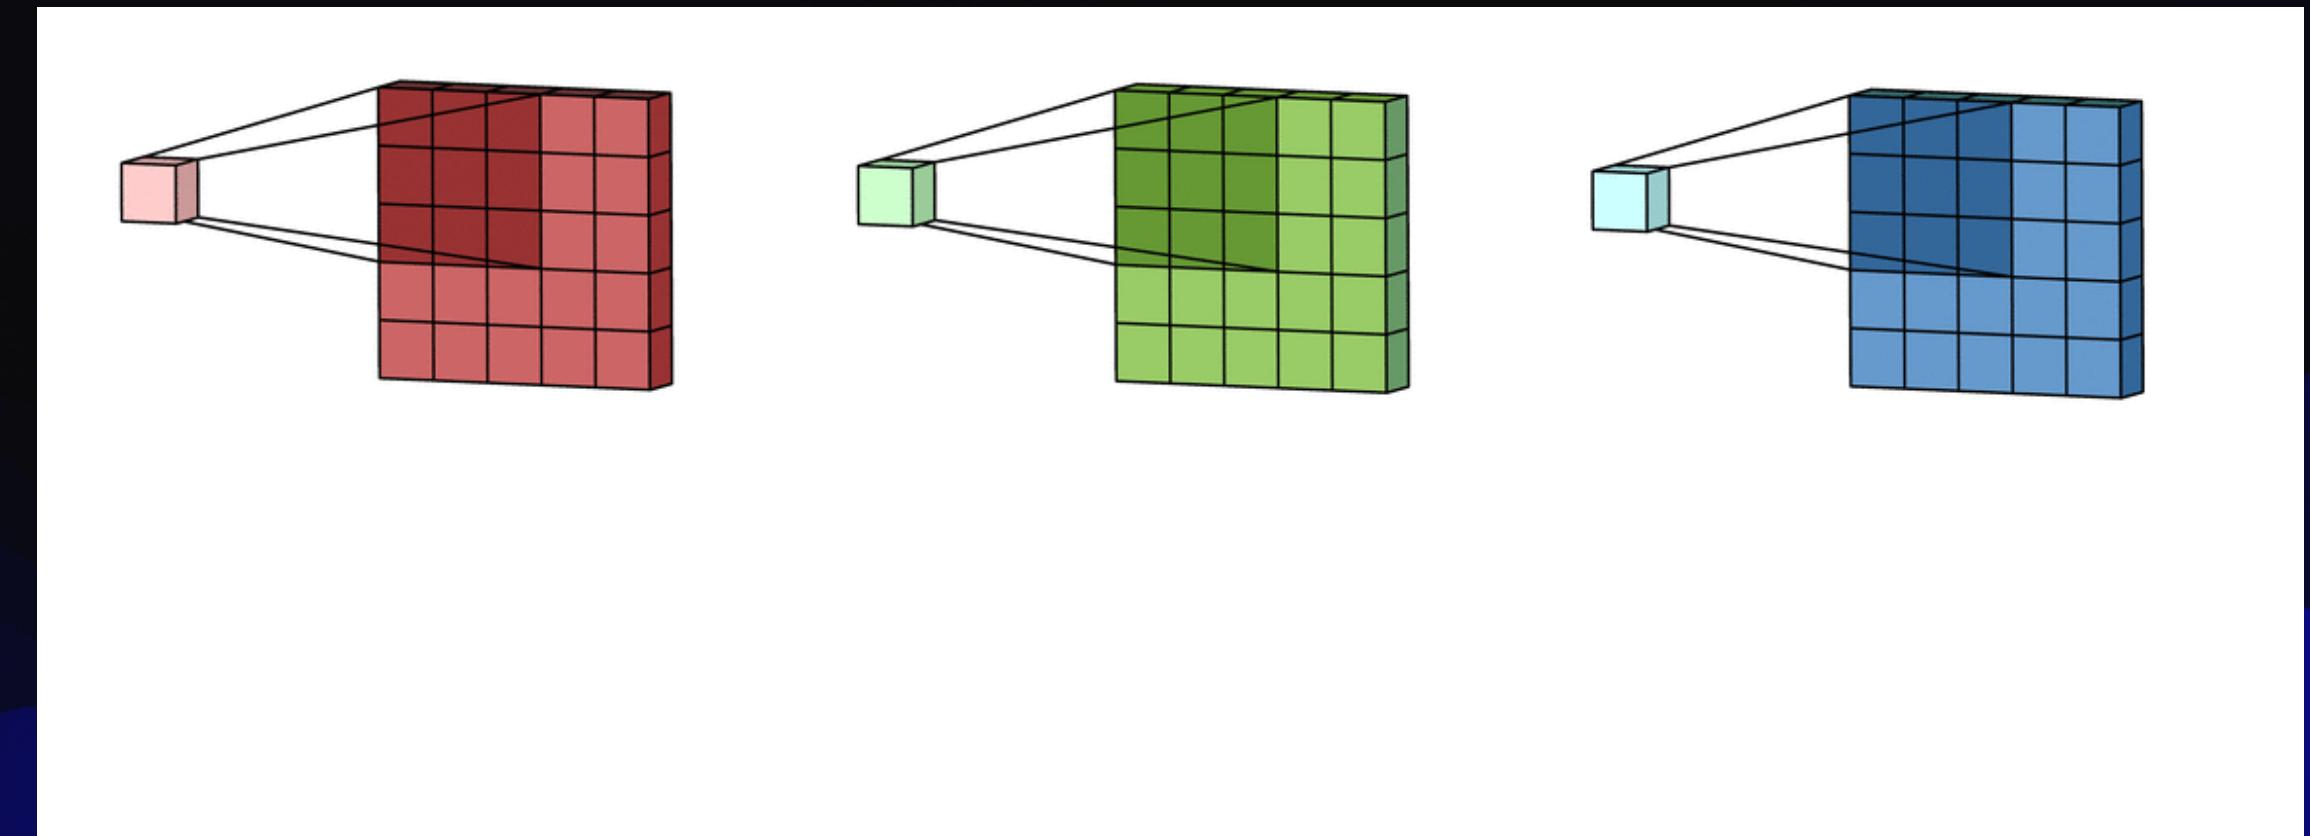
- Image → data on a grid
 - For square images, inputs are $I \in \mathbb{R}^{n_{\text{pixels}} \times n_{\text{pixels}} \times n_{\text{features}}}$
 - Example: 3x3 RGB frog images
- CNN maps image to probability predictions
 - Example: CNN(image) = P(frog)



Convolutional Neural Networks

Convolutional Layers

- Convolutional layers act like fully-connected layers in DNNs, but have several key differences:
 - The weights in the convolutional layer (*filter*) are applied to small patches of the image
 - Parameter sharing - the *same weights* are applied equally to every region of the input image!
 - Tuning the convolutional weights allows the network to look for different features - *feature extraction!*

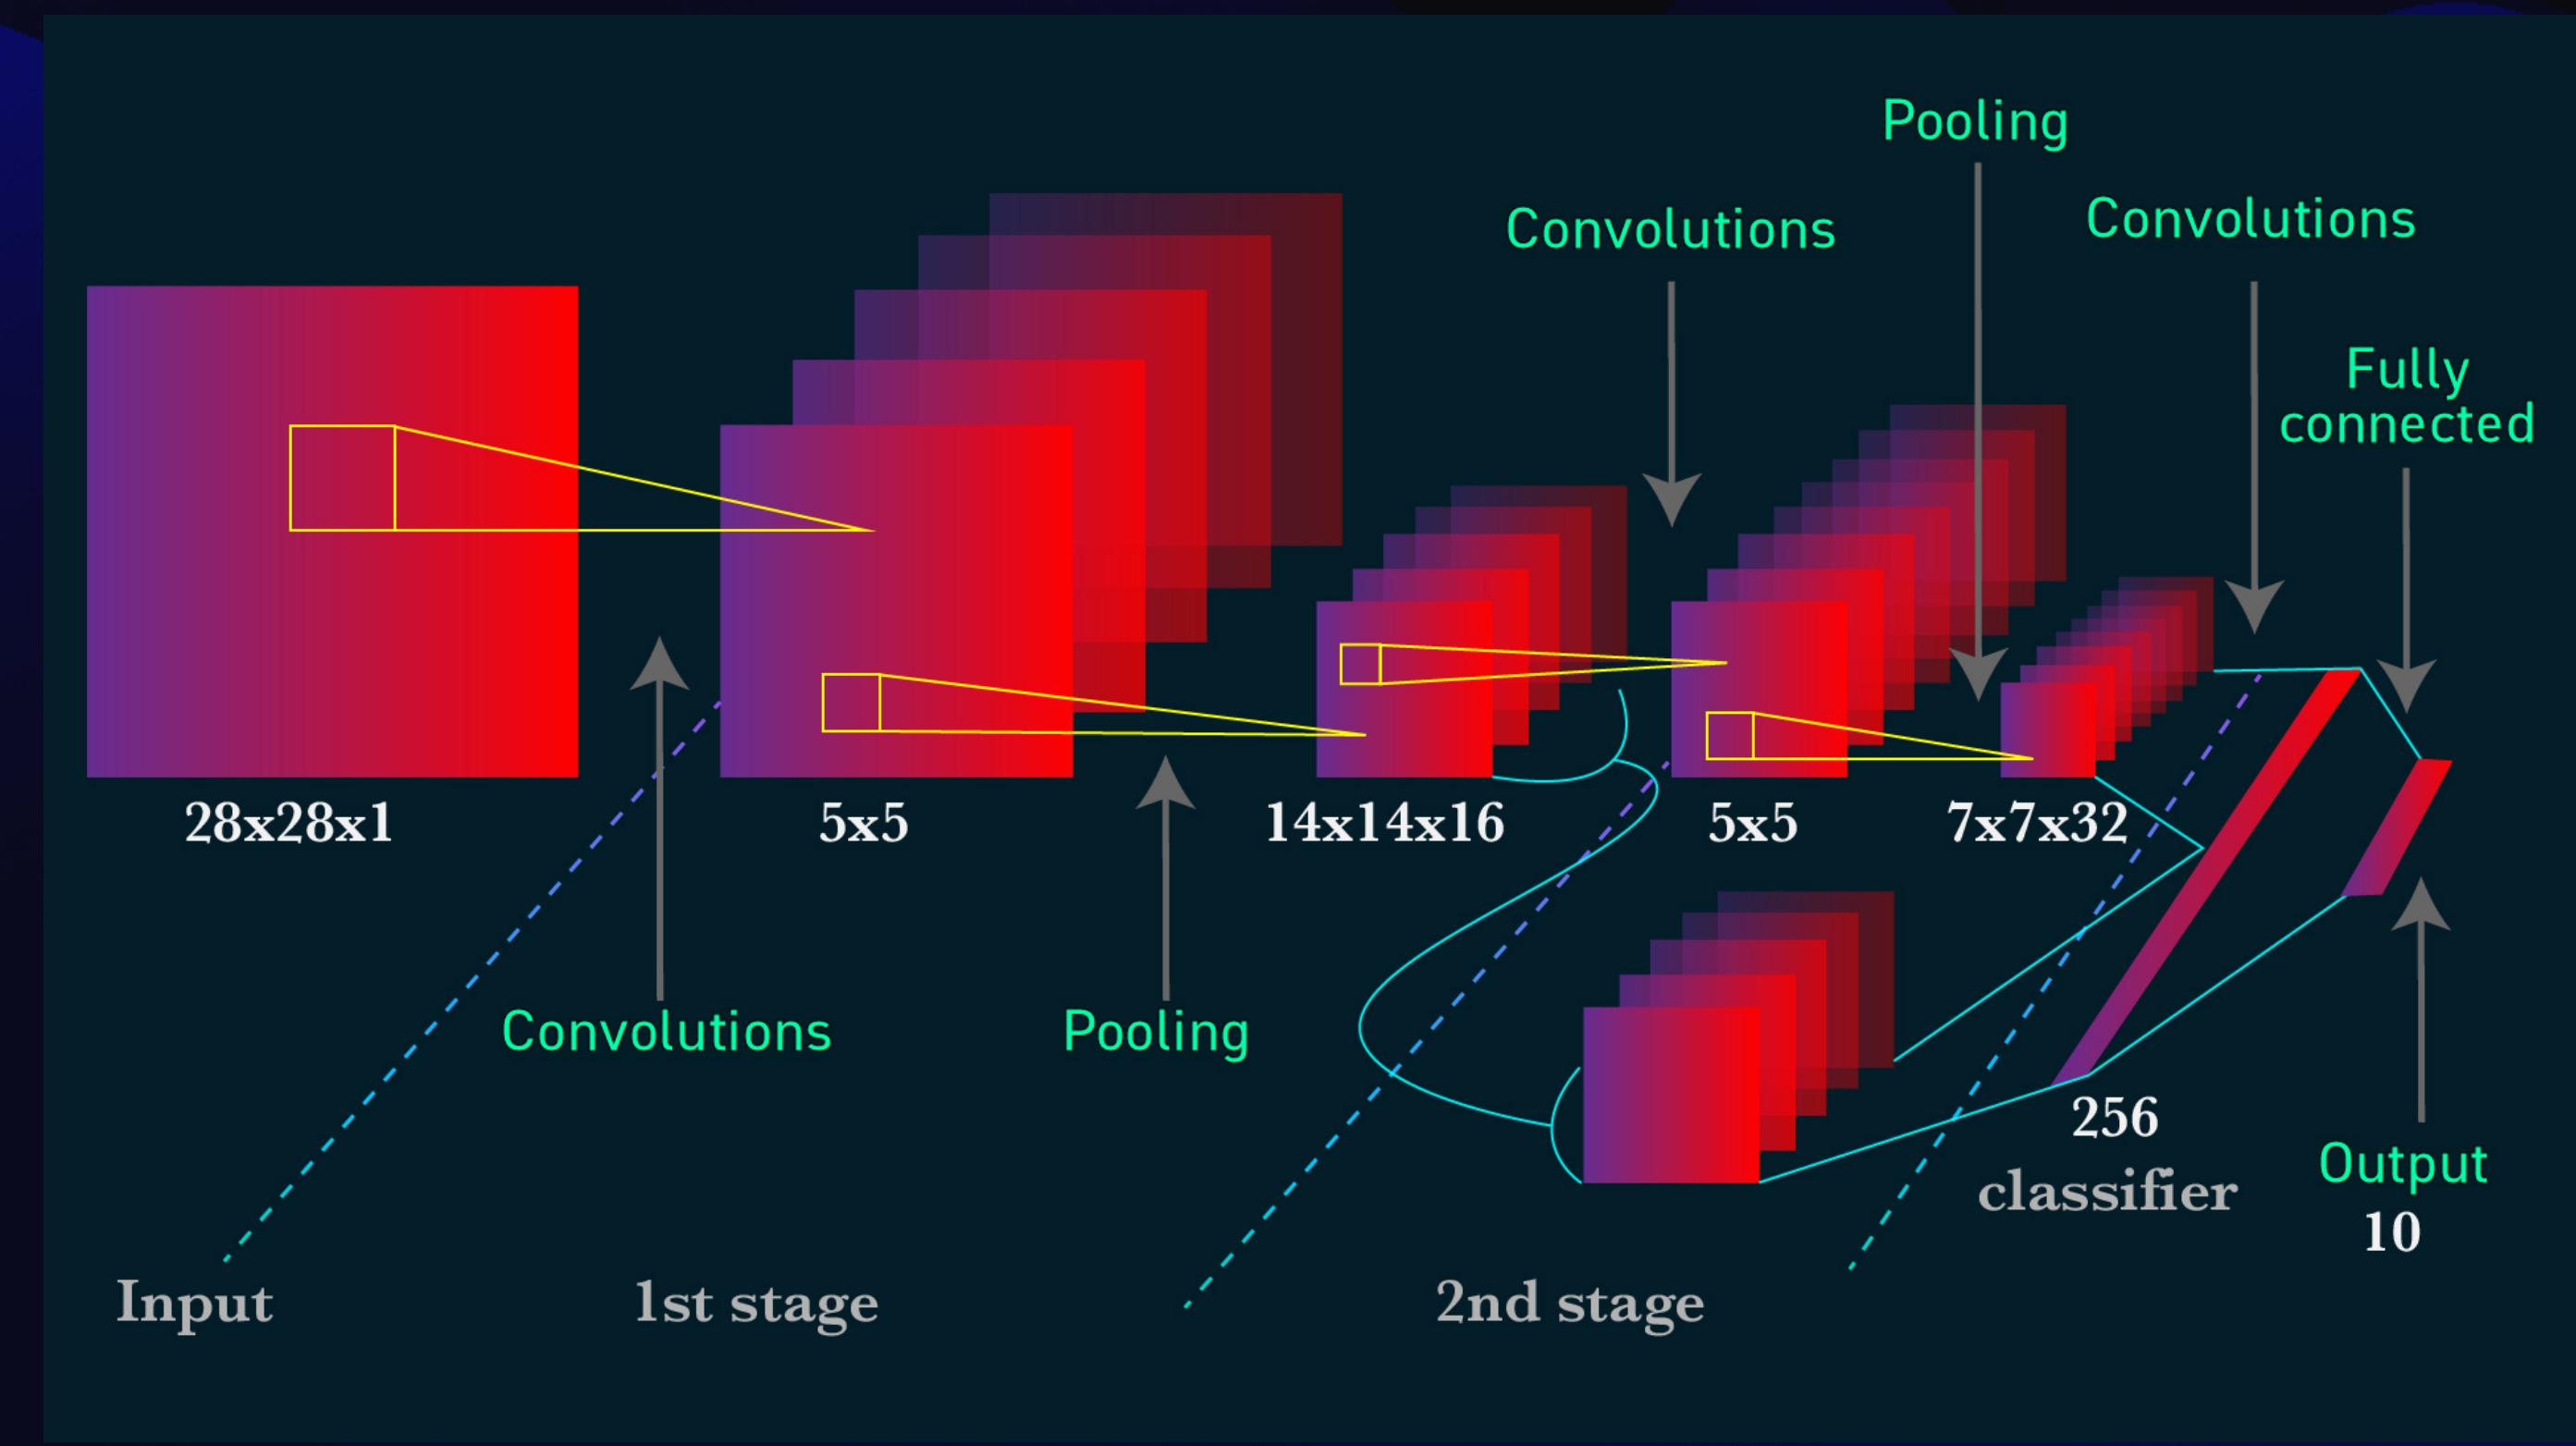
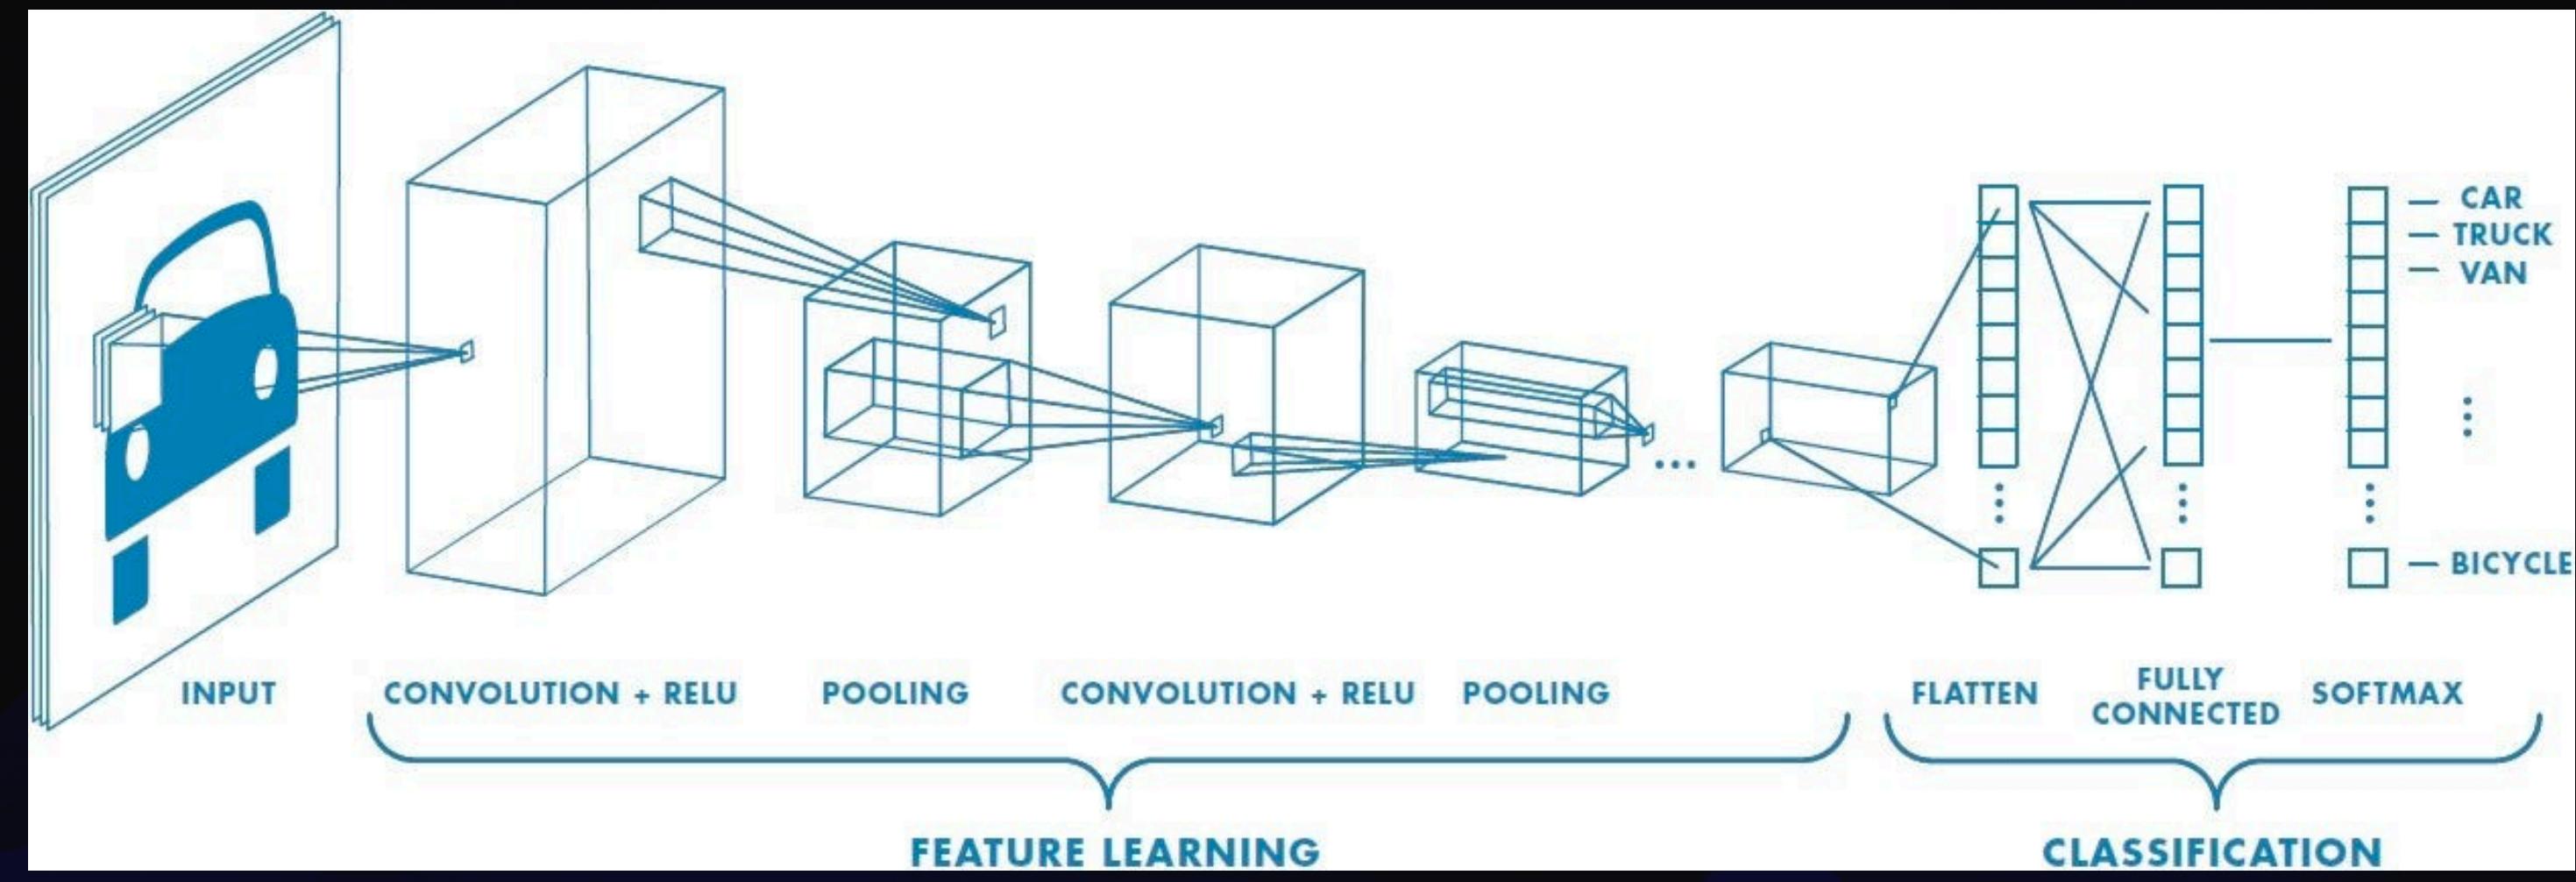


1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature



CNN Notebook

neural-network-zoo_CNNs.ipynb



<https://tinyurl.com/2fx8smzk>

<https://github.com/GageDeZoort/neural-network-zoo/tree/main>

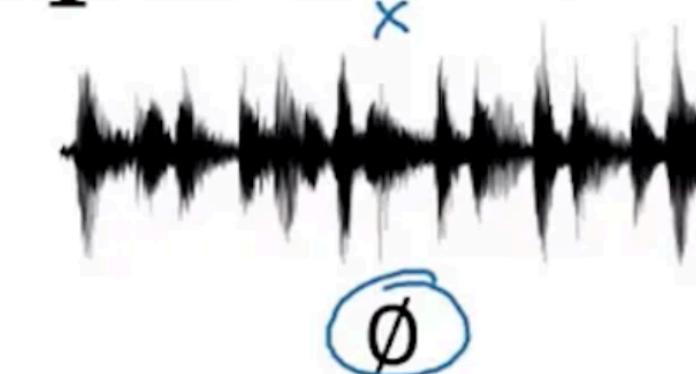
Recurrent Neural Networks

Recurrent Neural Networks

Predictions on Sequential Data

Examples of sequence data

Speech recognition



→ “The quick brown fox jumped over the lazy dog.”
y

Music generation



Sentiment classification

“There is nothing to like
in this movie.”



DNA sequence analysis

AGCCCCCTGTGAGGAACCTAG



AGCCCCTGTGAGGAACCTAG

Machine translation

Voulez-vous chanter avec
moi?



Do you want to sing with
me?

Video activity recognition



Running

Name entity recognition

Yesterday, Harry Potter
met Hermione Granger.

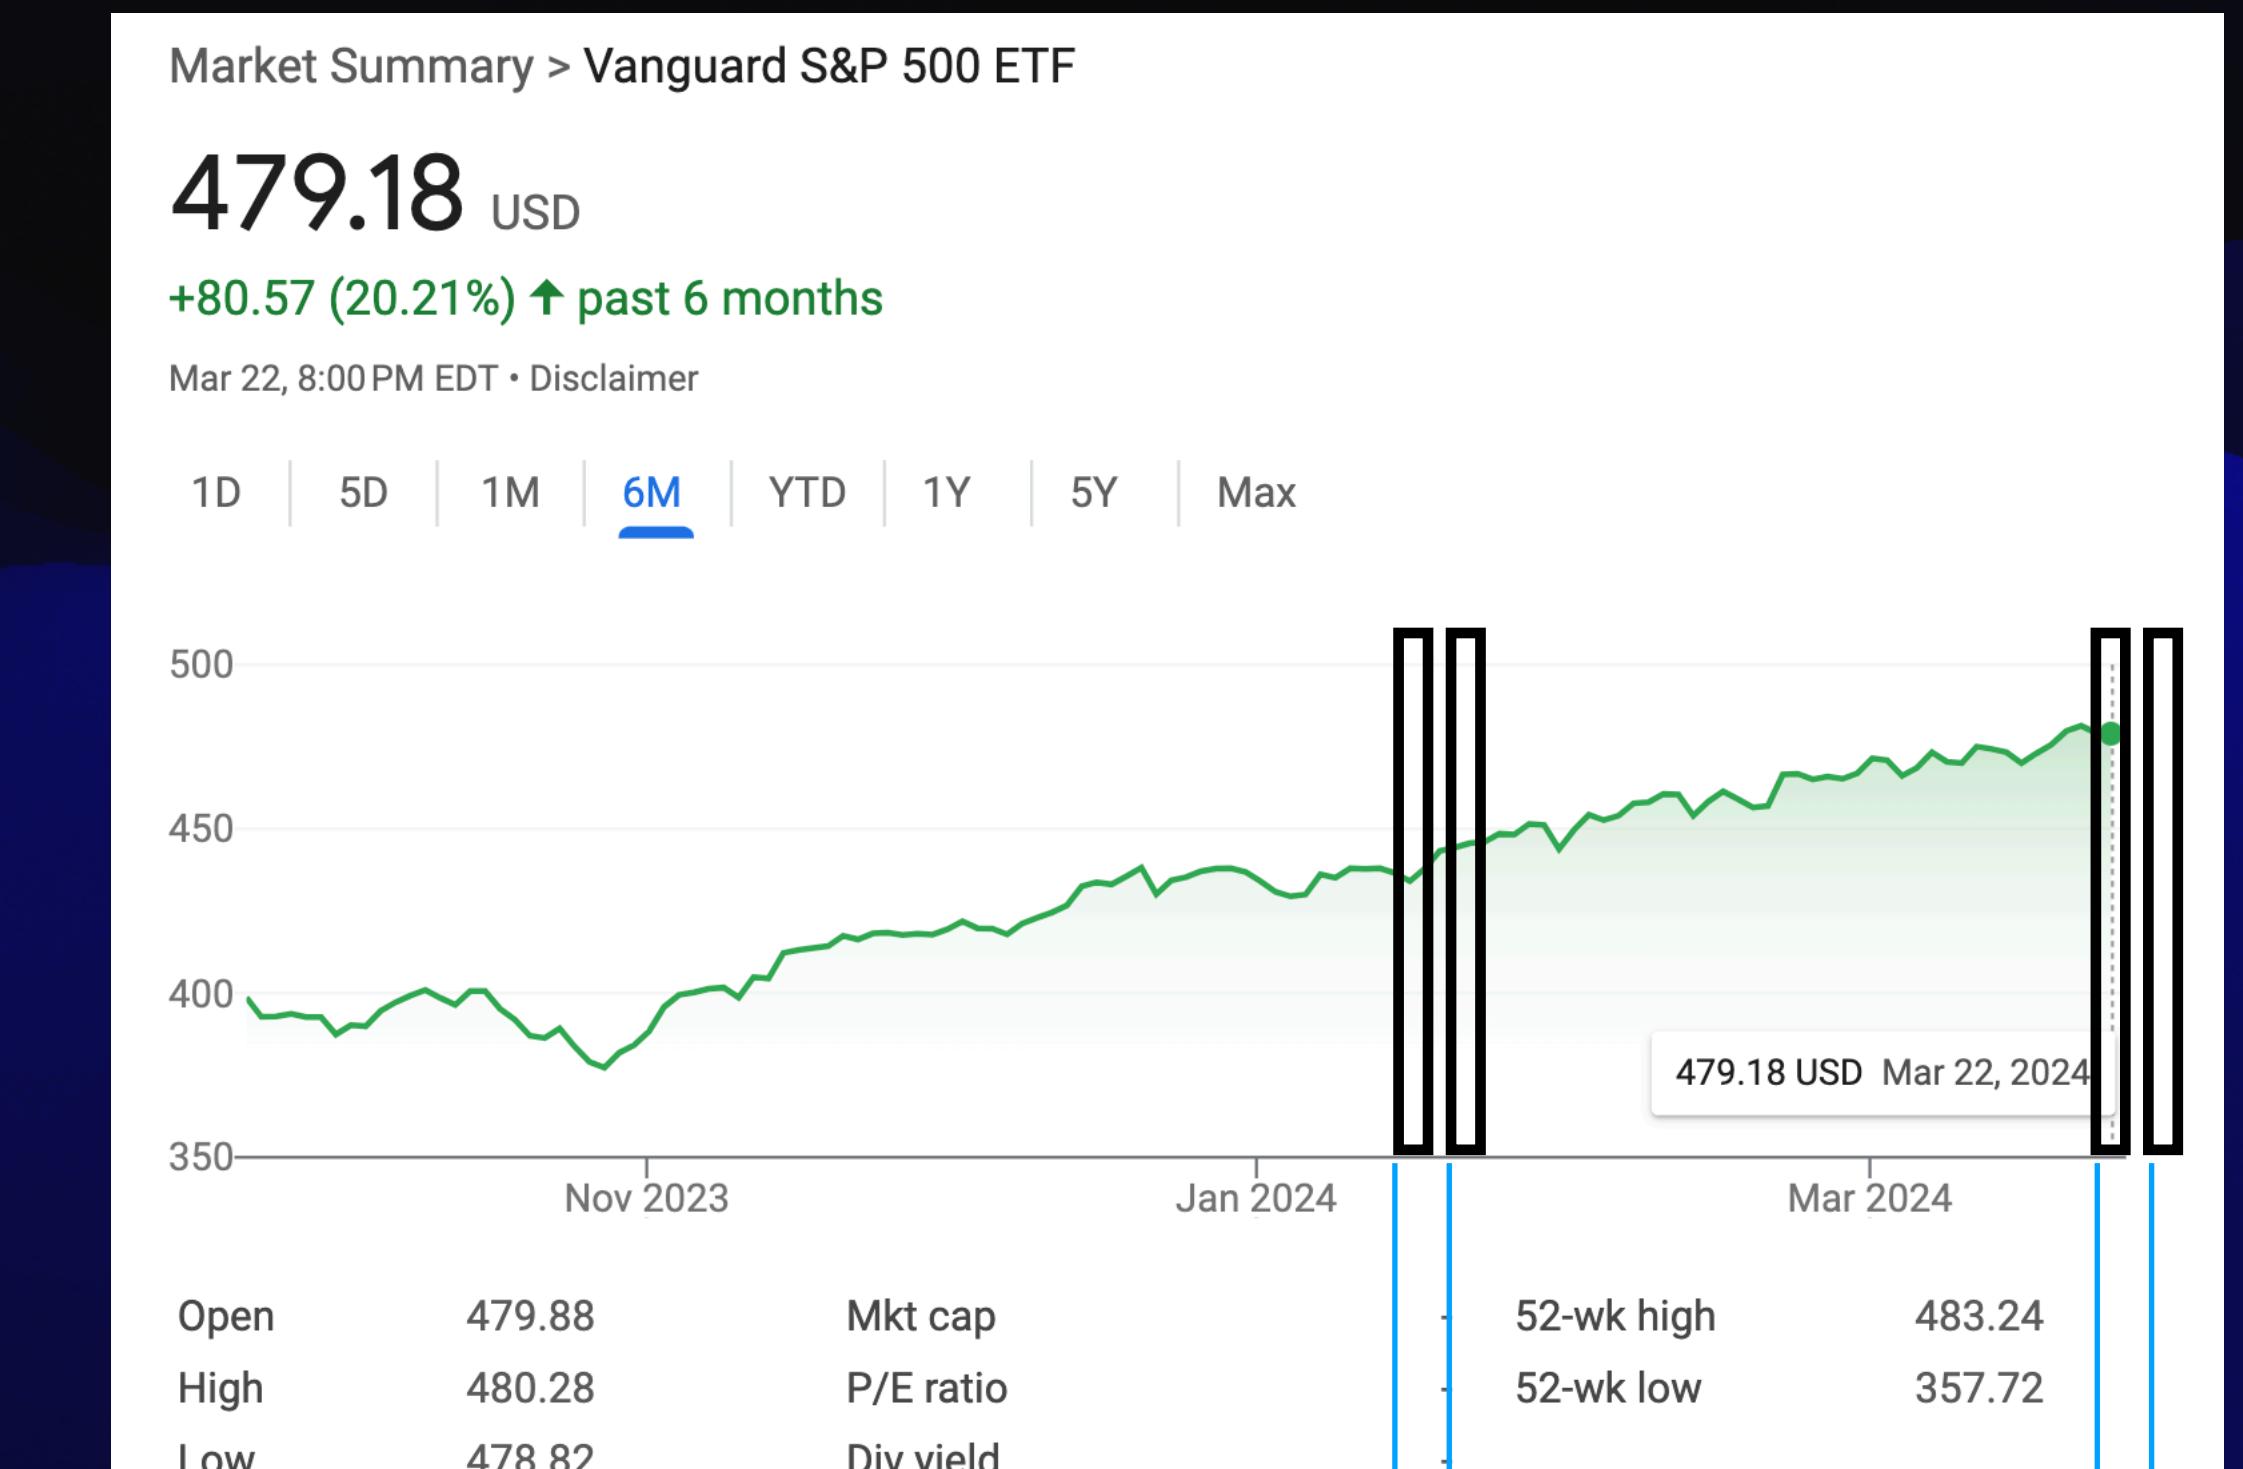


Yesterday, Harry Potter
met Hermione Granger.
Andrew Ng

Recurrent Neural Networks

Basic Idea

- Time-indexed inputs:
 $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(T)}$
 - word 1, word 2, word 3,...
 - stock price day 1, stock price day 2,...
- Parameter sharing: apply the *same function* (with learnable weights) at every time t



Function f

Function f

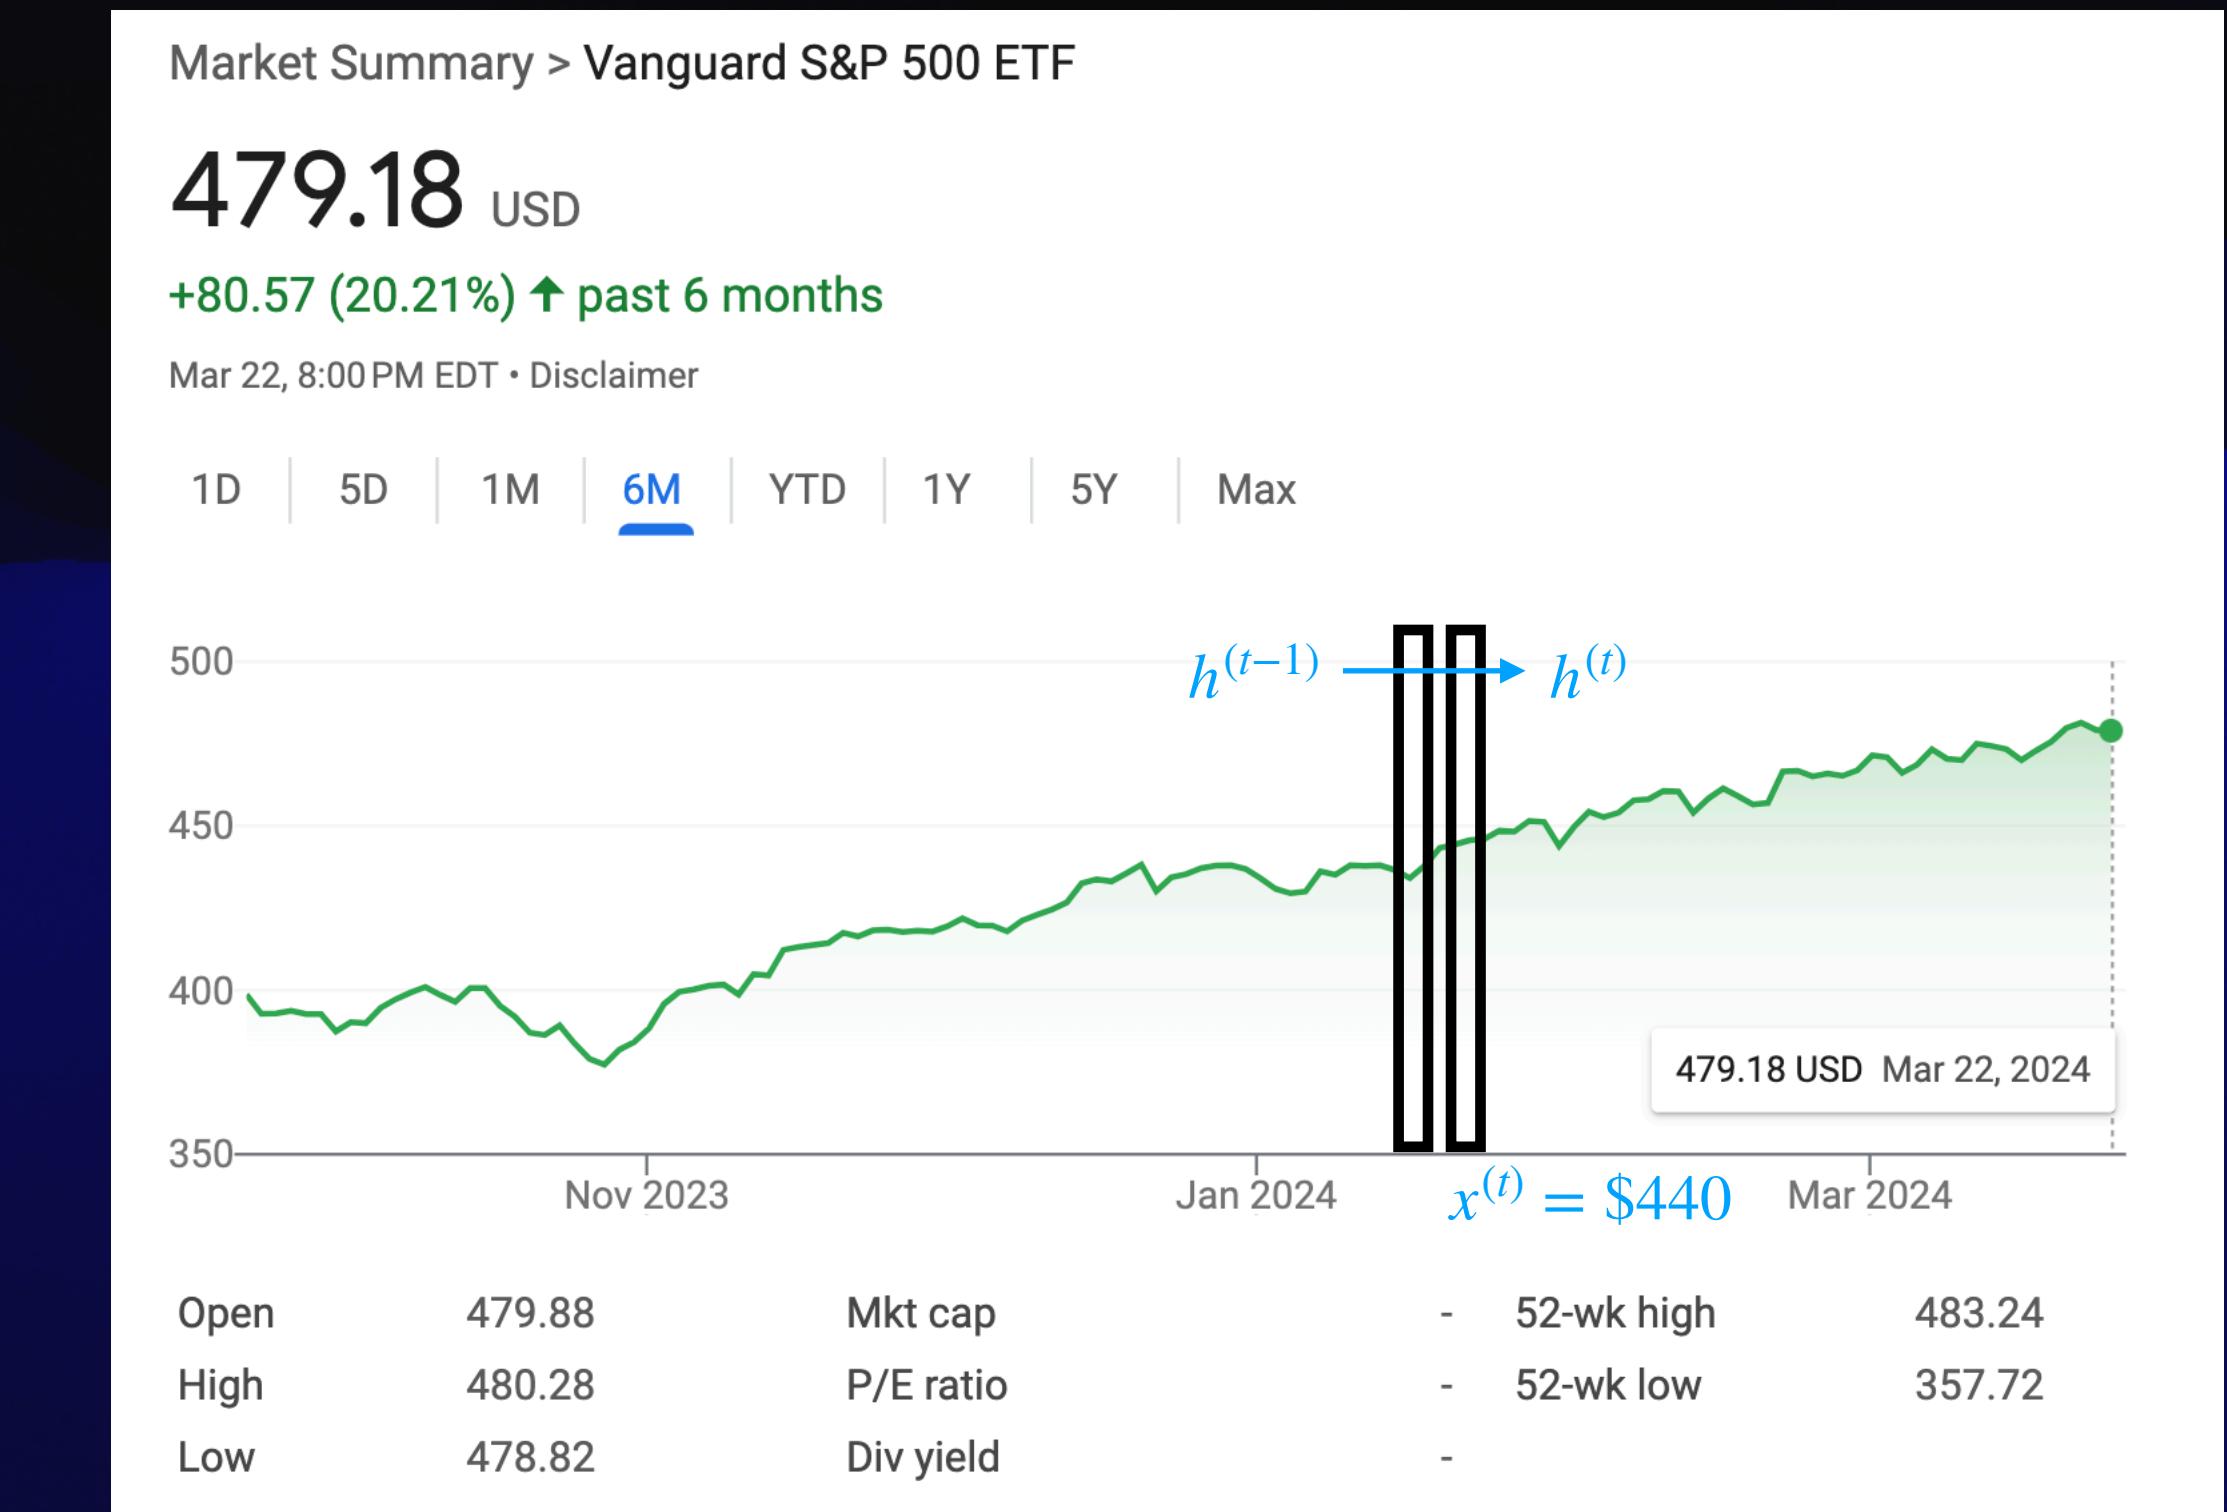
Recurrent Neural Networks

Basic Idea

Given a set of learnable parameters θ , the hidden units in many RNNs are calculated via

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)}; \theta)$$

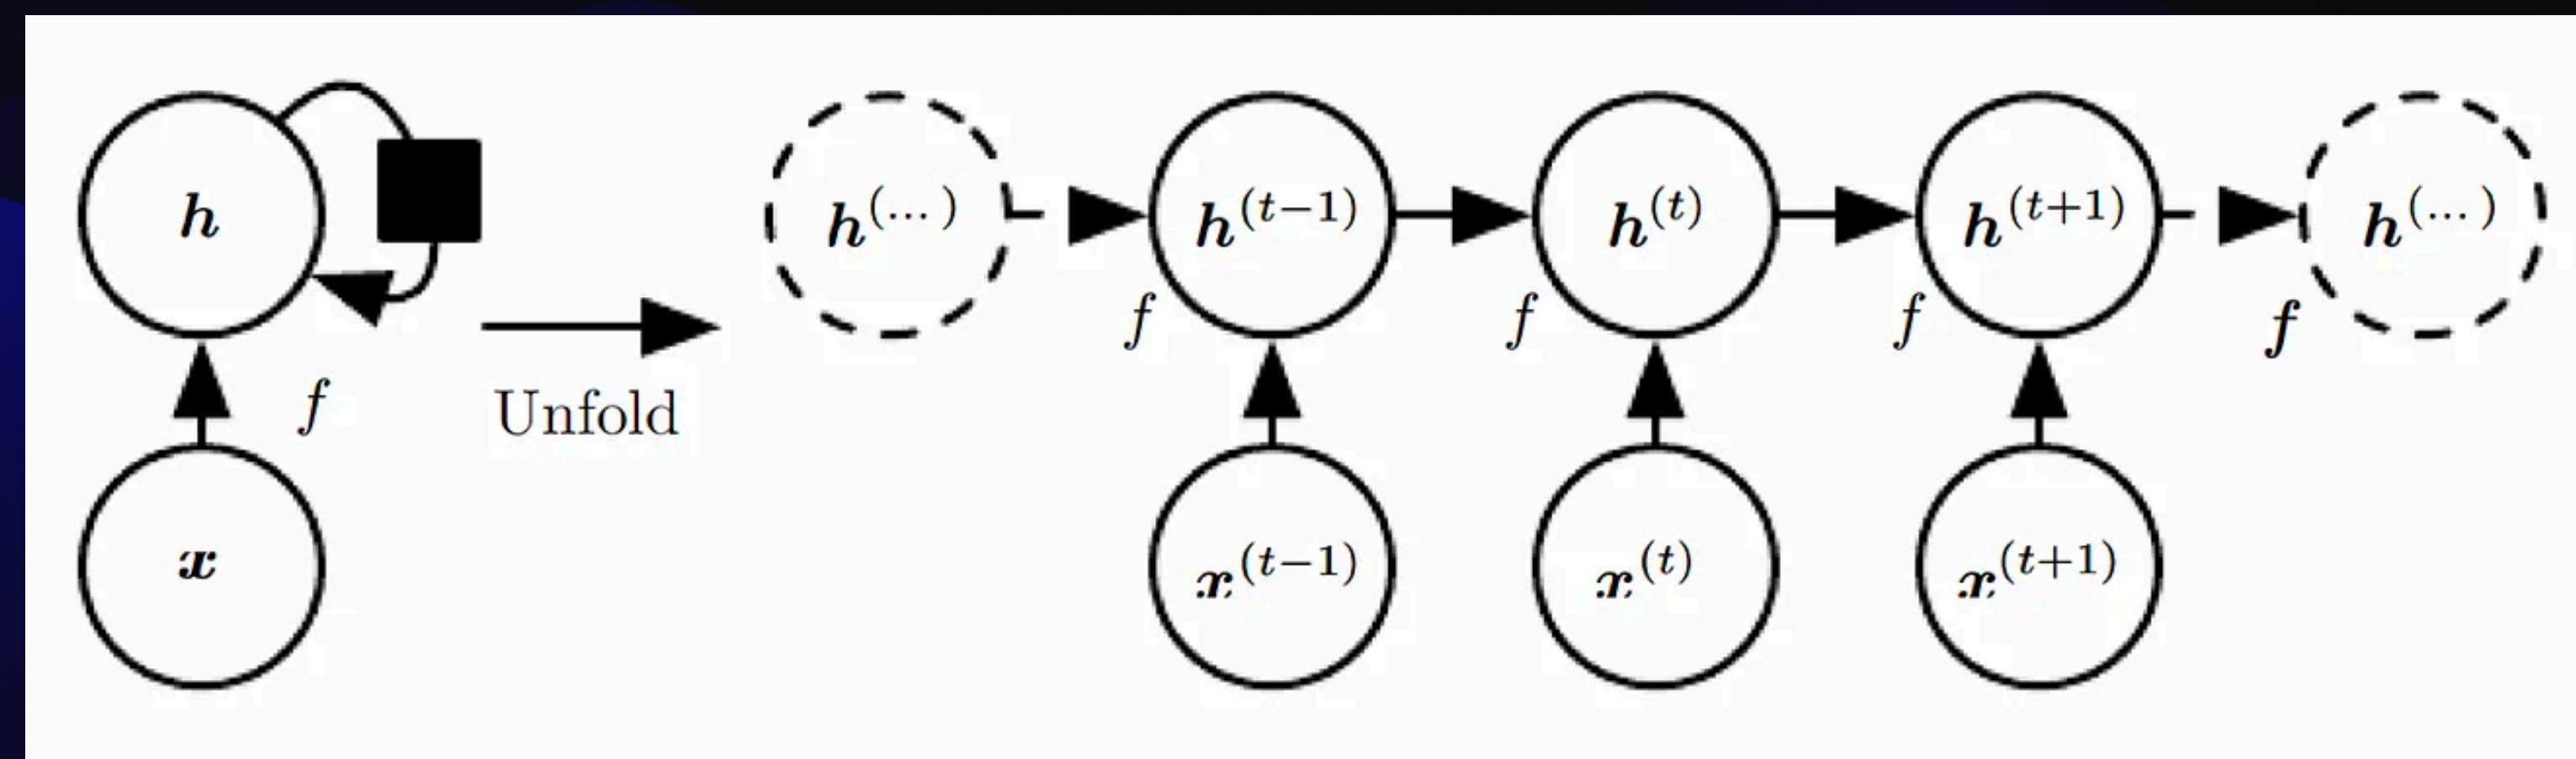
“Hidden Representations” $h(t)$ leveraged to make predictions



Recurrent Neural Networks

Basic Idea

Image from *Deep Learning* by Goodfellow et al.



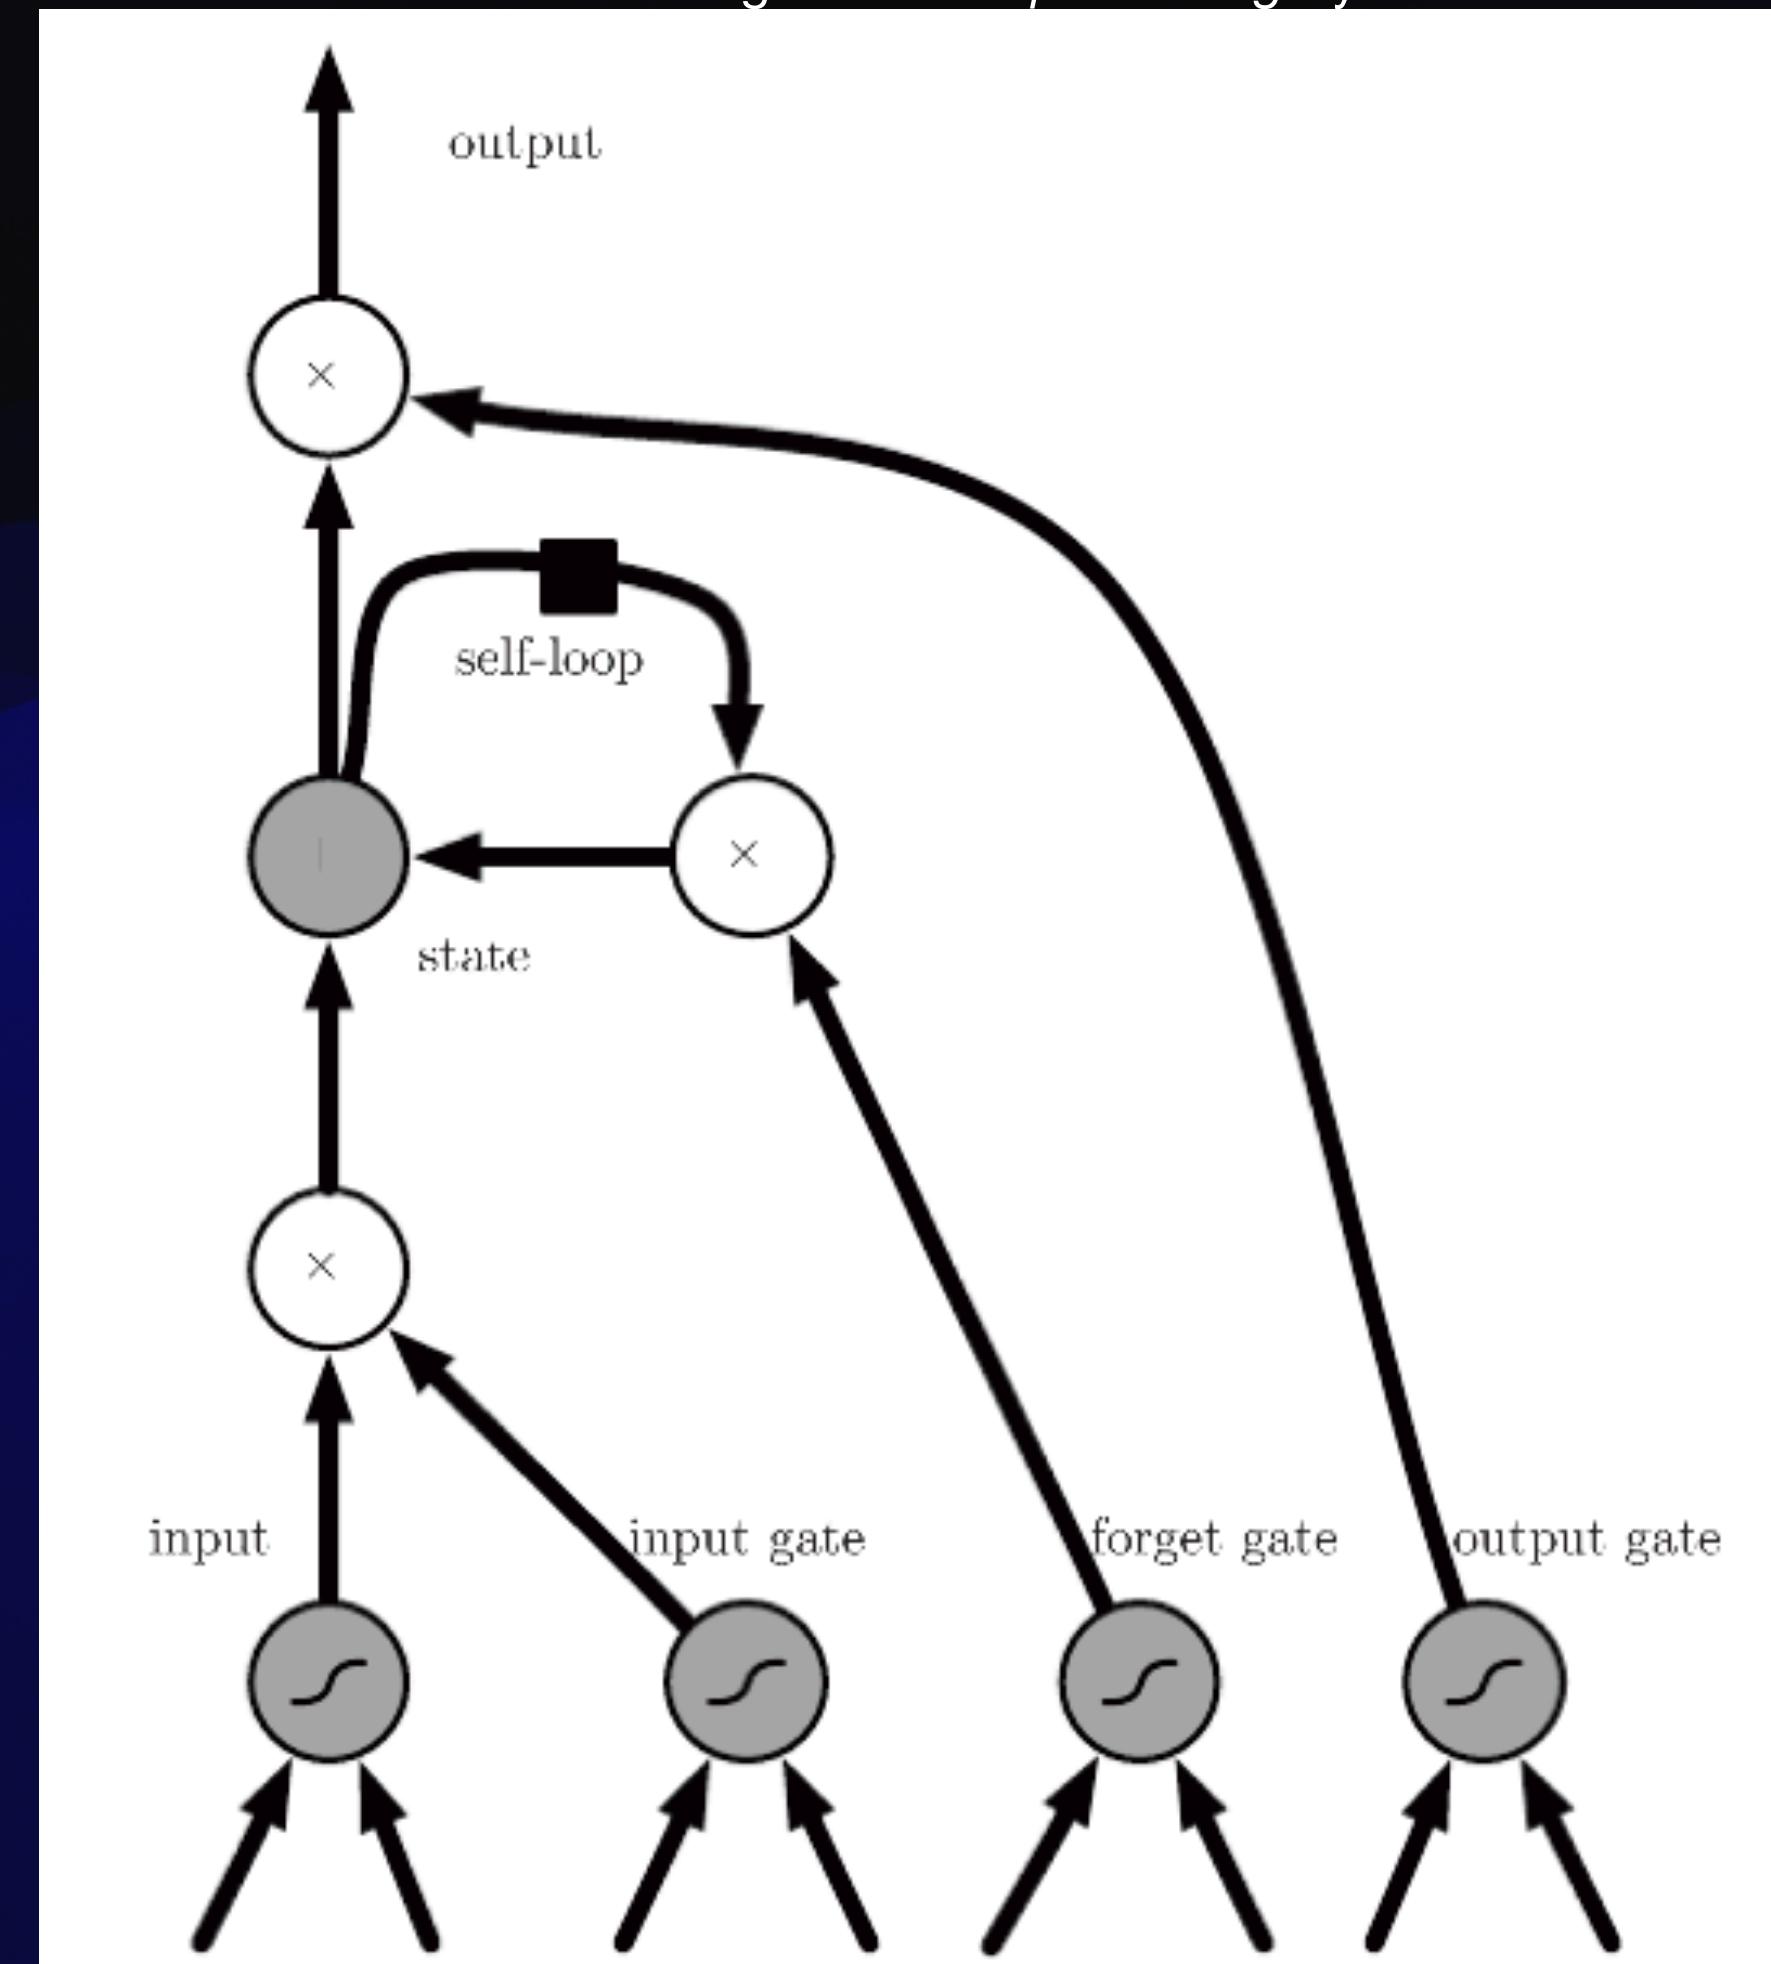
RNN with no outputs; information is processed sequentially, taking into account both $\mathbf{x}^{(t)}$ and $h^{(t-1)}$ but applying the *same function* $f(\cdot ; \theta)$ at each timestep

Long Short-Term Memory (LSTM)

An Upgraded RNN Module

- RNNs are finicky to train; they often suffer from exploding/vanishing gradients
- This has motivated the development of more advanced RNNs like LSTMs
- The LSTM is a recurrent “cell” that is applied to all timesteps equally

Image from *Deep Learning* by Goodfellow et al.



RNN Notebook

neural-network-zoo_RNNs.ipynb



<https://tinyurl.com/2fx8smzk>

<https://github.com/GageDeZoort/neural-network-zoo/tree/main>

Graph Neural Networks

Graph Neural Networks

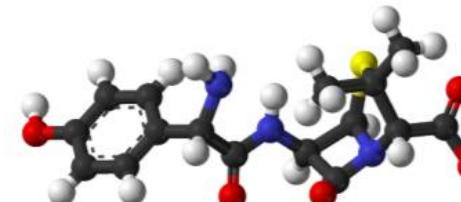
Intro to Graphs

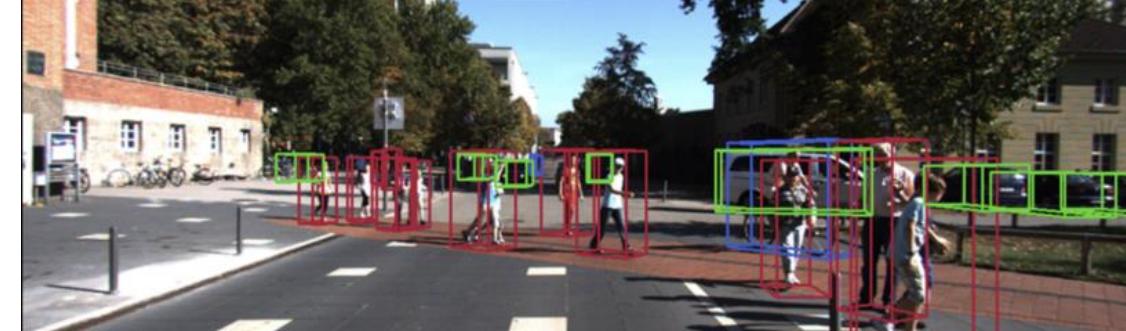
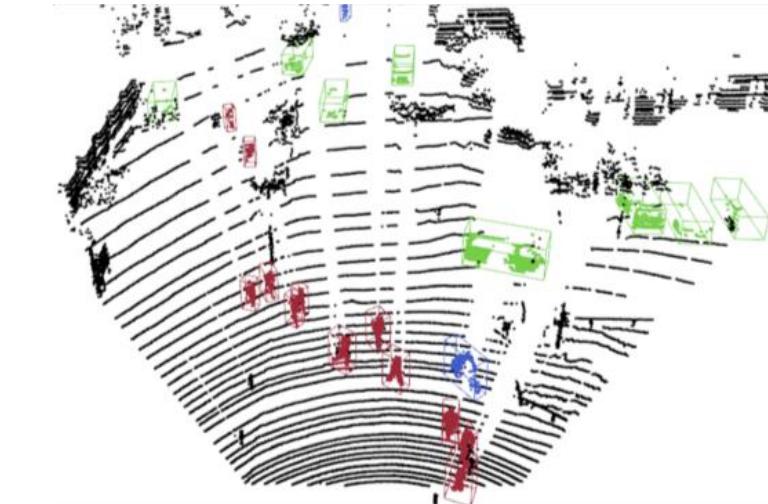
- Graphs: vertices (nodes) connected by edges (relationships)
- Examples:
 - Social networks of people connected by friendships
 - Molecules of atoms connected by nodes
 - Planets connected by mutual gravitational interactions



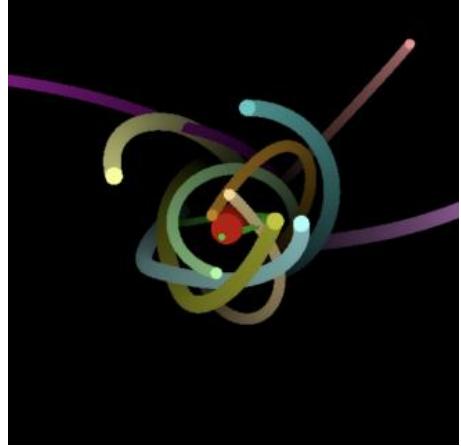
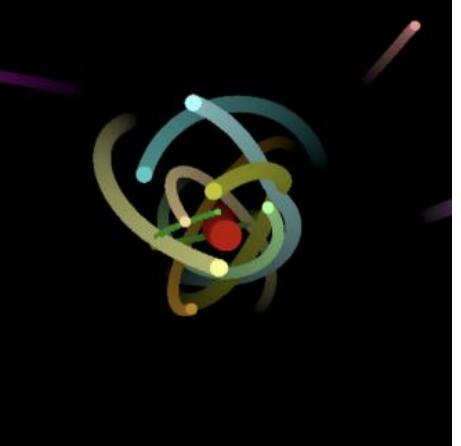
Graph Neural Networks

Graph Learning Tasks

DRUG DISCOVERY $GNN($  $) \rightarrow$ molecule property

INSTANCE SEGMENTATION $GNN($  $) \rightarrow$ 

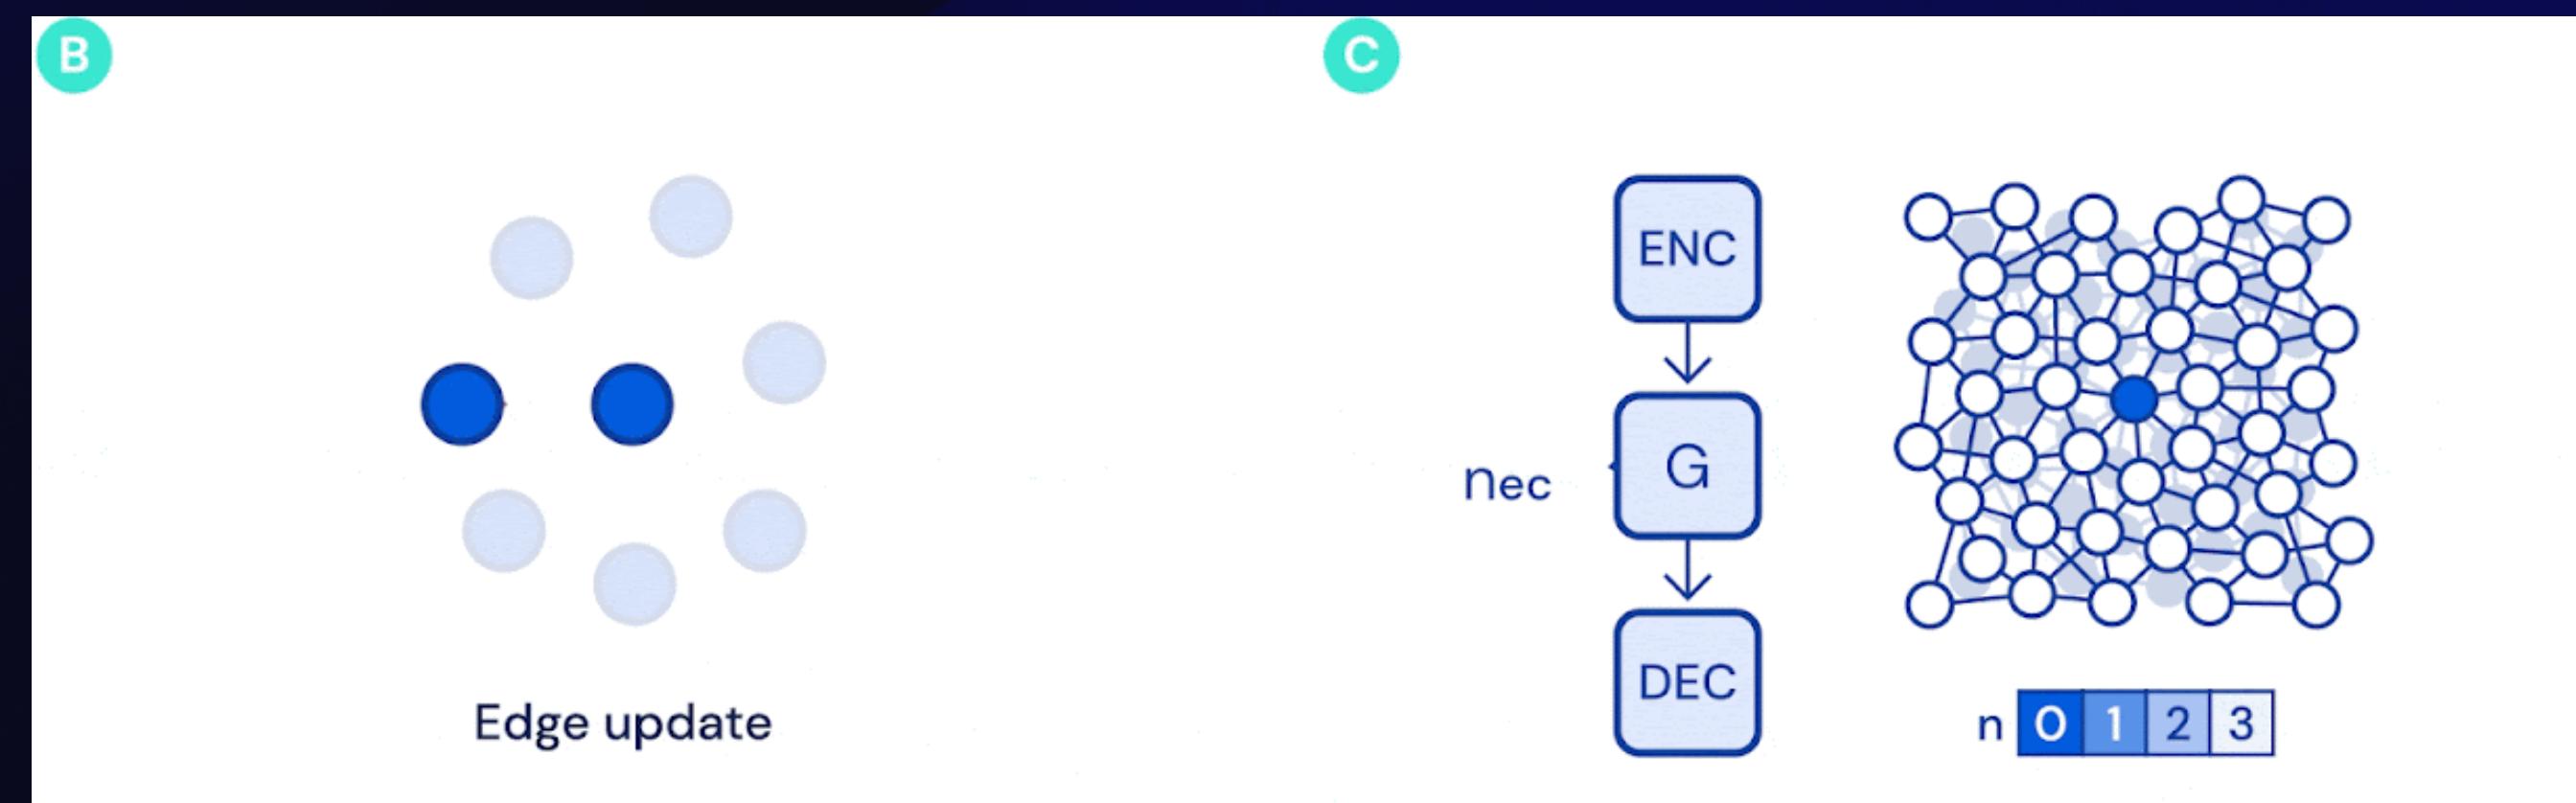
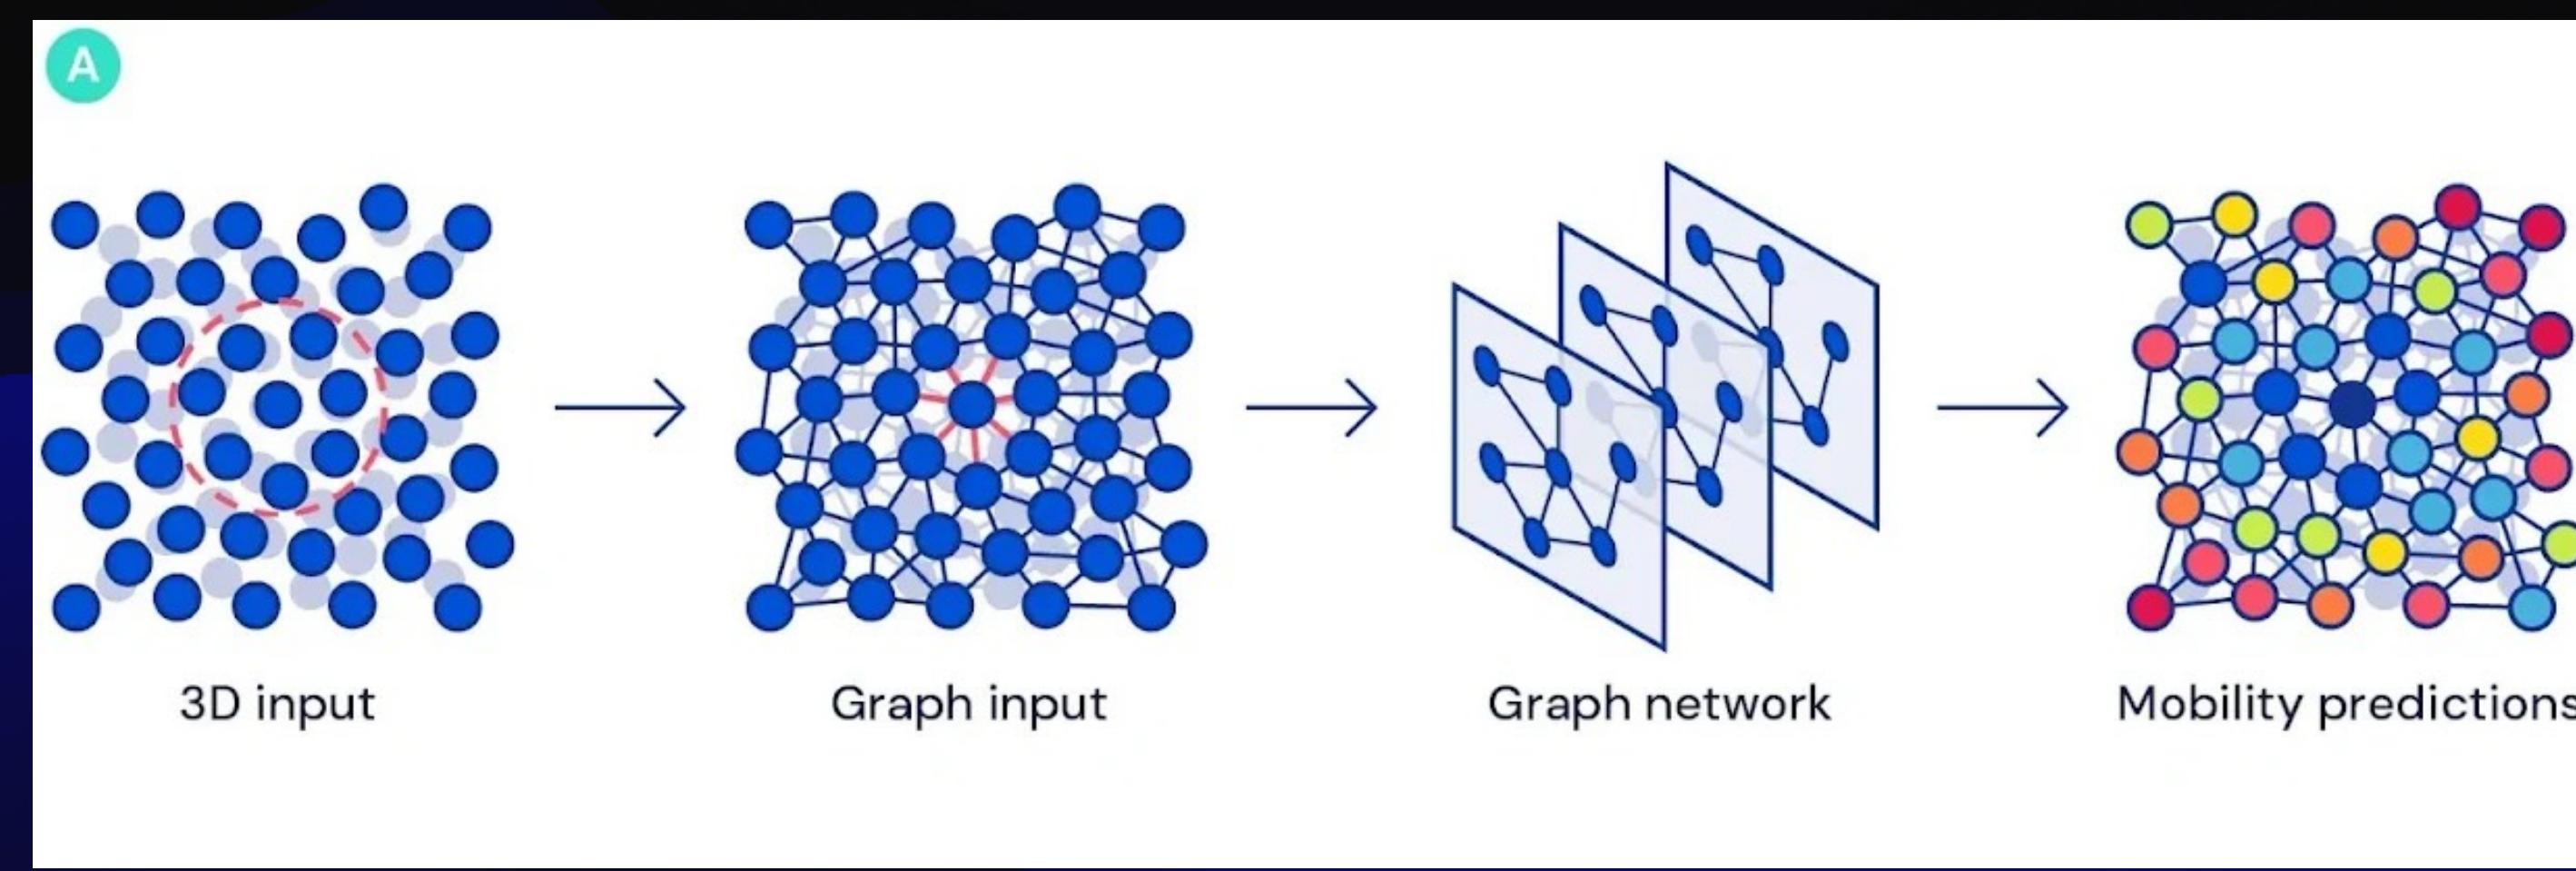
[2003.01251.pdf \(arxiv.org\)](https://arxiv.org/pdf/2003.01251.pdf)

PHYSICS SIMULATION $GNN($  $) \rightarrow$ 

[\[1612.00222\] Interaction Networks for Learning about Objects, Relations and Physics \(arxiv.org\)](https://arxiv.org/pdf/1612.00222.pdf)

Graph Neural Networks

Message Passing NNs (MPNNs)



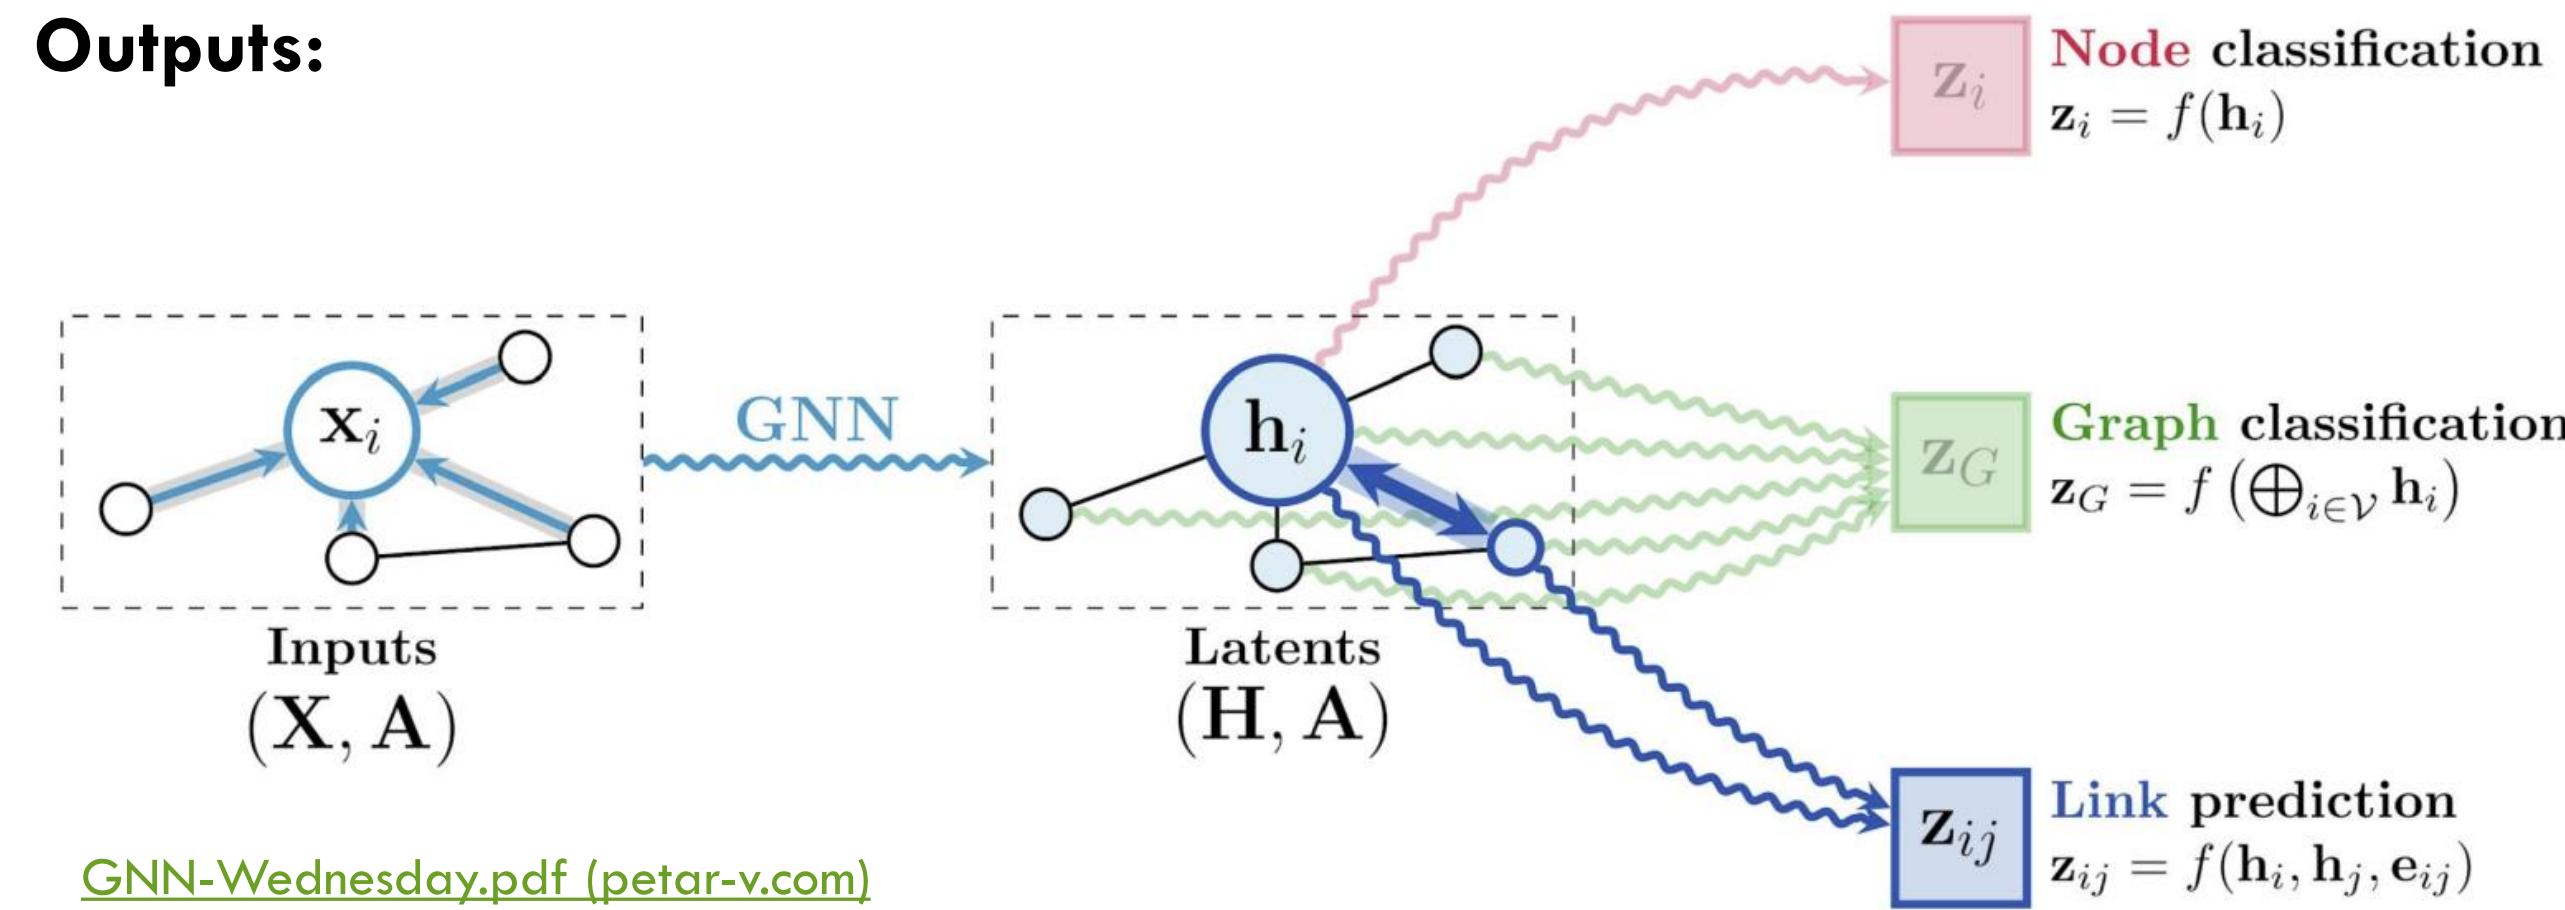
Graph Neural Networks

Message Passing NNs (MPNNs)

Generic MPNN Layers:

$$\mathbf{h}_u^{(k)} = \phi^{(k)} \left[\mathbf{h}_u^{(k-1)}, \bigoplus_{v \in N(u)} \psi^{(k)} \left(\mathbf{h}_u^{(k-1)}, \mathbf{h}_v^{(k-1)}, \mathbf{e}_{uv}^{(k-1)} \right) \right]$$

Outputs:



[GNN-Wednesday.pdf \(petar-v.com\)](#)



A dark blue background featuring three horizontal layers of wavy, translucent shapes. The top layer is a very dark navy blue, the middle layer is a medium navy blue, and the bottom layer is a bright, saturated navy blue. The waves are soft and fluid, creating a sense of depth and motion.

GNN Notebook!

GNN Notebook

neural-network-zoo_GNNs.ipynb



<https://tinyurl.com/2fx8smzk>

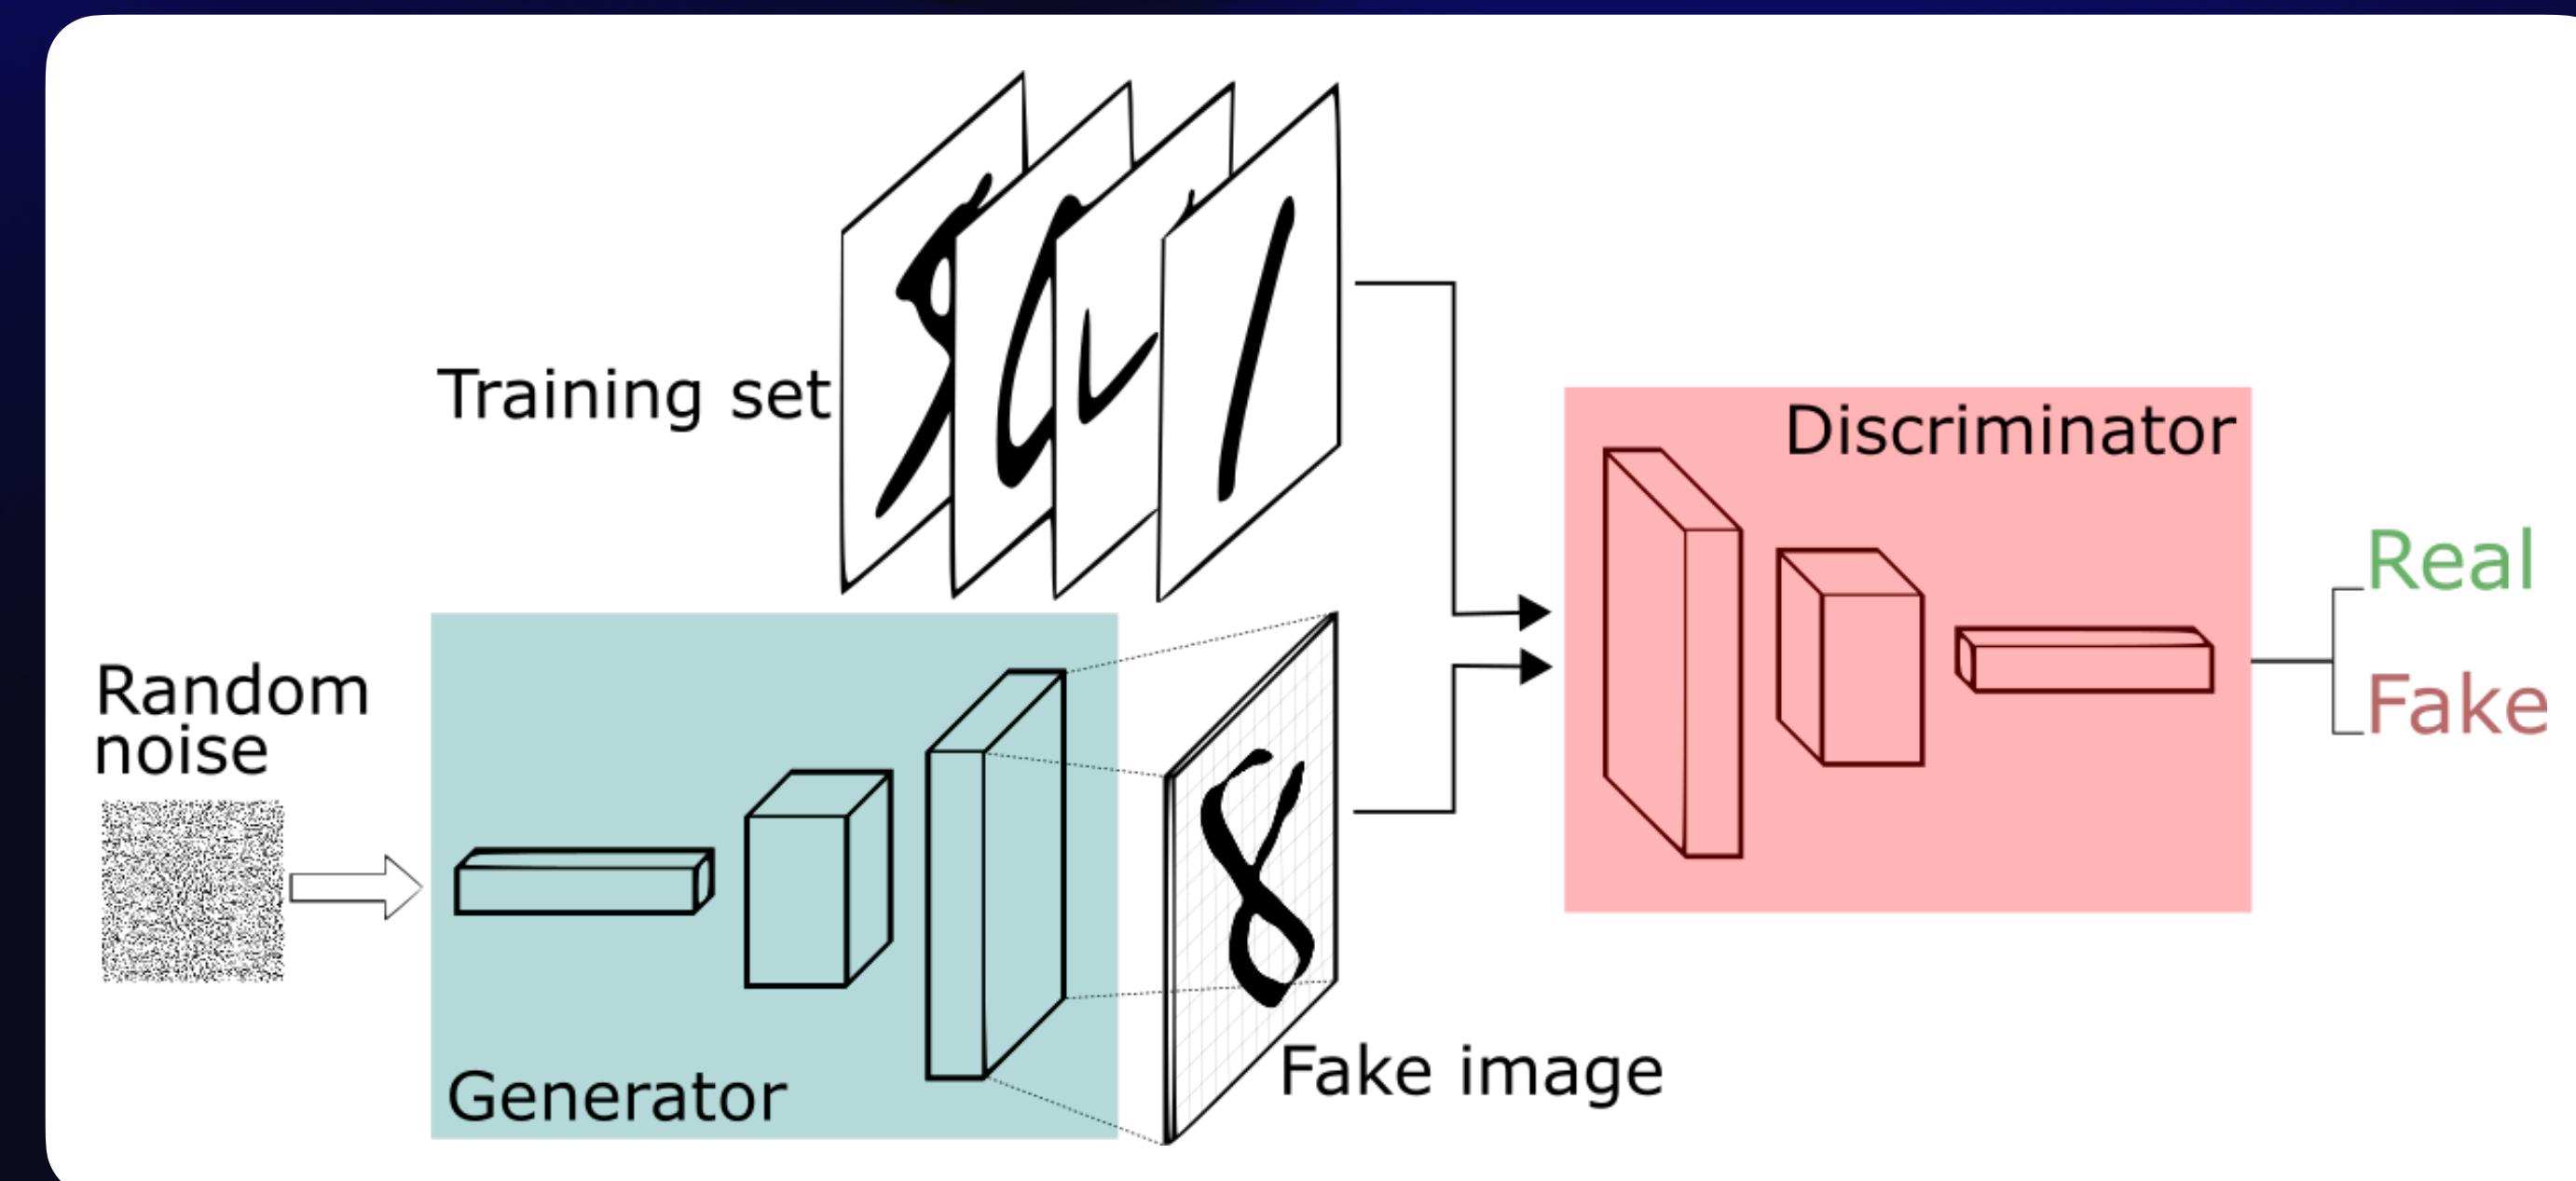
<https://github.com/GageDeZoort/neural-network-zoo/tree/main>

Generative Adversarial Networks

Generative Adversarial Networks (GANs)

Intro to GANs

- “Generative” AI: use ML to create new images, sounds, etc.
- GANs: two agents (the *generator* and the *discriminator*) are given competing tasks:



GAN Notebook

neural-network-zoo_GANs.ipynb



<https://tinyurl.com/2fx8smzk>

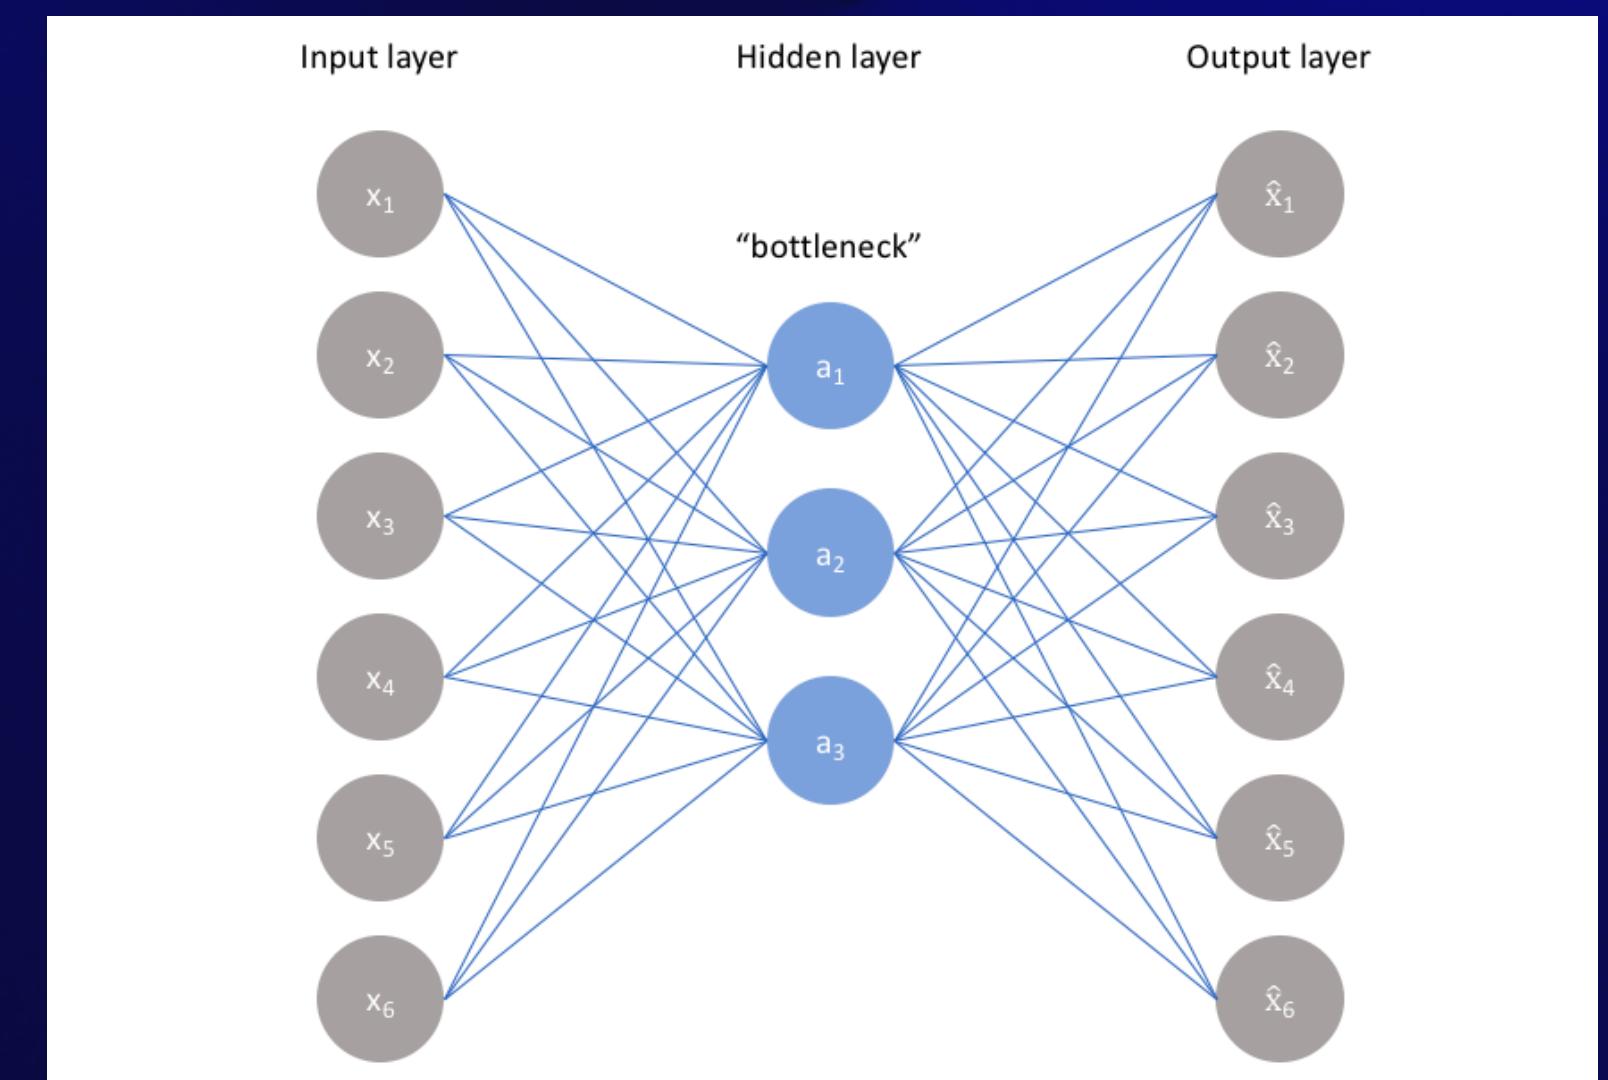
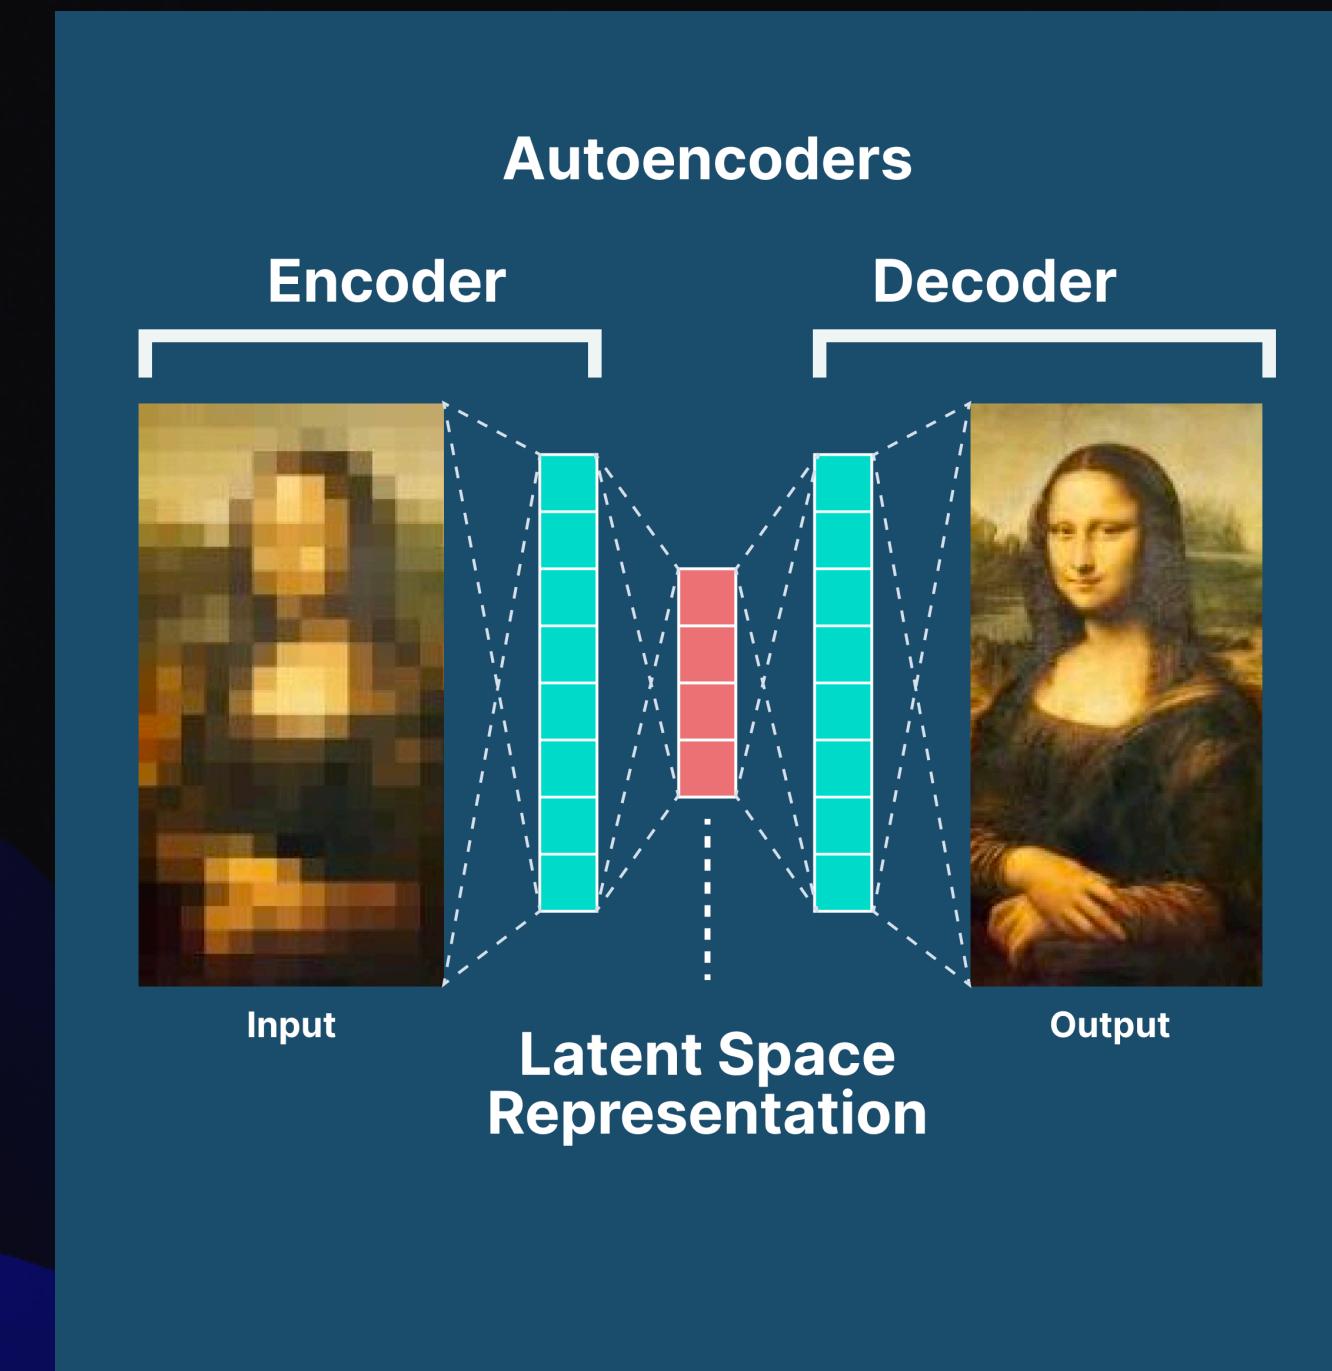
<https://github.com/GageDeZoort/neural-network-zoo/tree/main>

Autoencoders

Autoencoders

Learning Efficient Codings

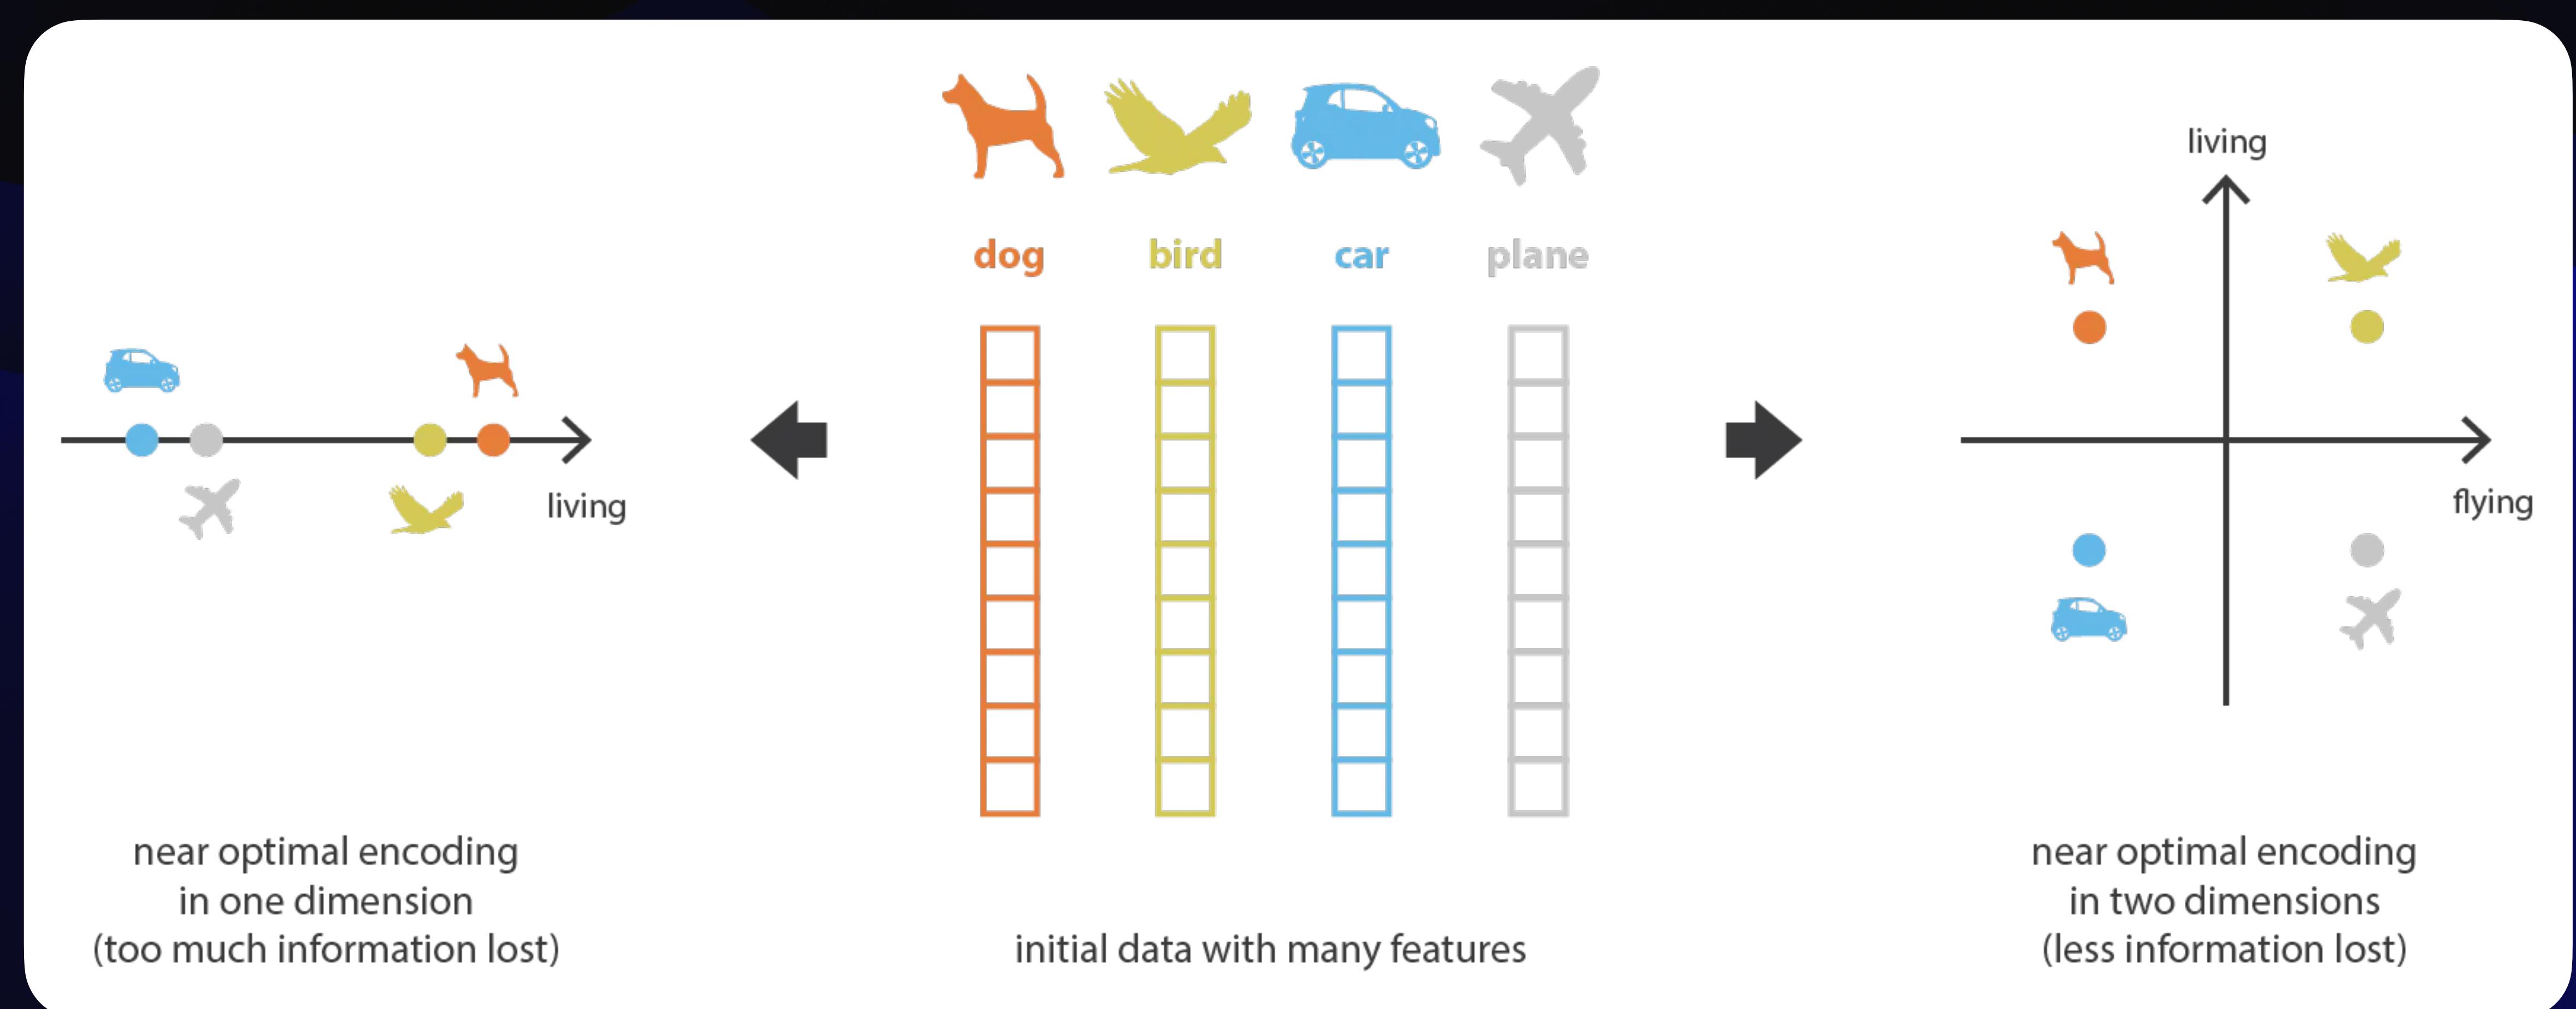
- Autoencoders are used to produce compressed data representations (*dimensionality reduction*)
 - **Encoder:** produces a lower-dimensional (compressed) “latent” representation of the input data
 - **Decoder:** given the compressed representation, reconstruct the original data
- Decoded representations typically less noisy,
- Uses: efficient encoding, image denoting, generative modeling, anomaly detection



<https://www.v7labs.com/blog/autoencoders-guide#:~:text=An%20autoencoder%20is%20an%20unsupervised,even%20generation%20of%20image%20data.>

Autoencoders

Learning Efficient Codings





Autoencoder Notebook!

AE Notebook

neural-network-zoo_AEs.ipynb



<https://tinyurl.com/2fx8smzk>

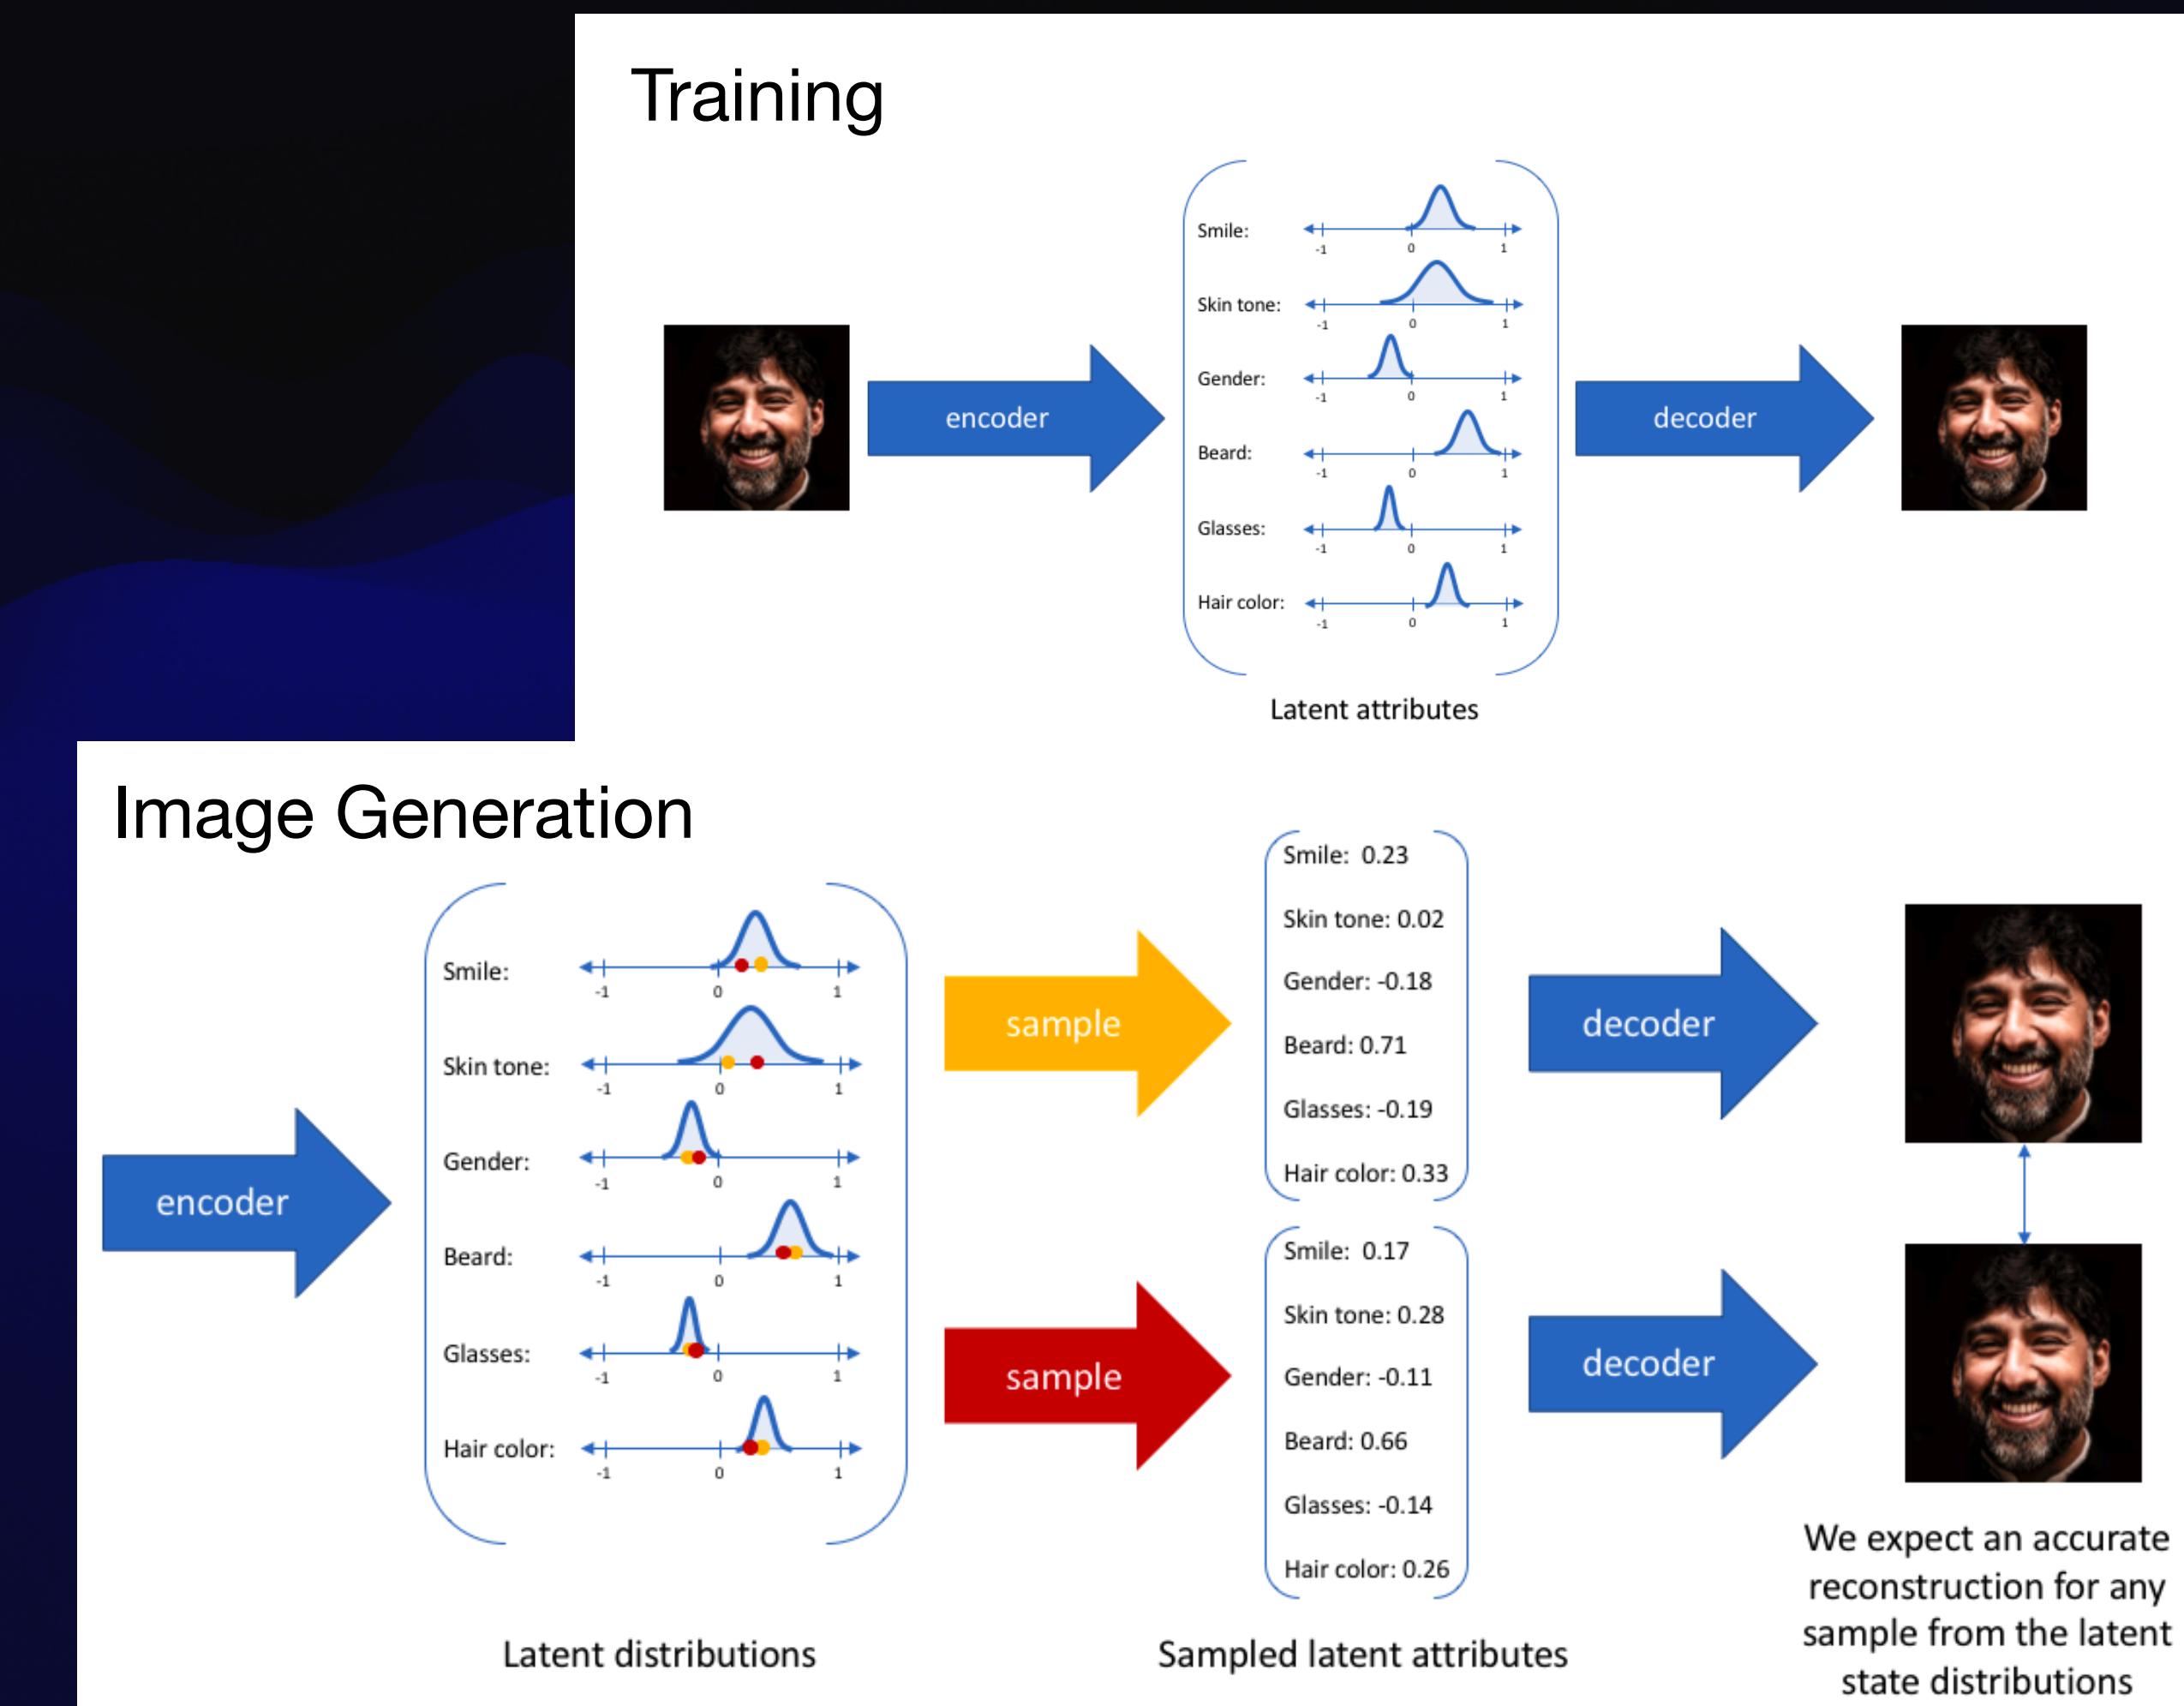
<https://github.com/GageDeZoort/neural-network-zoo/tree/main>

Variational Autoencoders

Variational Autoencoder (VAEs)

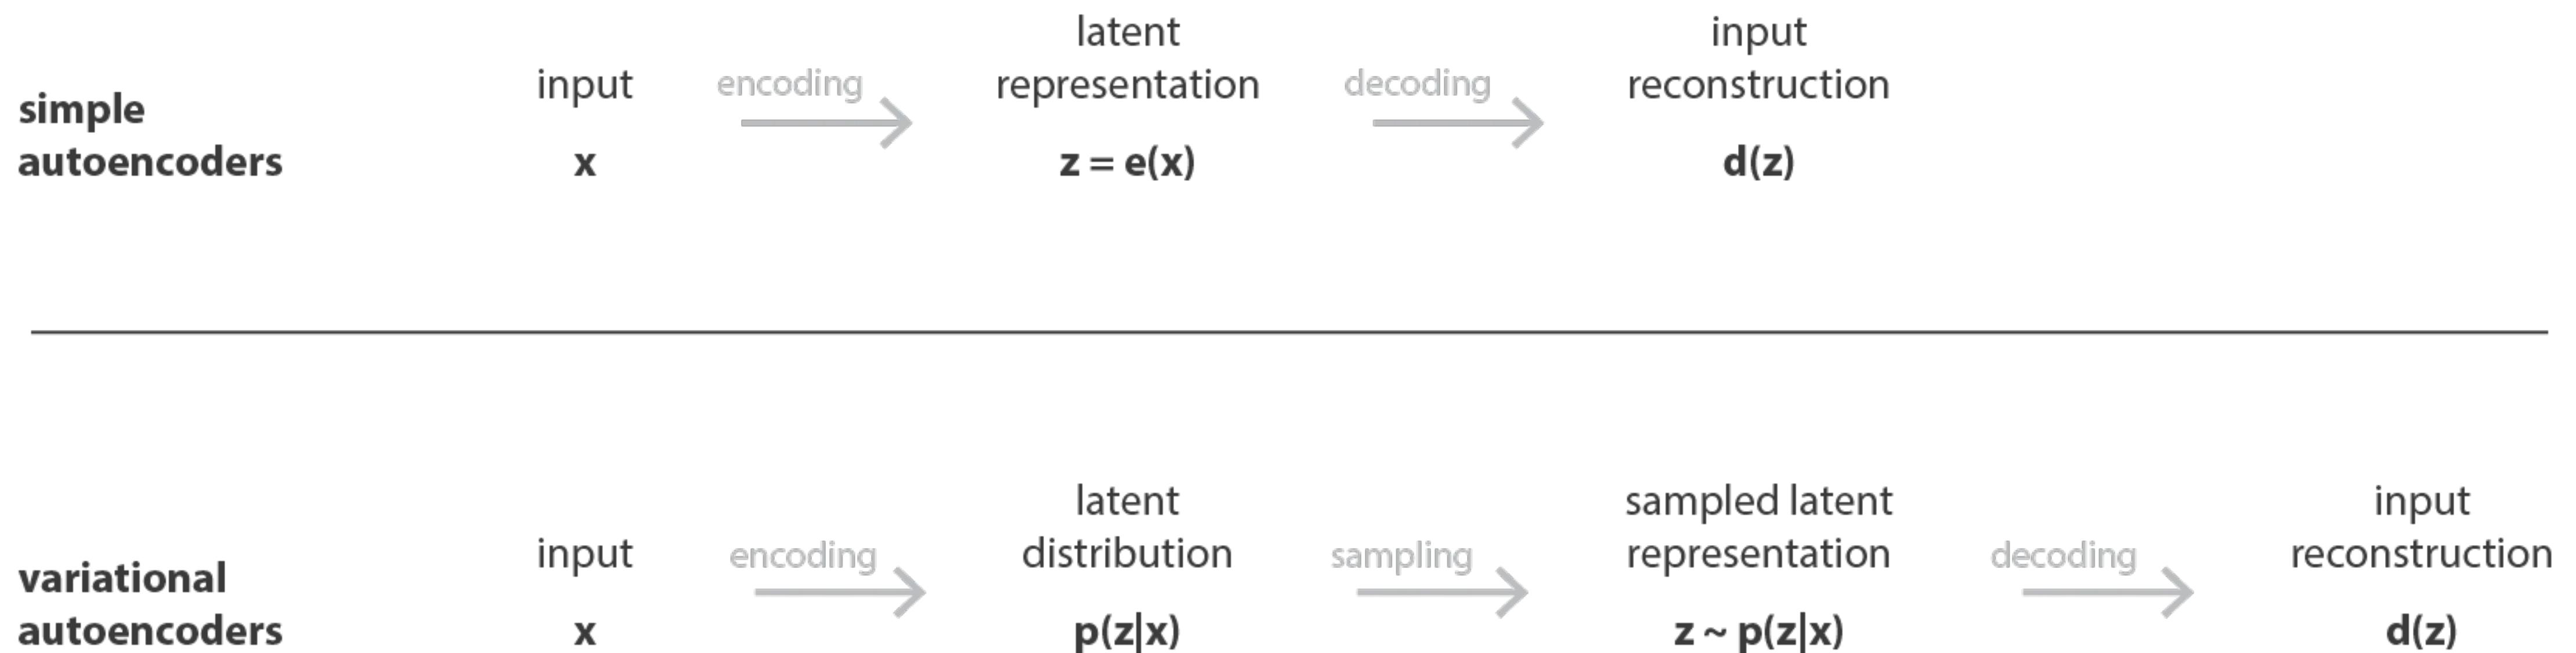
Generative Modeling via Autoencoders

- Generate realistic images from random noise
 - **Encoder:** predict means and standard deviations of a *probability distribution* over the latent features
 - **Decoder:** given a random sample from the latent distributions, produce the corresponding output



Variational Autoencoder (VAEs)

Generative Modeling via Autoencoders



VAE Notebook

neural-network-zoo_VAEs.ipynb



<https://tinyurl.com/2fx8smzk>

<https://github.com/GageDeZoort/neural-network-zoo/tree/main>